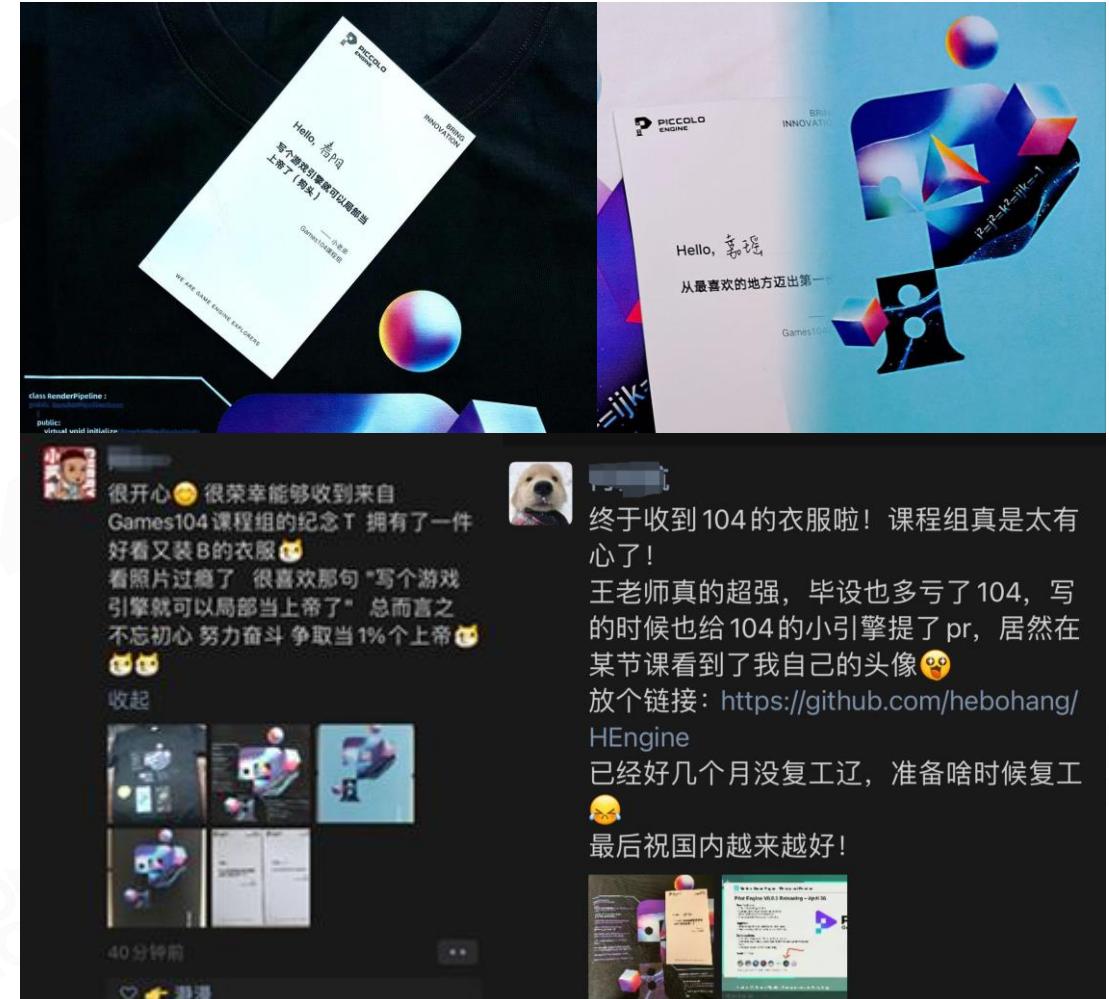
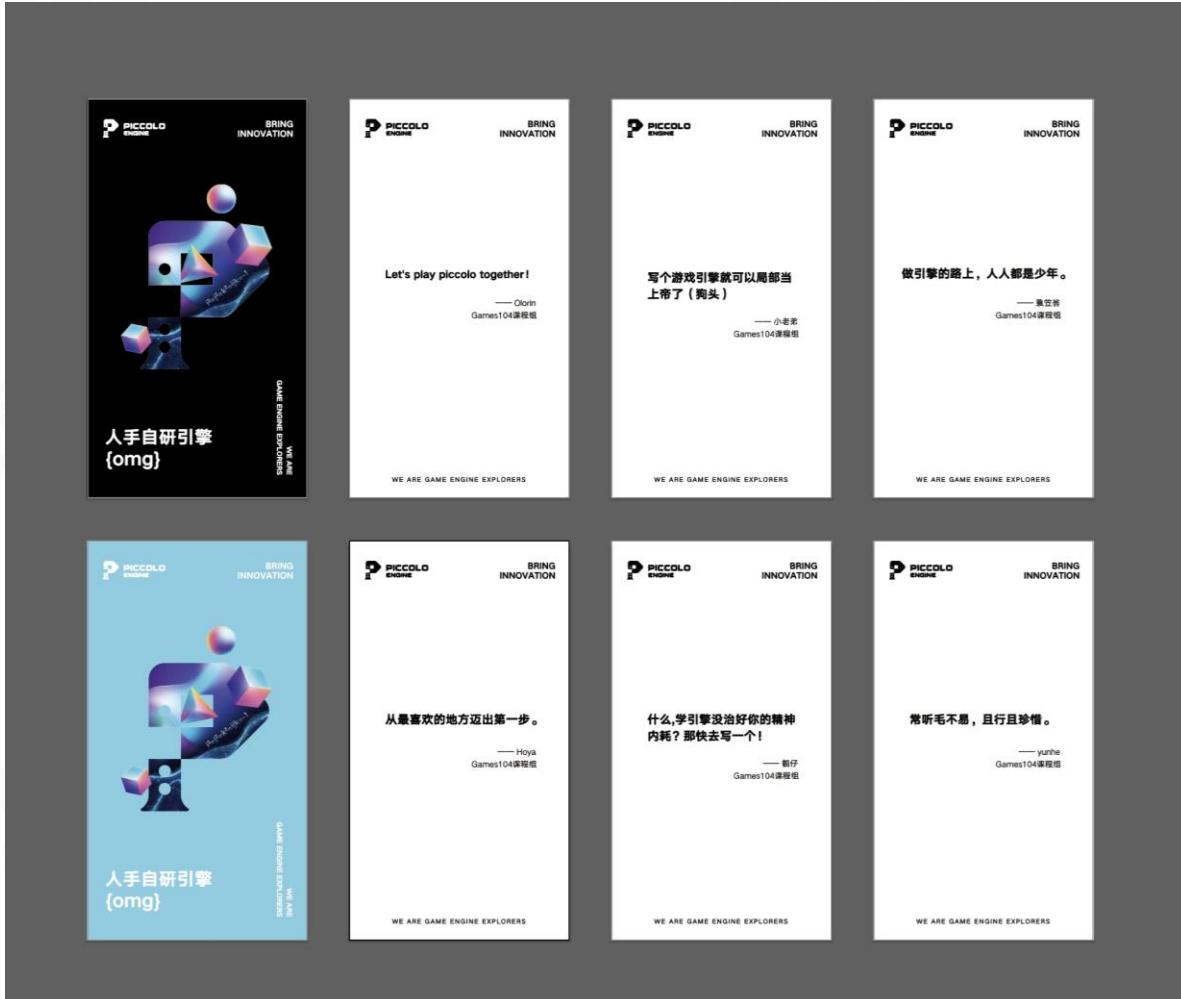




## Our Best Wishes to Every Explorer





## Piccolo Engine v0.0.8 Released – 12 September

- GPU-based Particle System!



- Piccolo Code Explained - 10 October

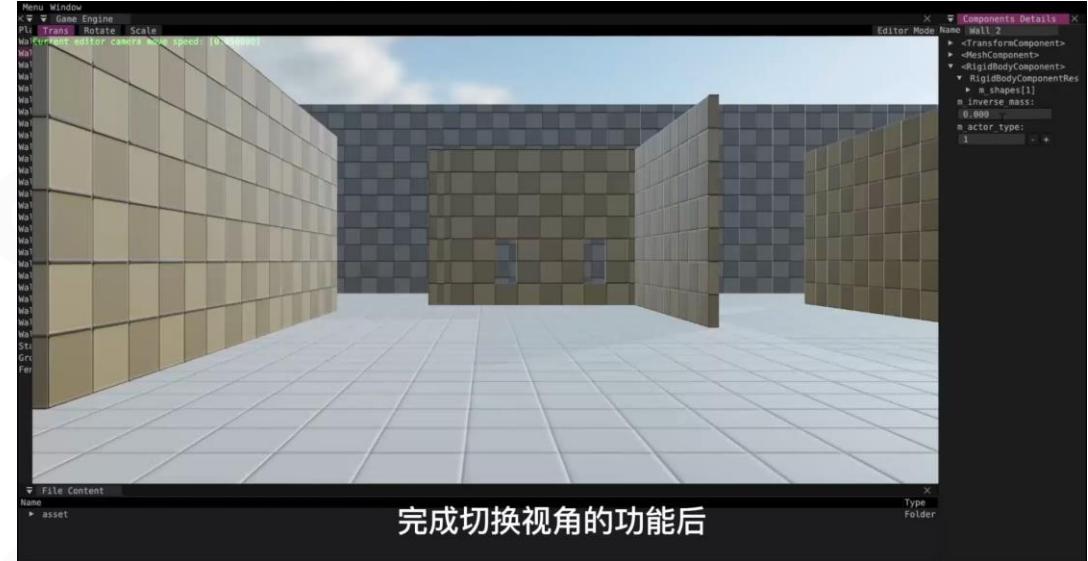
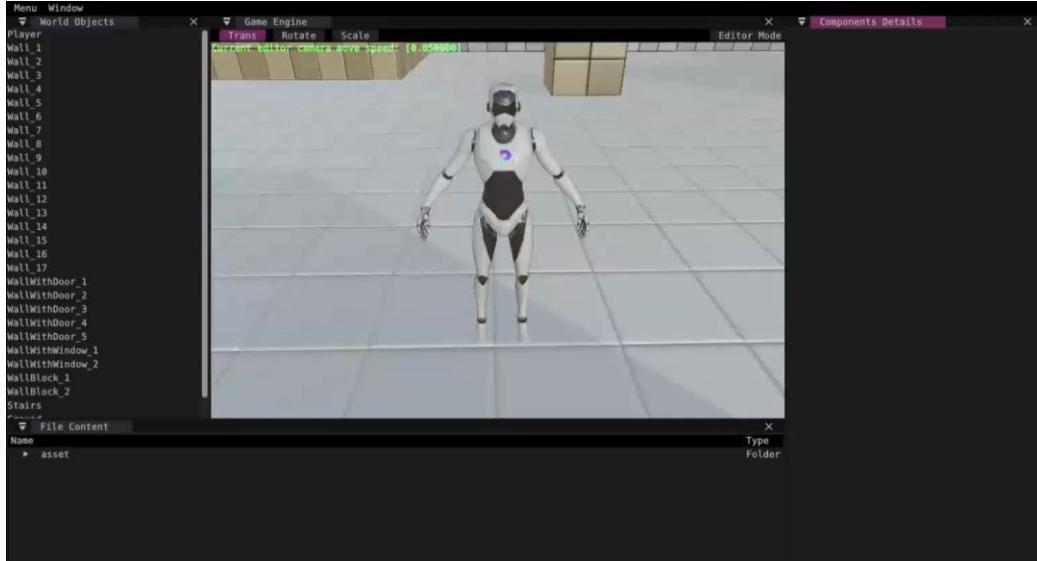


关注公众号，回复【入群】，加入我们的社区





## Homework Showcase (1/2)



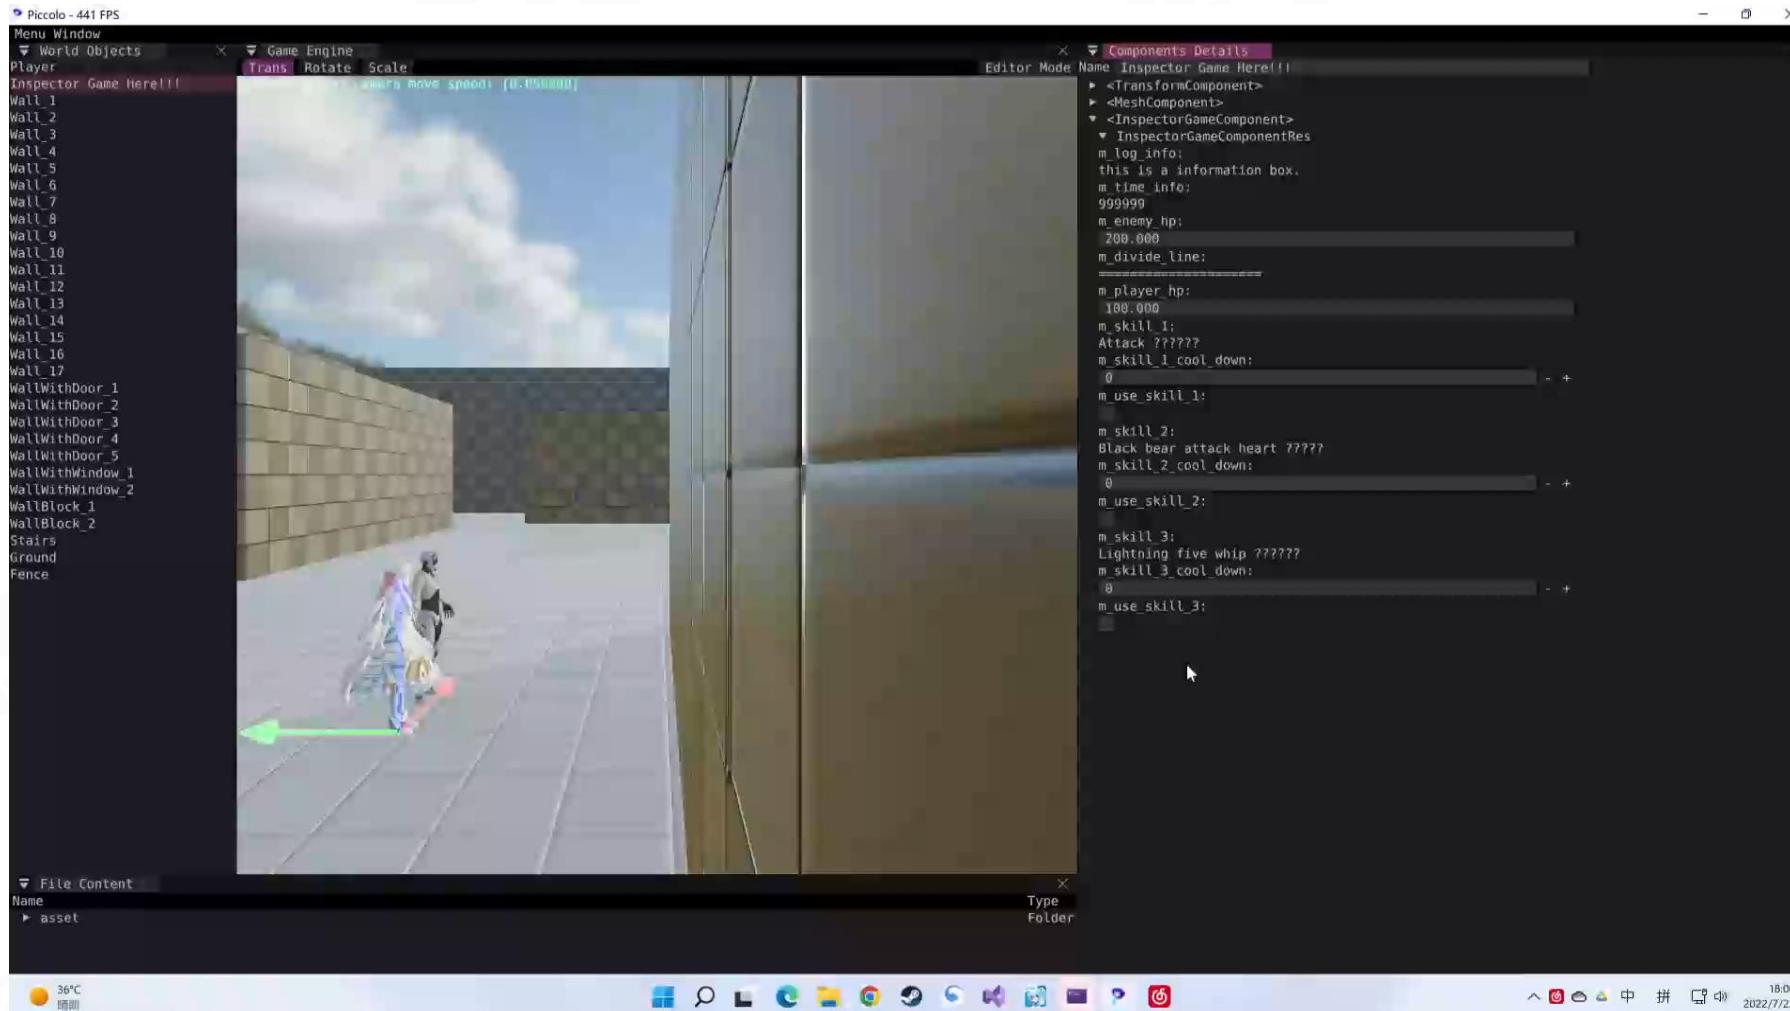
1. Added color type reflection UI
2. Added changing mesh base color function

Added camera mode change function



## Homework Showcase (2/2)

A mini game!





## Q&A

- Q1: How does ECS handle destroy of entities?
- Q2: How can we measure Cache miss?
- Q3: How should we provide tools for designers to design functions under DOP architecture?



## Lecture 21

# Dynamic Global Illumination and Lumen

---

Advanced Topics



## Global Illumination(GI)

# The Rendering Equation

James Kajiya, "The Rendering Equation."

SIGGRAPH 1986.

Energy equilibrium:

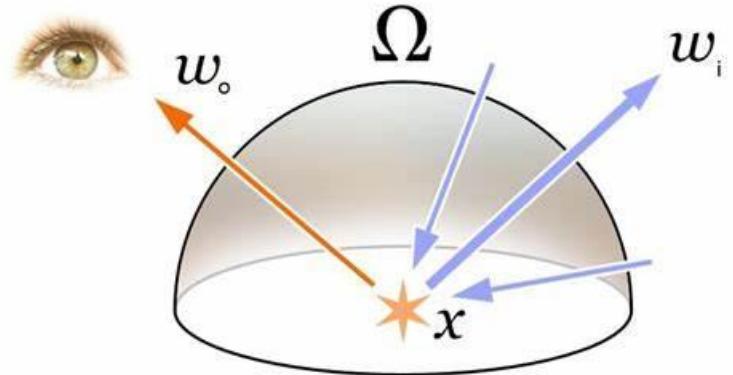


$$L_o(x, \omega_o) = L_e(x, \omega_o) + \int_{H^2} f_r(x, \omega_o, \omega_i) L_i(x, \omega_i) \cos \theta_i d\omega_i$$

outgoing            emitted            reflected

Radiance and Irradiance

# The Rendering Equation

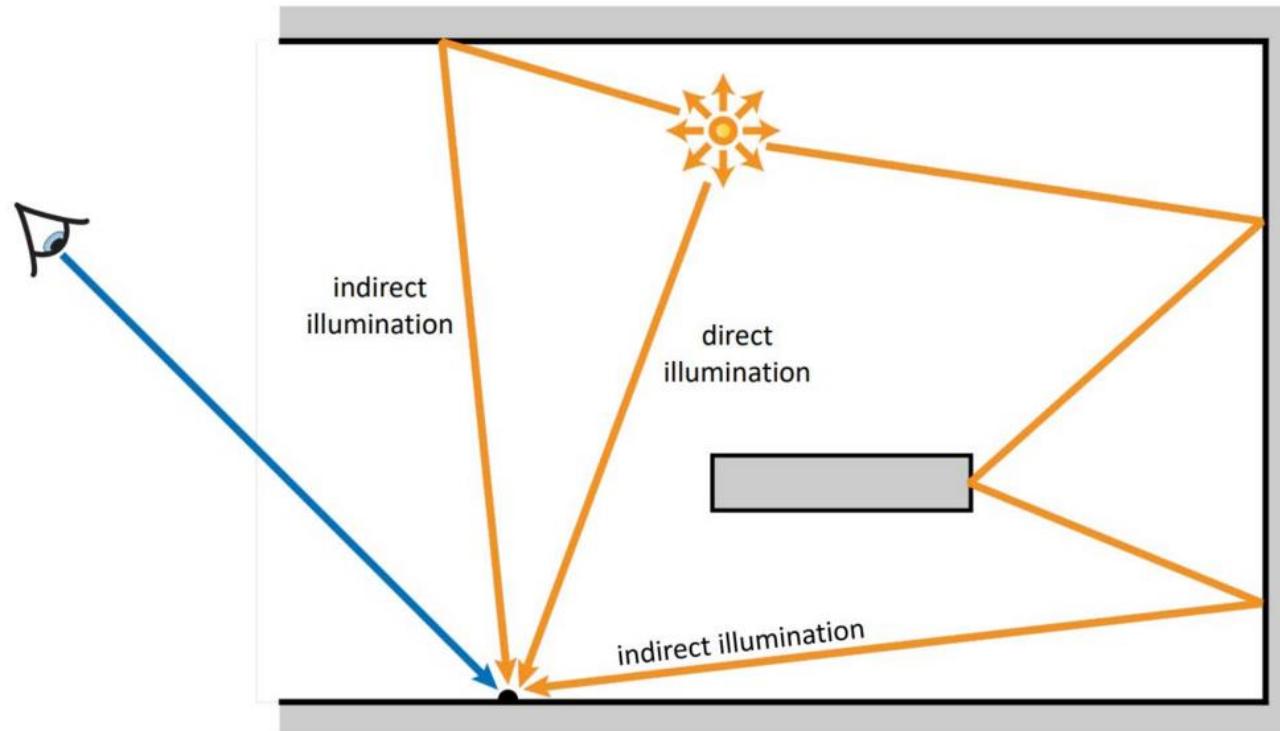


$$L_o(x, \omega_o) = L_e(x, \omega_o) + \int_{H^2} f_r(x, \omega_o, \omega_i) L_i(x, \omega_i) \cos \theta_i d\omega_i$$

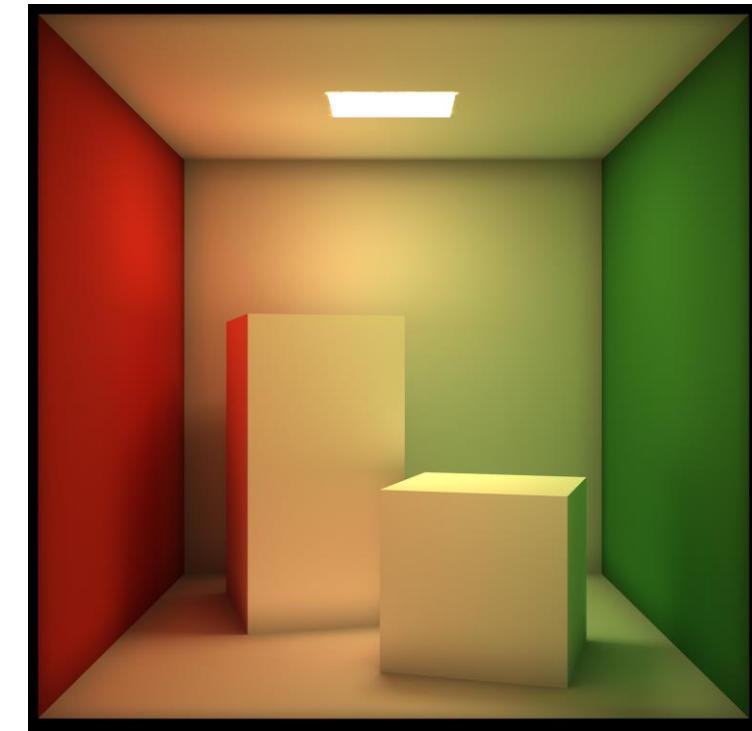
Annotations for the rendering equation:

- outgoing/observed radiance
- emitted radiance (e.g., light source)
- point of interest
- direction of interest
- all directions in hemisphere
- scattering function
- incoming radiance
- angle between incoming direction and normal
- incoming direction

# Global Illumination: Billions of Light Source



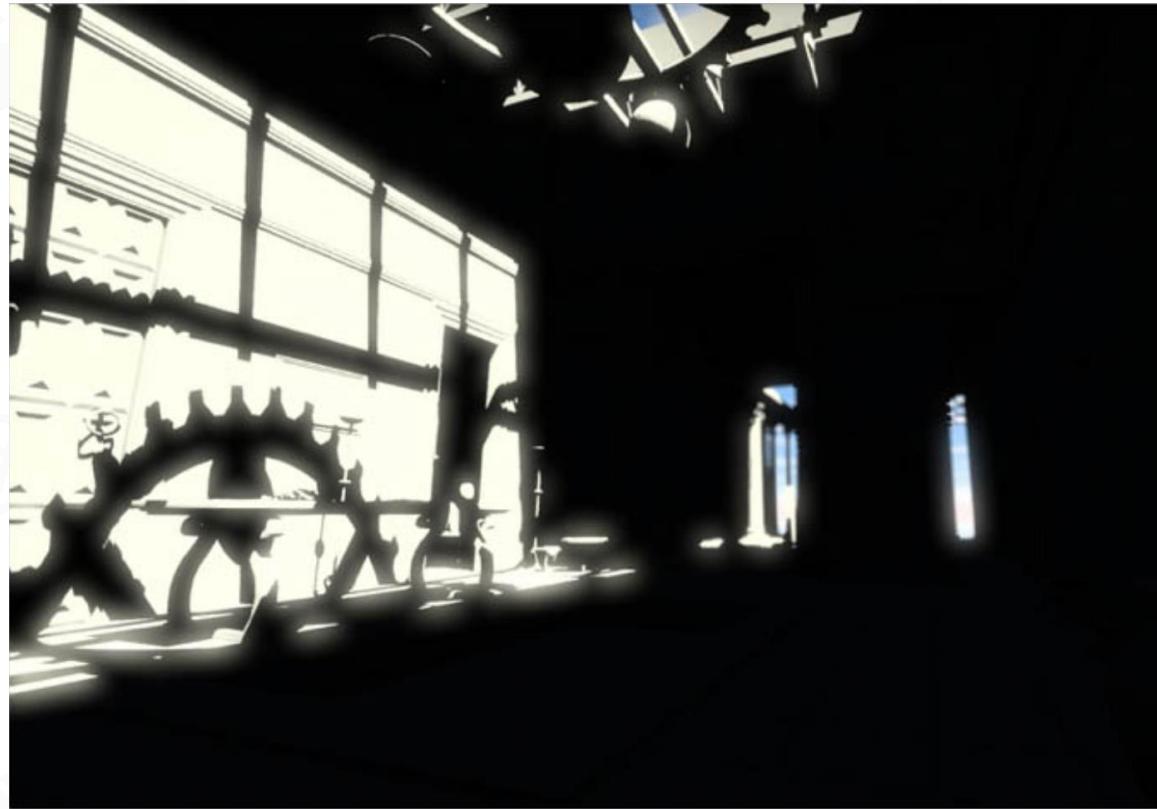
Direct vs. Indirect Illumination



Global Illumination (GI)

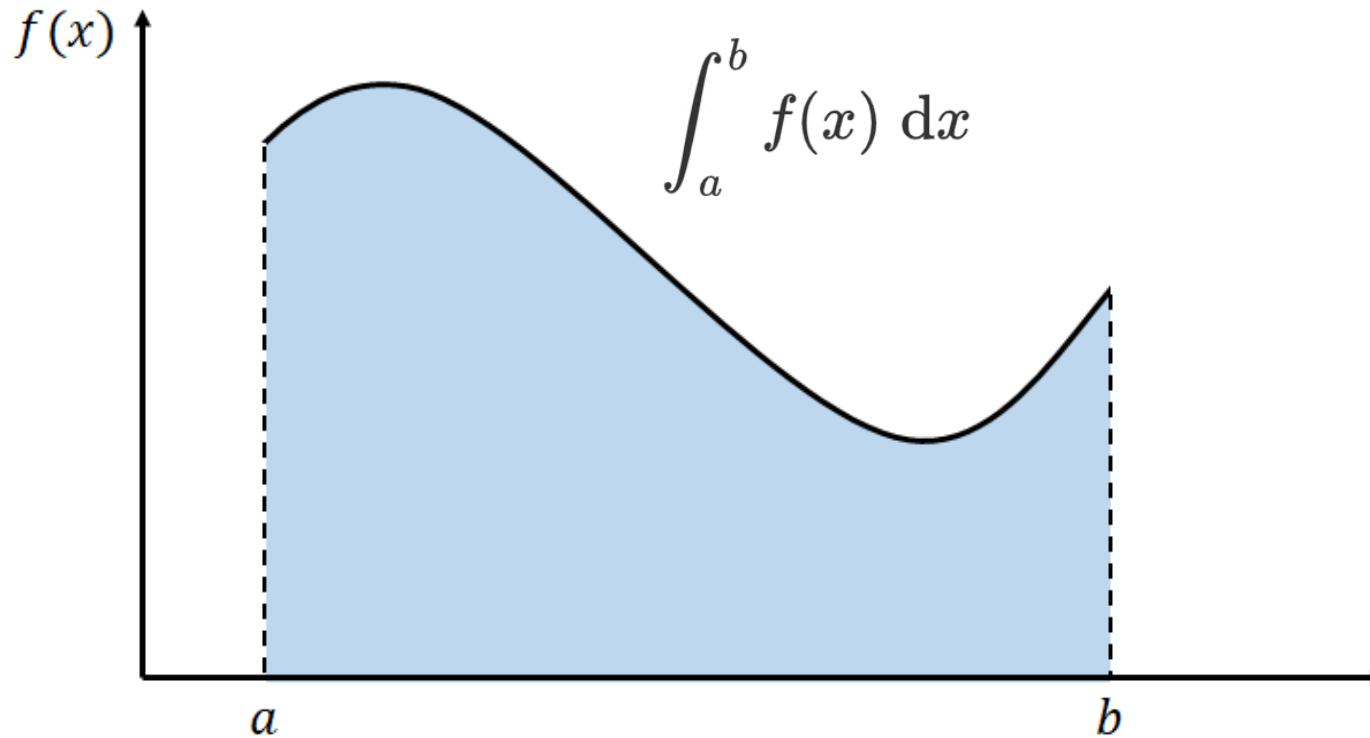


## Global Illumination is Matter for Gaming





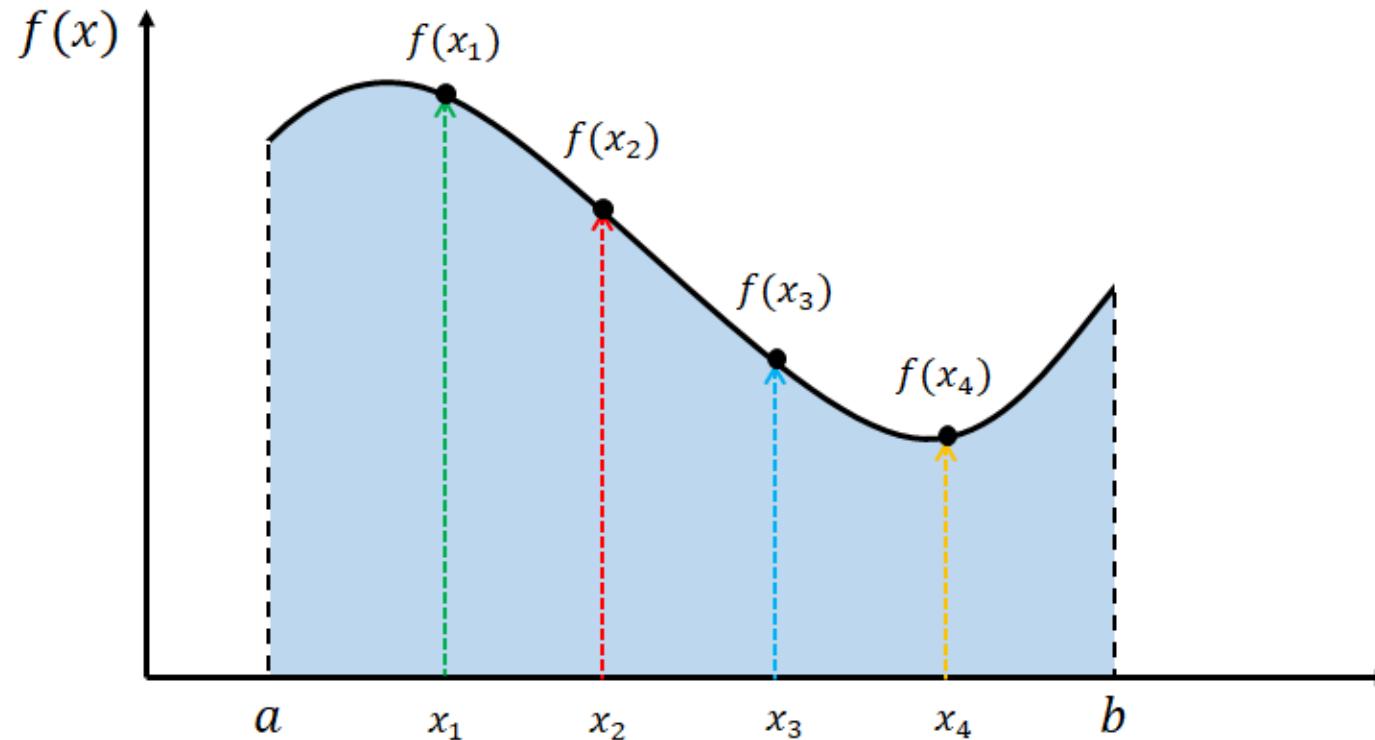
## Monte Carlo Integration



- How to solve an integral, when it's too hard to solve it analytically?



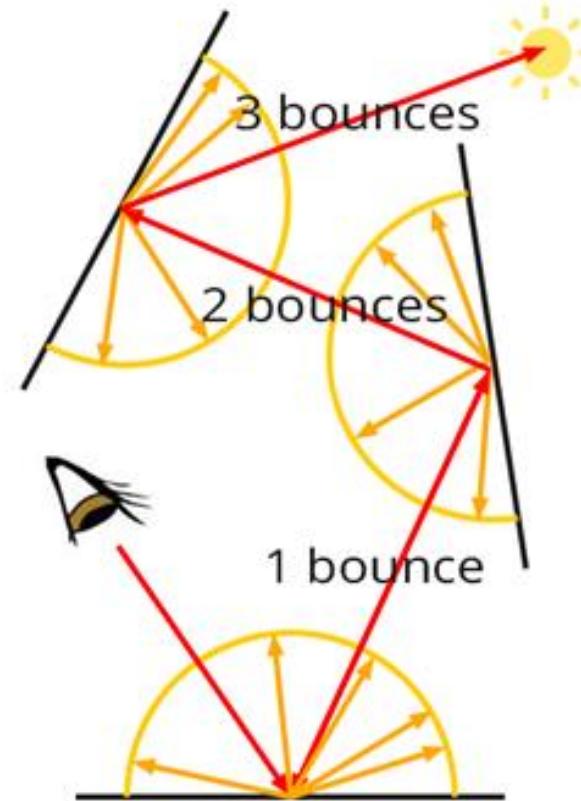
# Monte Carlo Integration



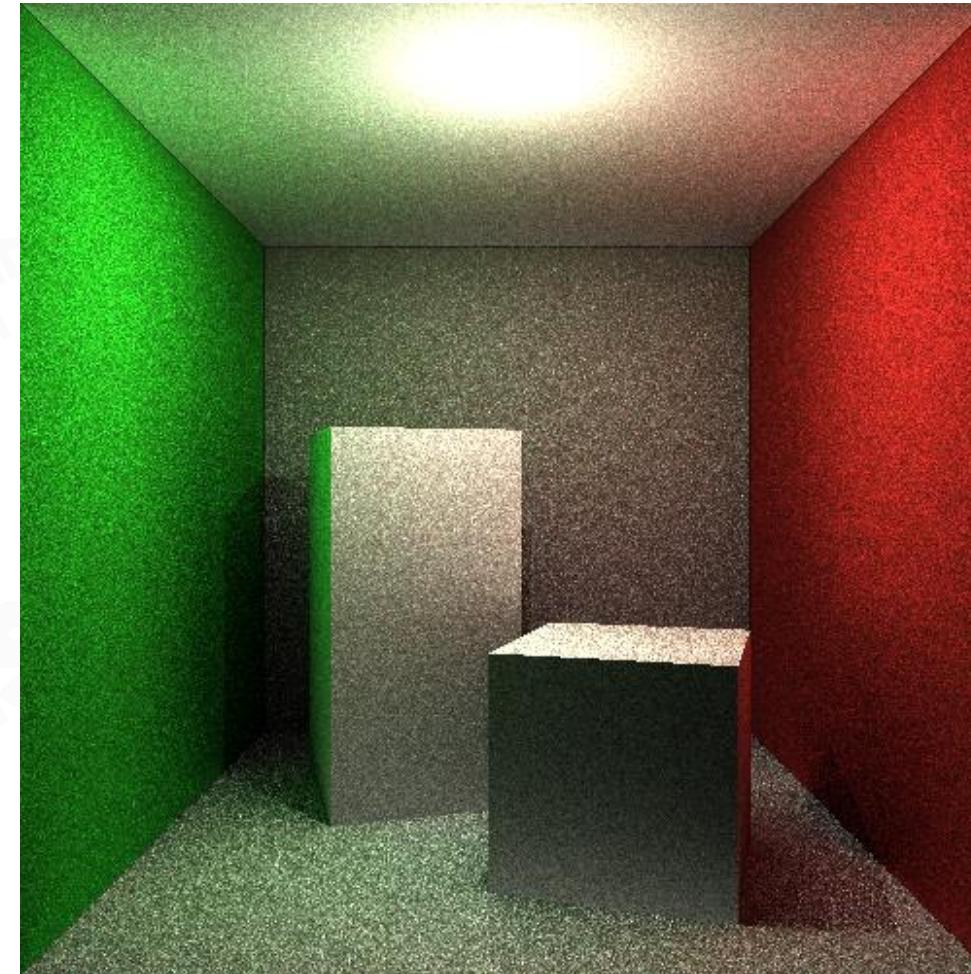
- Approximate integral with the average of randomly sample values



## Monte Carlo Ray Tracing (Offline)



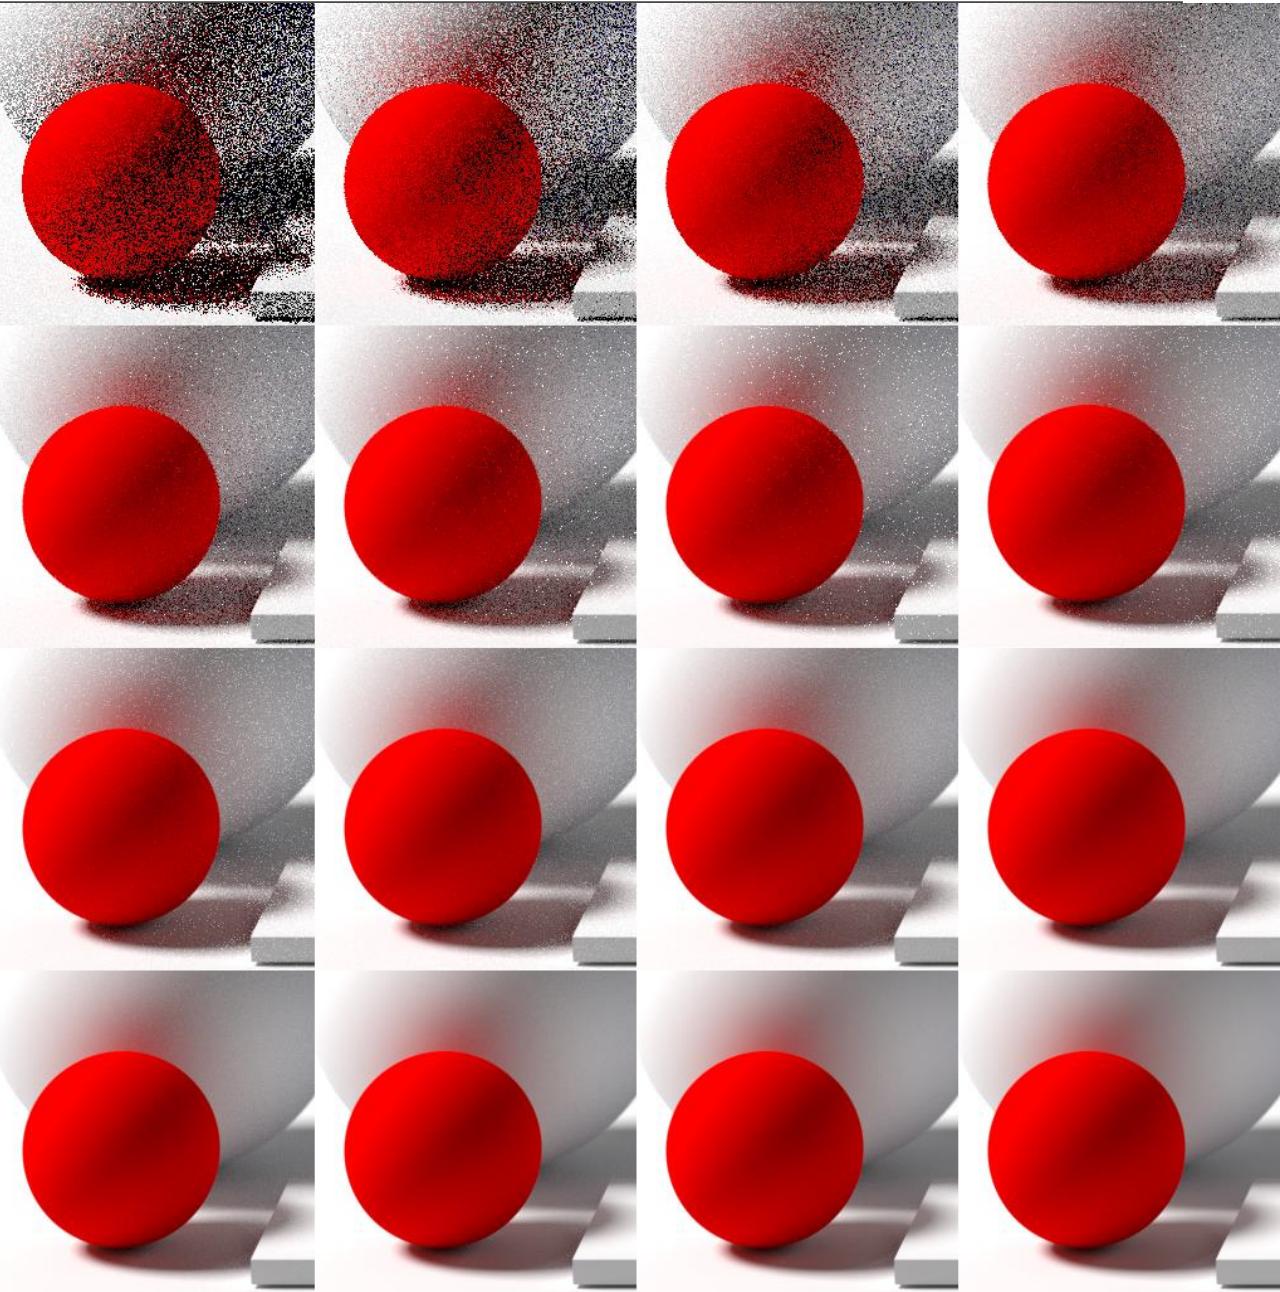
© www.scratchapixel.com





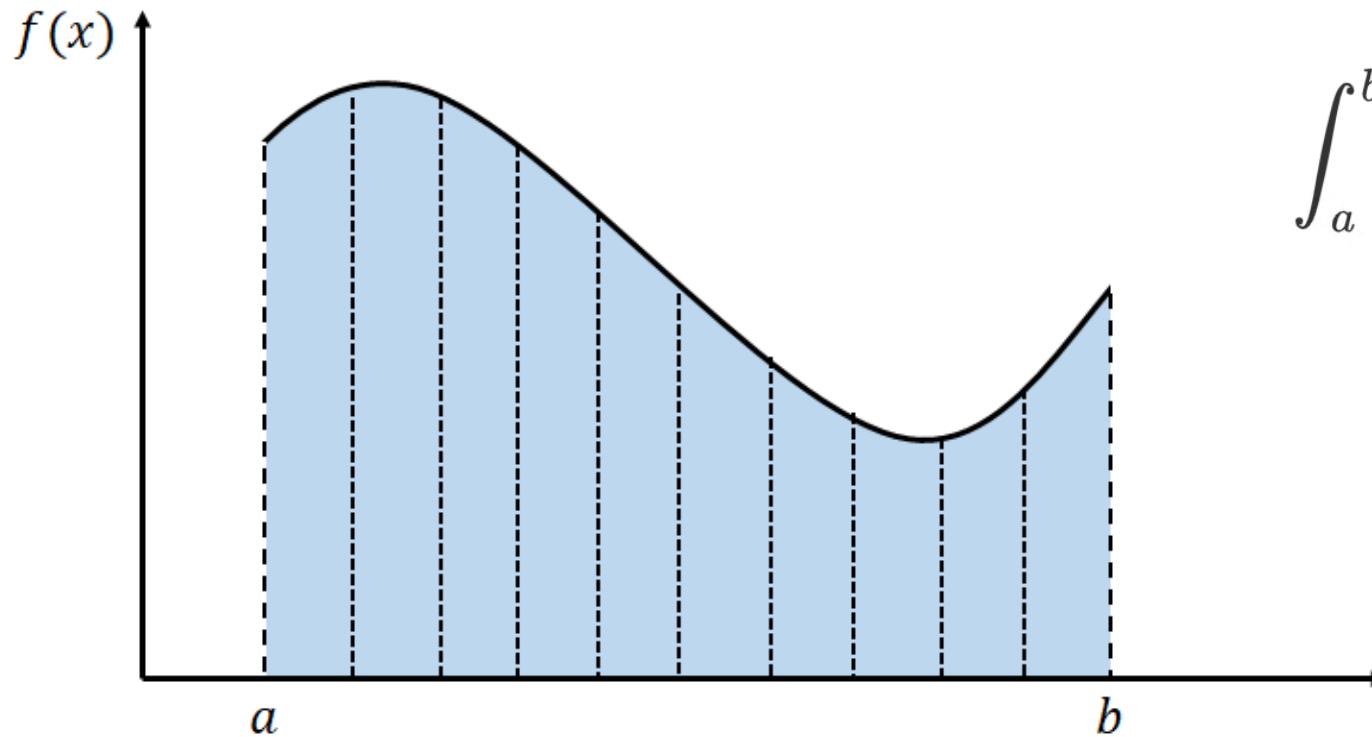
## Sampling is the Key

- Noise decreases as the number of samples per pixel increases. The top left shows 1 sample per pixel, and doubles from left to right each square.





## Sampling : Uniform Sampling



$$\int_a^b f(x) \, dx = \lim_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n f(x_i)(b - a)$$

$$= \lim_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n \frac{f(x_i)}{\frac{1}{b-a}}$$

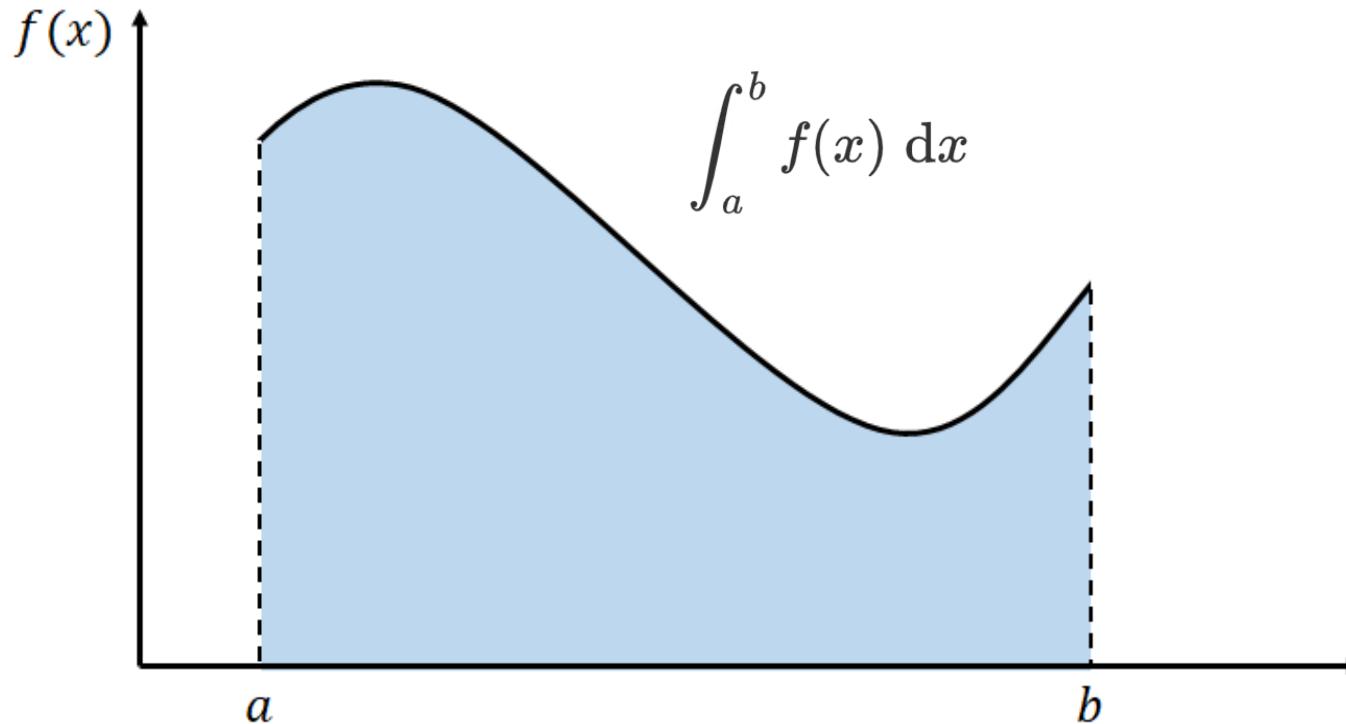


Probability Density Function

- We are doing uniform random sample, so we have  $\frac{1}{b-a}$  factor here



# Probability Distribution Function



$$\int_a^b f(x)dx \sim F_n(X) = \frac{1}{n} \sum_{k=1}^n \frac{f(X_k)}{\textcolor{red}{PDF}(X_k)}$$

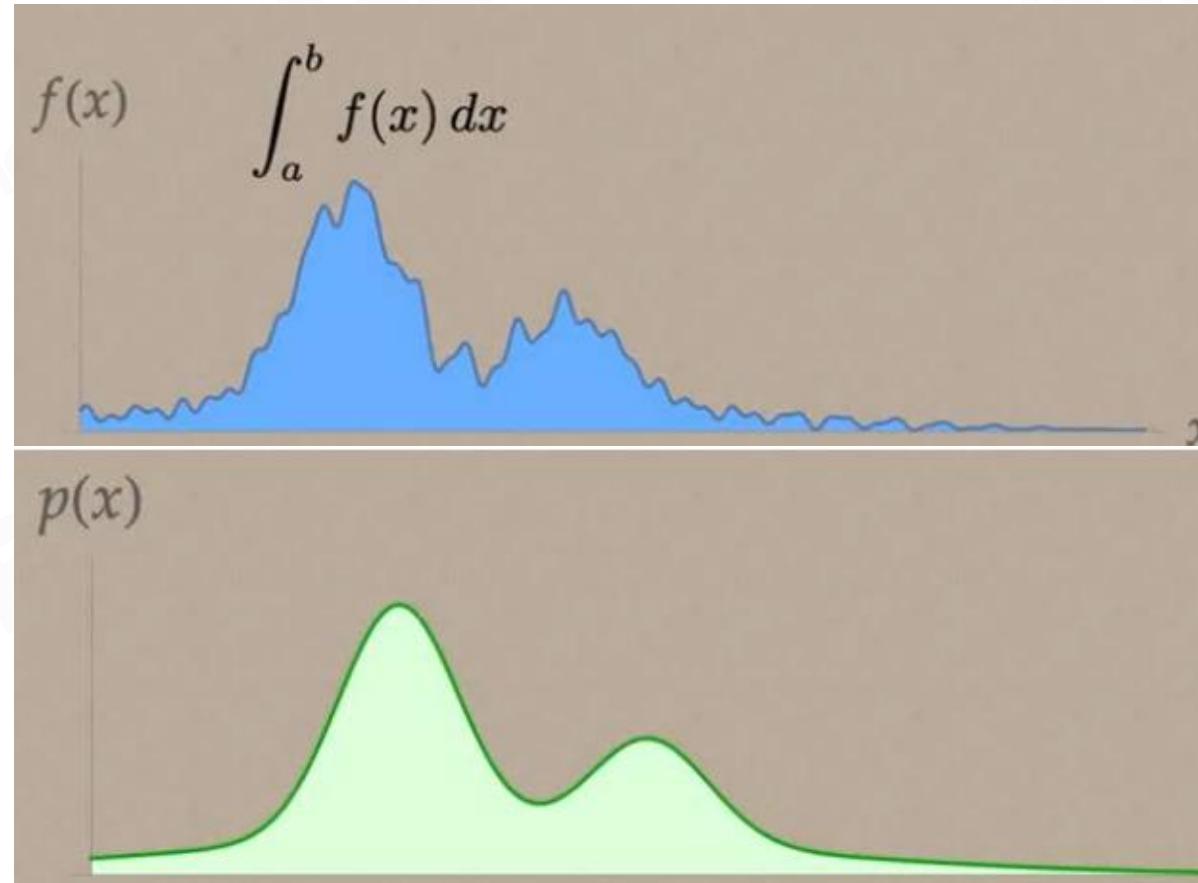
## Probability Distribution Function

- Describes the relative likelihood for this random variable to take on a given value
- Higher means more possible to be chosen



## Importance Sampling

The PDF can be arbitrary, but which is the best?





# Importance Sampling : Best PDF for Rendering?

- Rendering equation:

$$L_o(p, \omega_o) = \int_{\Omega^+} L_i(p, \omega_i) f_r(p, \omega_i, \omega_o) (n \cdot \omega_i) d\omega_i$$

- Monte Carlo Integration:

$$L_o(p, \omega_o) \approx \frac{1}{N} \sum_{i=1}^N \frac{L_i(p, \omega_i) f_r(p, \omega_i, \omega_o) (n \cdot \omega_i)}{p(\omega_i)}$$

- What's our  $f(x)$  ?

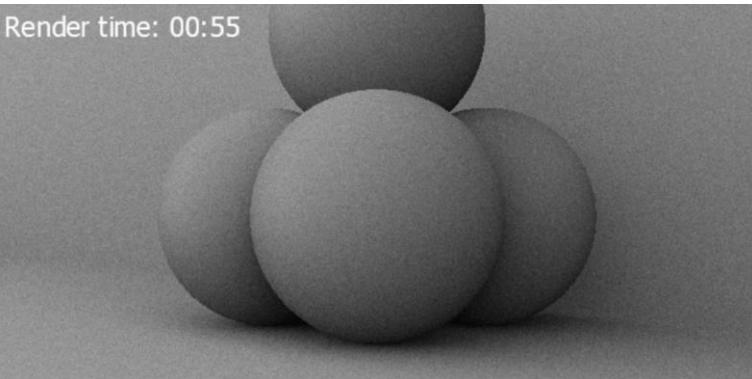
$$L_i(p, \omega_i) f_r(p, \omega_i, \omega_o) (n \cdot \omega_i)$$

- What's our pdf ?

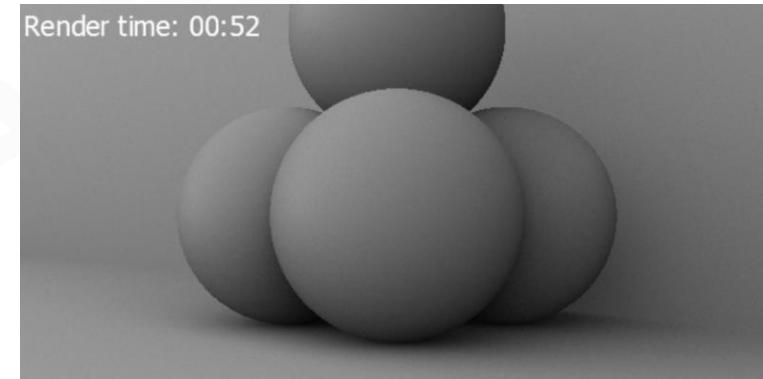
- Uniform:  $p(\omega_i) = \frac{1}{2\pi}$
- Other pdf ? (cosine-weight, GGX)



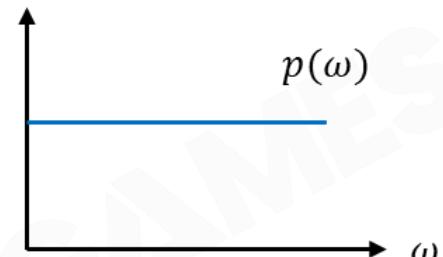
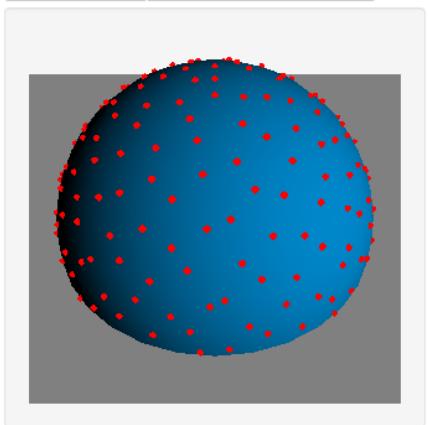
## Importance Sampling : PDF is Matter



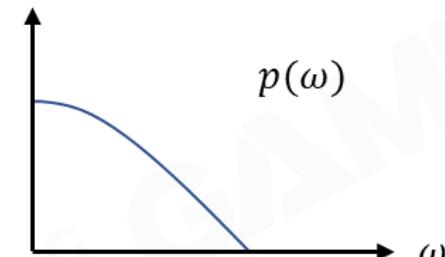
uniform sampling 256spp



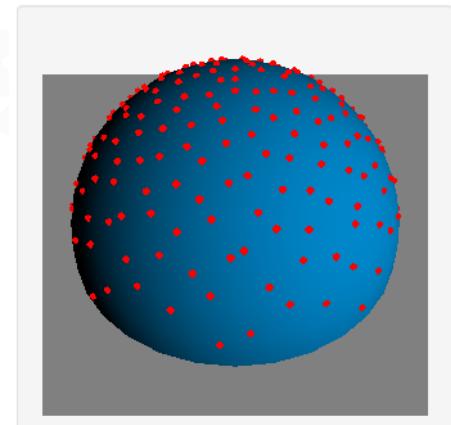
cosine weights importance sampling 256spp



$$p(\omega) = \frac{1}{2\pi}$$



$$p(\omega) = \frac{\cos \theta}{\pi}$$



spp: samples per pixel



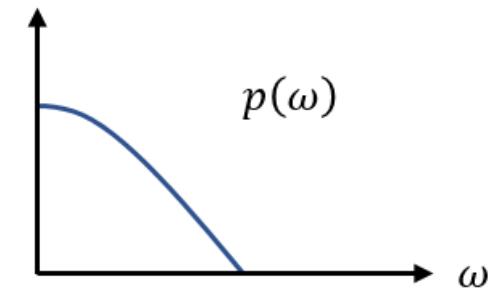
## Importance Sampling : Cosine and GGX PDF



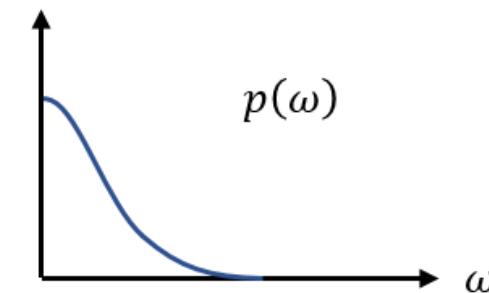
Before



After



$$p(\omega) = \frac{\cos \theta}{\pi}$$



$$p(\omega) = \frac{\alpha^2 \cos \theta}{\pi((\alpha^2 - 1) \cos^2 \theta + 1)^2}$$

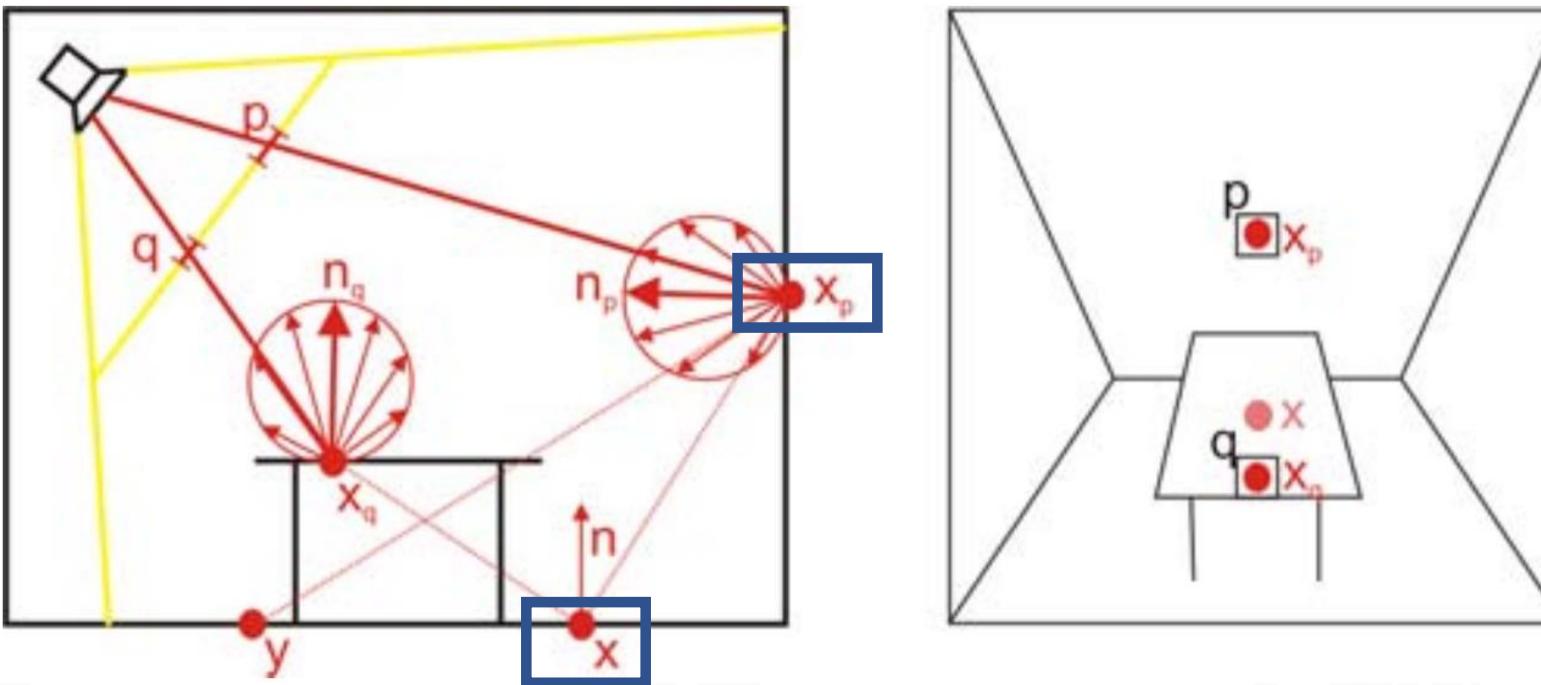


## Reflective Shadow Maps (RSM, 2005)

Let's inject light in. (Photon Mapping?)



- Each pixel on the shadow map is a indirect light source



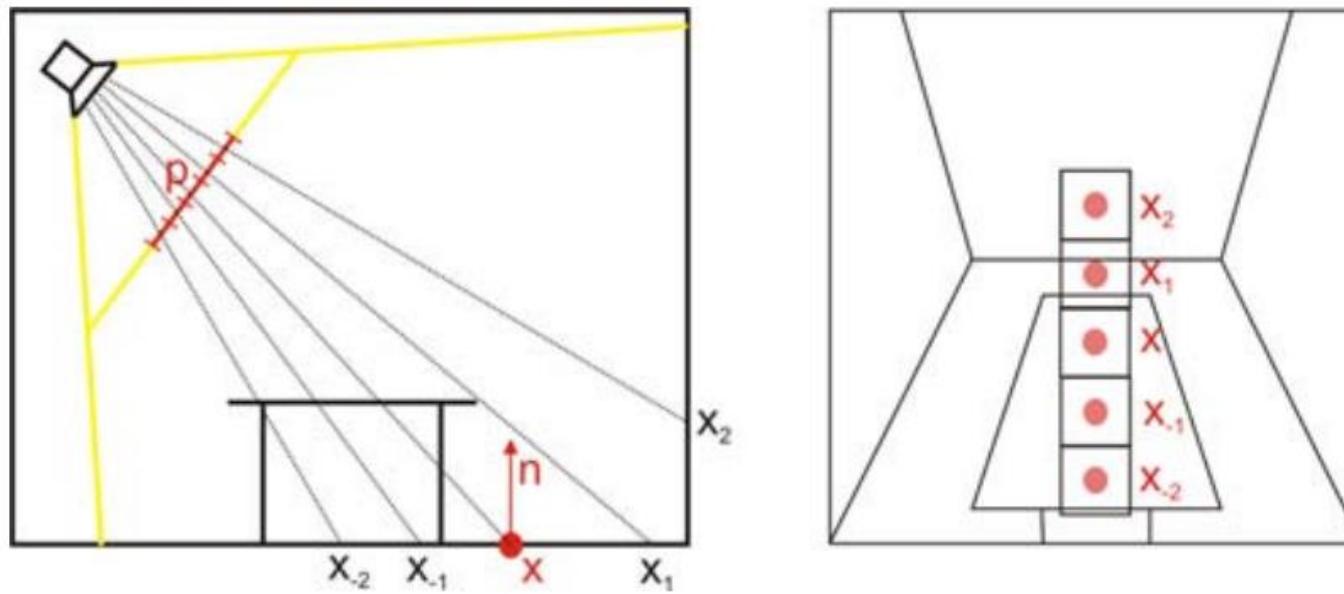
- How the RSM pixel  $x_p$  illuminates position  $x$ ?

$$E_p(x, n) = \phi_p \frac{\max\{0, \langle n_p | x - x_p \rangle\} \max\{0, \langle n | x_p - x \rangle\}}{\|x - x_p\|^4}$$



- The indirect irradiance at a surface point  $x$  can be approximated by summing up the illumination due to all pixel lights.
- Do not consider occlusion.

$$E(x, n) = \sum_{\text{pixels } p} E_p(x, n)$$





# Cone Tracing with RSM

- Gathering Indirect Illumination
  - random sampling RSM pixels
  - precompute such a sampling pattern and reuse it for all indirect light computations
    - 400 samples were sufficient
    - use Poisson sampling to obtain a more even sample distribution

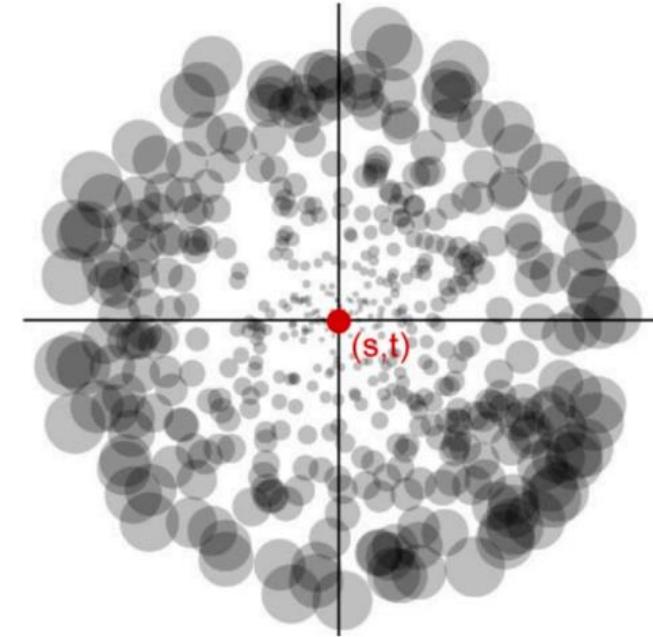
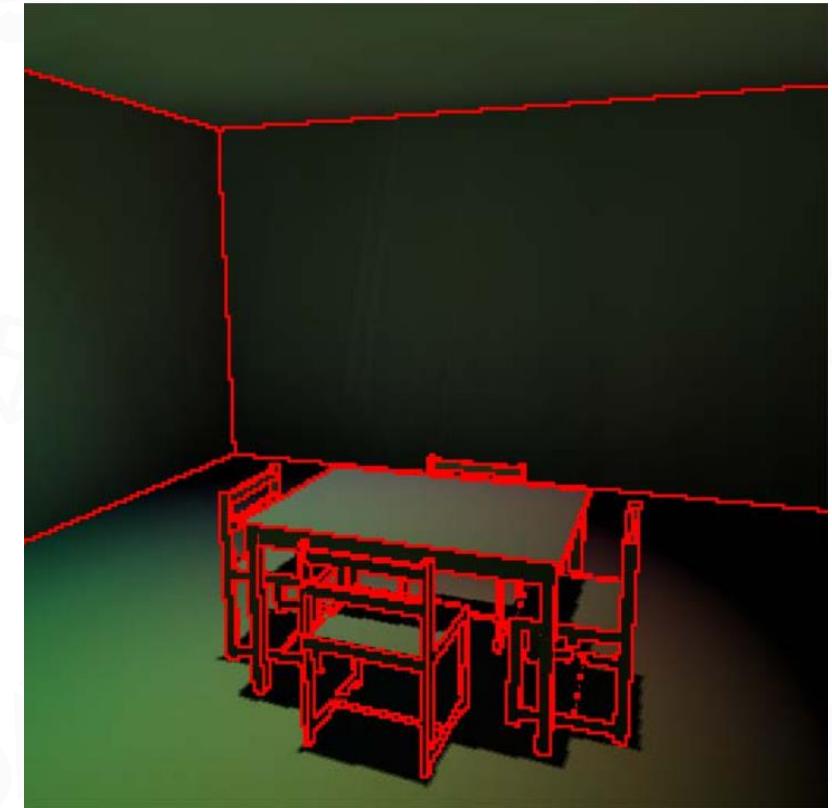


Figure 4: Sampling pattern example. The sample density decreases and the sample weights (visualized by the disk radius) increases with the distance to the center.



## Acceleration with Low-Res Indirect Illumination

- Compute the indirect illumination for a low resolution image
- For each pixel on full resolution:
  - get its four surrounding low-res samples
  - validate by comparing normal and world space position
  - bi-linear interpolation
- Recompute the left (**red** pixels)





Gears of War 4, Uncharted 4, The Last of US, etc





## Thanks, RSM

### Cool Ideas

- Easy to be implemented
- Photon Injection with RSM
- Cone sampling in mipmap
- Low-res Indirect illumination with error check

### Cons

- Single bounce
- No visibility check for indirect illumination

### Reflective Shadow Maps

Carsten Dachsbacher\*  
University of Erlangen-Nuremberg

Marc Stamminger†  
University of Erlangen-Nuremberg

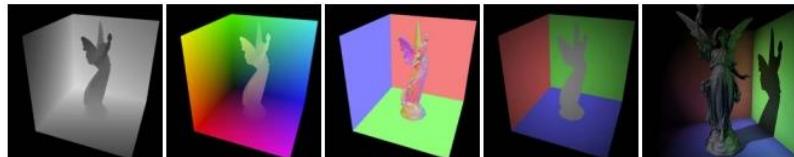


Figure 1: This figure shows the components of the reflective shadow map (depth, world space coordinates, normal, flux) and the resulting image rendered with indirect illumination from the RSM. Note that the angular decrease of flux is shown exaggerated for visualization.

#### Abstract

In this paper we present "reflective shadow maps", an algorithm for interactive rendering of plausible indirect illumination. A reflective shadow map is an extension to a standard shadow map, where every pixel is considered as an indirect light source. The illumination due to these indirect lights is evaluated on-the-fly using adaptive sampling in a fragment shader. By using screen-space interpolation of the indirect lighting, we achieve interactive rates, even for complex scenes. Since we mainly work in screen space, the additional effort is largely independent of scene complexity. The resulting indirect light is approximate, but leads to plausible results and is suited for dynamic scenes. We describe an implementation on current graphics hardware and show results achieved with our approach.

**CR Categories:** I.3.3 [Computer Graphics]: Three-Dimensional Graphics and Realism—Color, shading, shadowing, and texture; I.3.3 [Computer Graphics]: Hardware Architecture—Graphics processors

**Keywords:** indirect illumination, hardware-assisted rendering

#### 1 Introduction

Interactive computer graphics has developed enormously over the last years, mainly driven by the advance of graphics acceleration hardware. Scenes of millions of polygons can be rendered in real-time on consumer-level PC cards nowadays. Programmability allows the inclusion of sophisticated lighting effects. However, these effects are only simple subcases of global illumination, e.g. reflections of distant objects or shadows of point lights. Real global illu-

mination, however, generates subtle, but also important effects that are mandatory to achieve realism.

Unfortunately, due to their global nature, full global illumination and interactivity are usually incompatible. Ray Tracing and Radiosity—just to mention the two main classes of global illumination algorithms—require minutes or hours to generate a single image with full global illumination. Recently, there has been remarkable effort to make ray tracing interactive (e.g. [Wald et al. 2003]). Compute clusters are necessary to achieve interactivity at good image resolution and dynamic scenes are difficult to handle, because they require to update the ray casting acceleration structures for every frame. Radiosity computation times are even further from interactive. Anyhow, a once computed radiosity solution can be rendered from arbitrary view points quickly, but, as soon as objects move, the update of the solution becomes very expensive again.

It has been observed that for many purposes, global illumination solutions do not need to be precise, but only plausible. In this paper, we describe a method to compute a rough approximation for the one-bounce indirect light in a scene. Our method is based on the idea of the shadow map. In a first pass, we render the scene from the view of the light source (for now, we assume that we have only one spot or parallel light source in our scene). The resulting depth buffer is called *shadow map*, and can be used to generate shadows. In a *reflective shadow map*, with every pixel, we additionally store the light reflected off the hit surface. We interpret each of the pixels as a small area light source that illuminates the scene. In this paper, we describe how the illumination due to this large set of light sources can be computed efficiently and coherently, resulting in approximate, yet plausible and coherent indirect light.

#### 2 Previous Work

Shadow maps [Williams 1978; Reeves et al. 1987] and shadow volumes [Crow 1977] are the standard shadowing algorithms for interactive applications. Recently, there have been extensions of both approaches to area lights [Assarsson and Akenine-Möller 2003; Chan and Durand 2003; Wyman and Hansen 2003]. Sometimes, such soft shadows are already referred to as "global illumination". In this paper, we concentrate on indirect illumination from point lights, but our approach can easily be combined with any of these soft shadow techniques.

\*e-mail: dachsbaucher@cs.fau.de

†e-mail: stamminger@cs.fau.de

Copyright © 2005 by the Association for Computing Machinery, Inc.  
Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers, or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from Permissions Dept, ACM Inc., fax +1 (212) 869-0481 or e-mail: [permissions@acm.org](mailto:permissions@acm.org).  
© 2005 ACM 1-59593-013-2/05/0004 \$5.00



## Light Propagation Volumes (LPV)

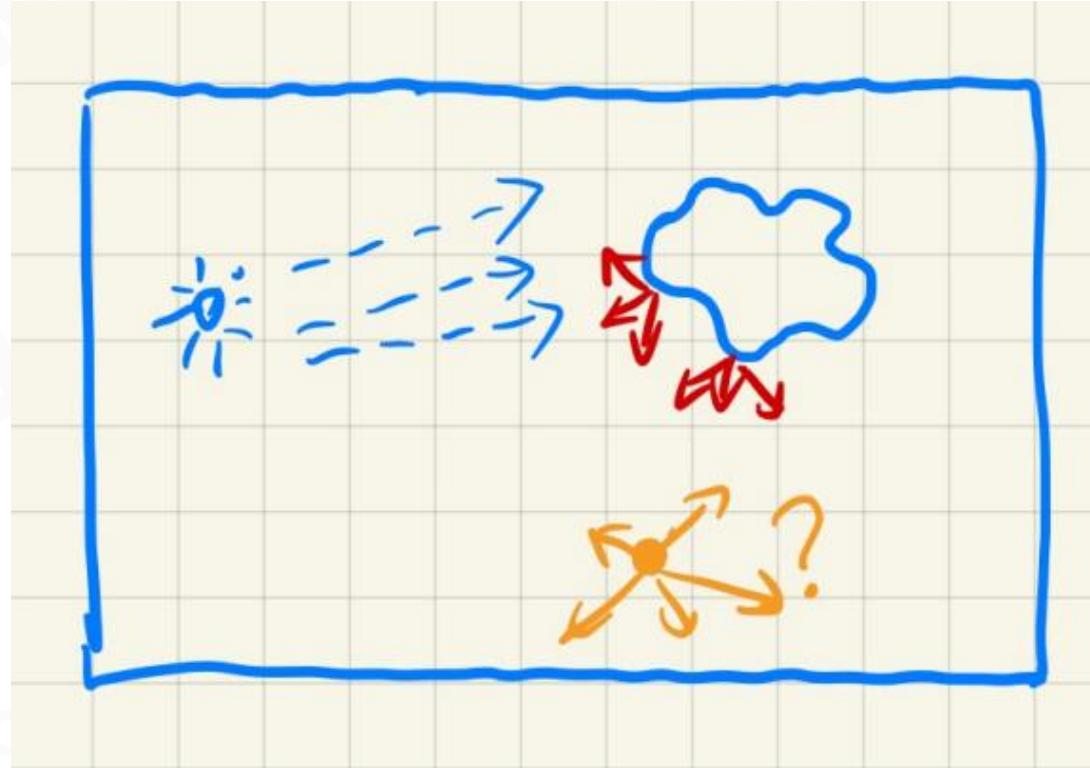


First introduced in CryEngine 3 (SIGGRAPH 2009)





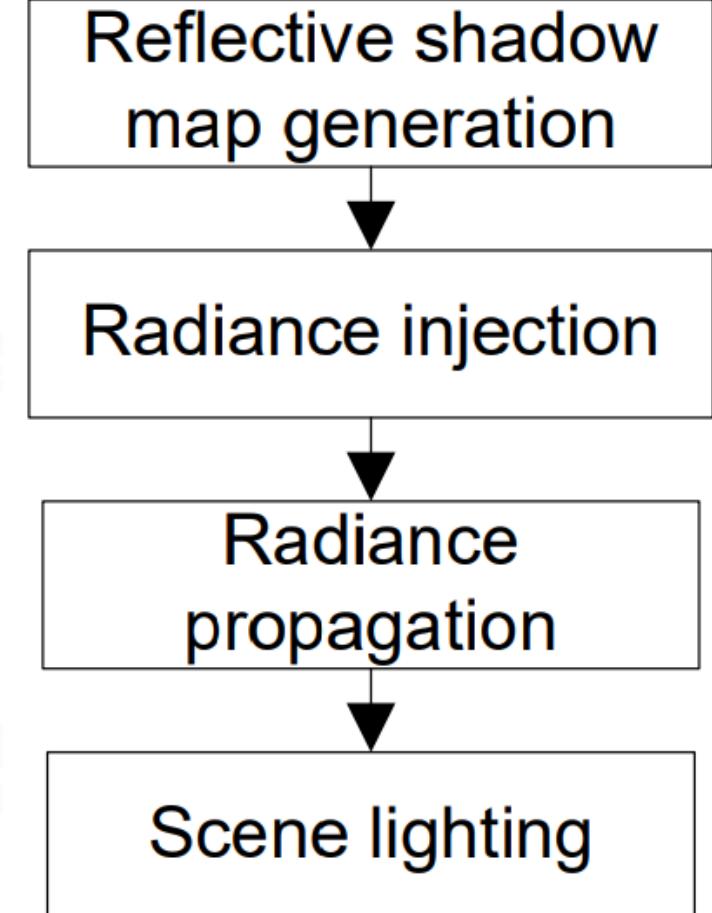
- Key Idea
  - Use a 3D grid to propagate radiance from directly illuminated surfaces to anywhere else





## Steps

1. Generation of radiance point set scene representation
2. Injection of point cloud of virtual light sources into radiance volume
3. Volumetric radiance propagation
4. Scene lighting with final light propagation volume

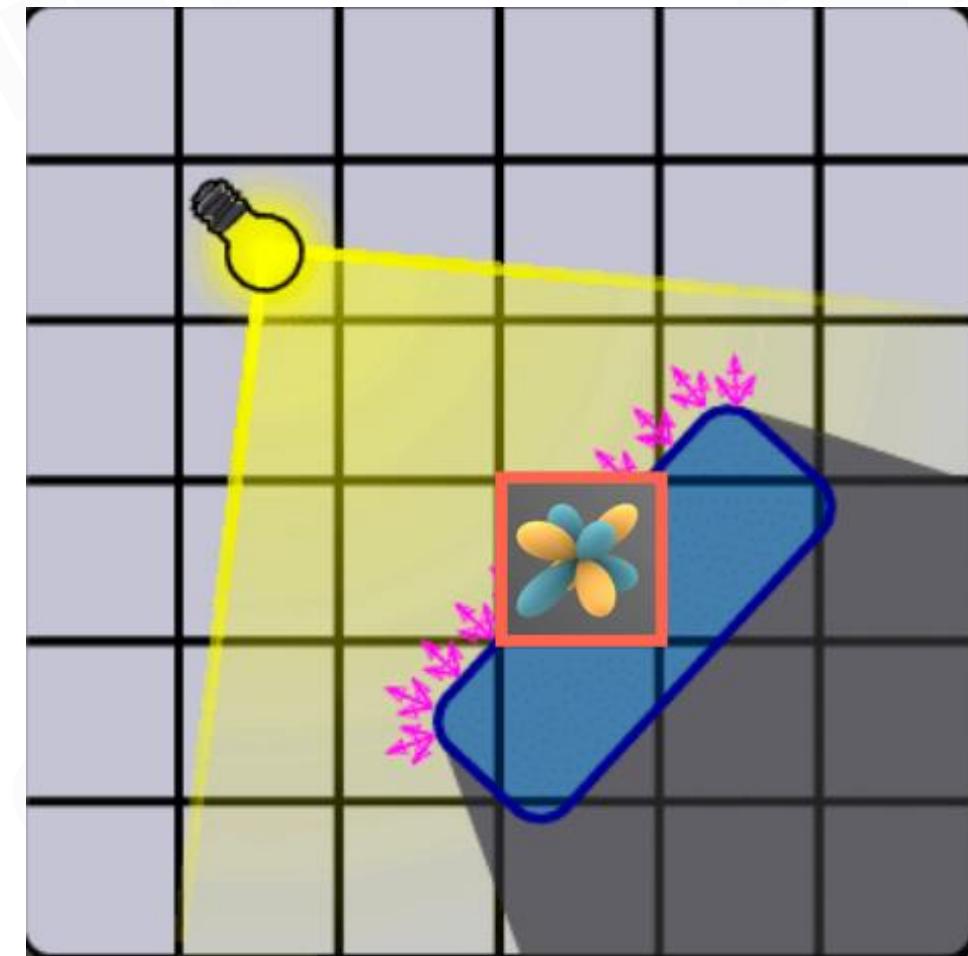




## “Freeze” the Radiance in Voxel

### Light Injection

- Pre-subdivide the scene into a 3D grid
- For each grid cell, find enclosed virtual light sources
- Sum up their directional radiance distribution
- Project to first 2 orders of SHs (4 in total)





# Radiance Propagation

- For each grid cell, collect the radiance received from each of its 6 faces
- Sum up, and again use SH to represent
- Repeat this propagation several times till the volume becomes stable

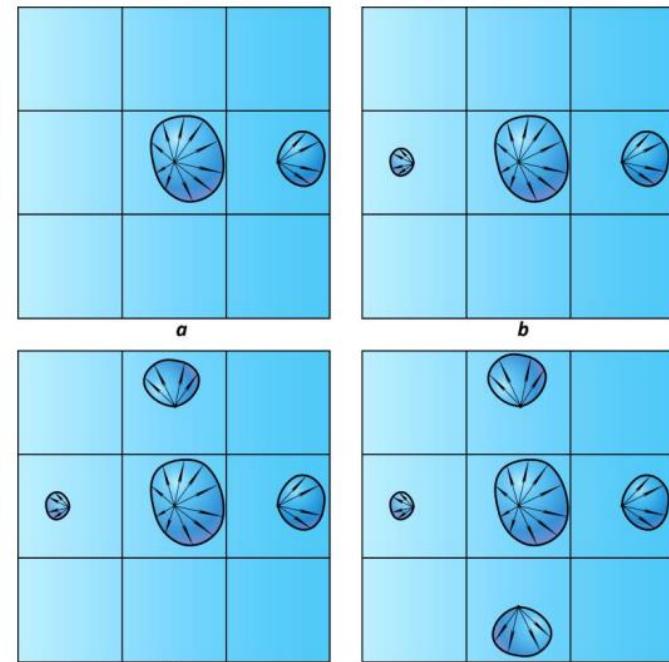
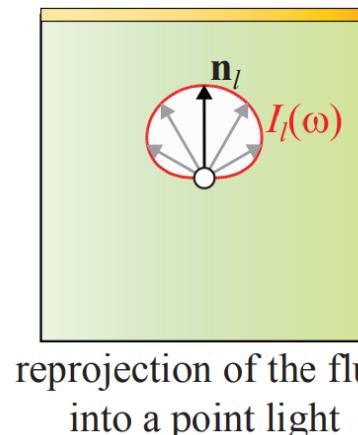
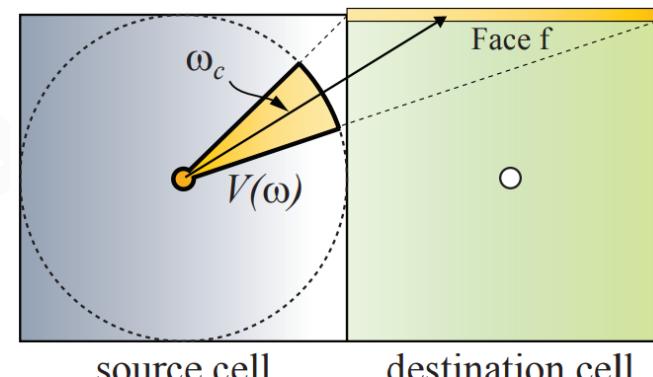
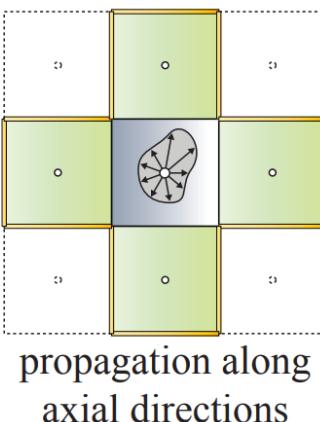
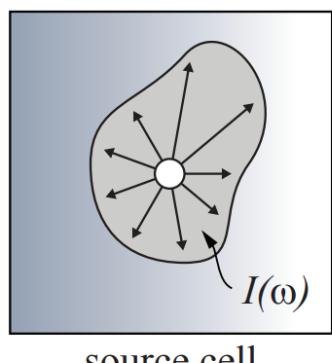
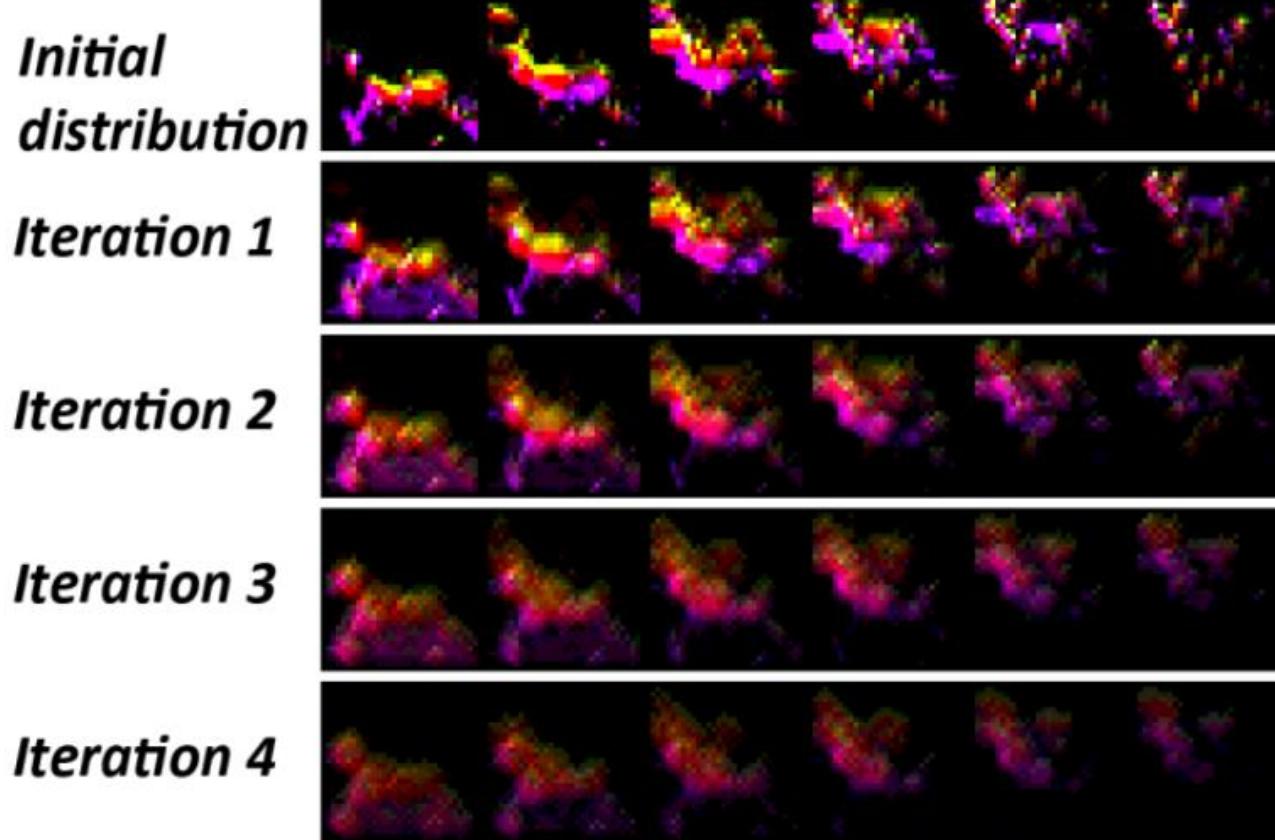


Figure 4. Radiance propagation iteration





## Light with “Limit Speed”?

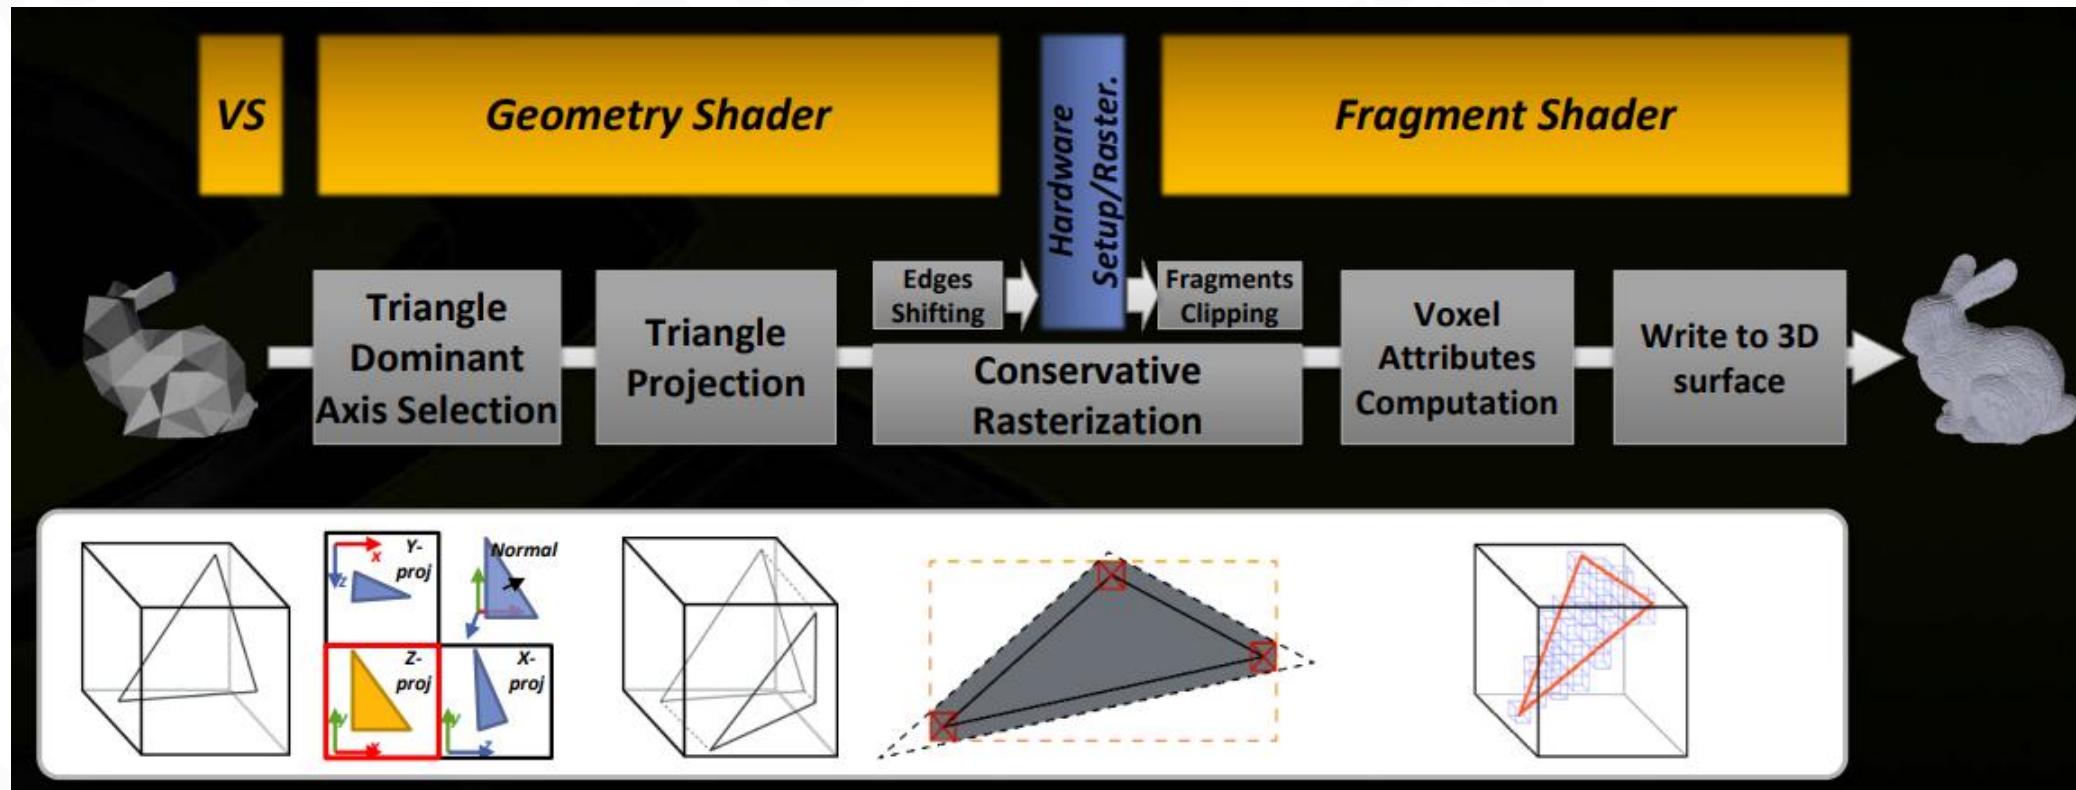




## Sparse Voxel Octree for Real-time Global Illumination (SVOGI)

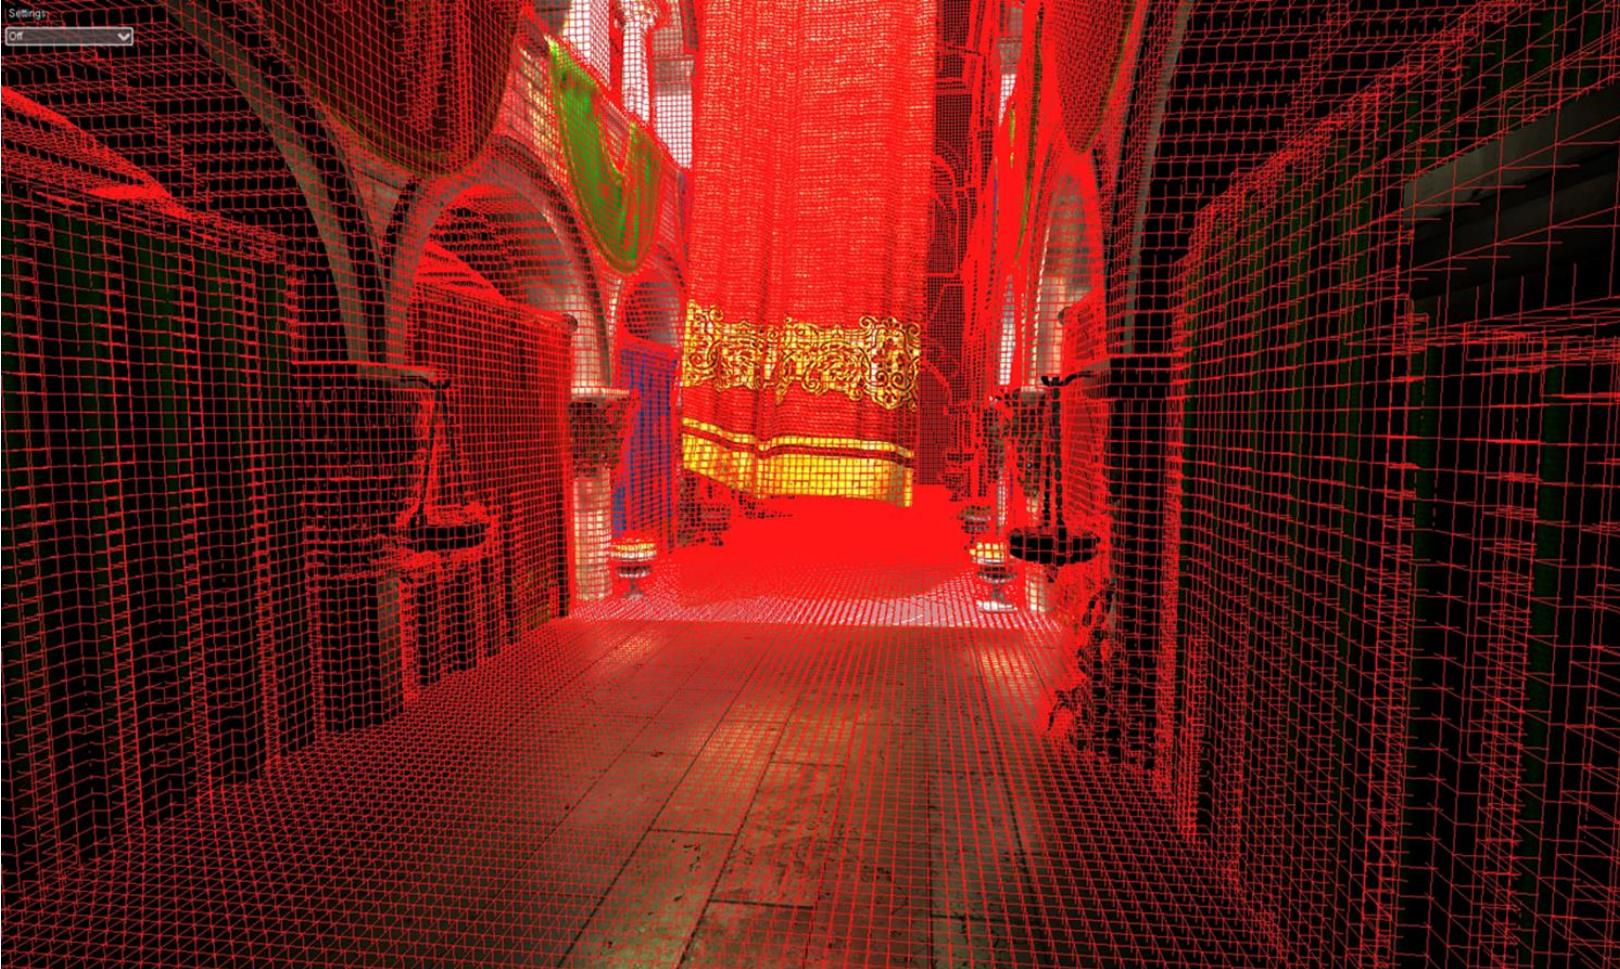


## Voxelization Pass



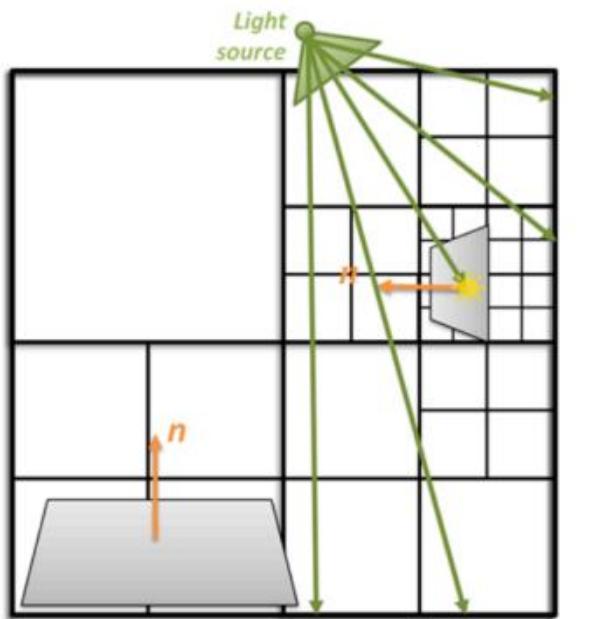


## Collect Surface Voxels

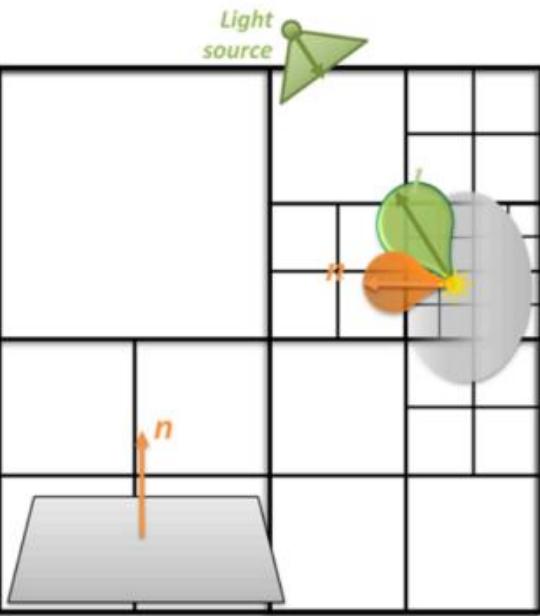




- Inject Irradiance into voxels from light
- Filter irradiance inside the octree

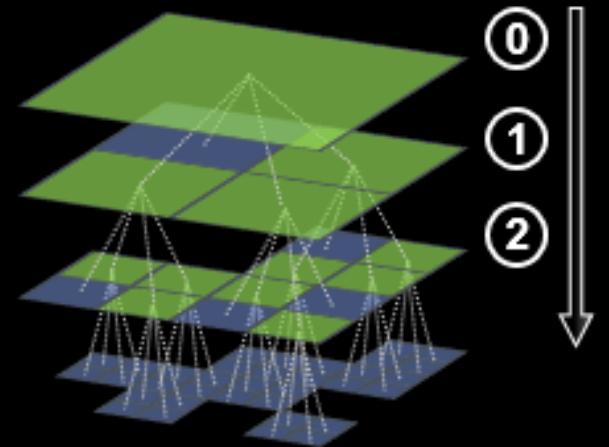


*Step 1: Render from light sources.  
Bake incoming radiance and light  
direction into the octree*

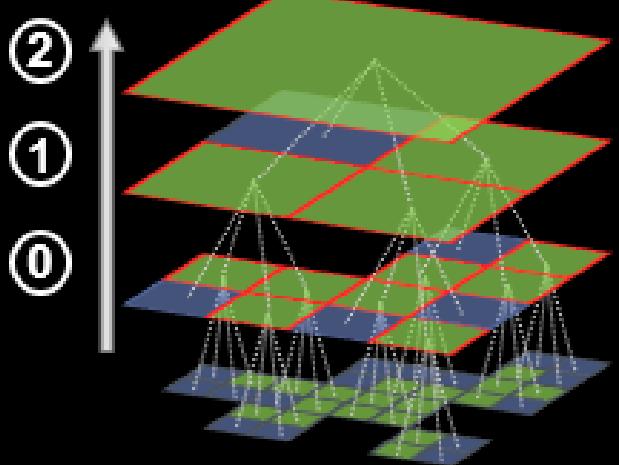


*Step 2: Filter irradiance values and  
light directions inside the octree*

## 1. Octree subdivision



## 2. Values MIP-mapping

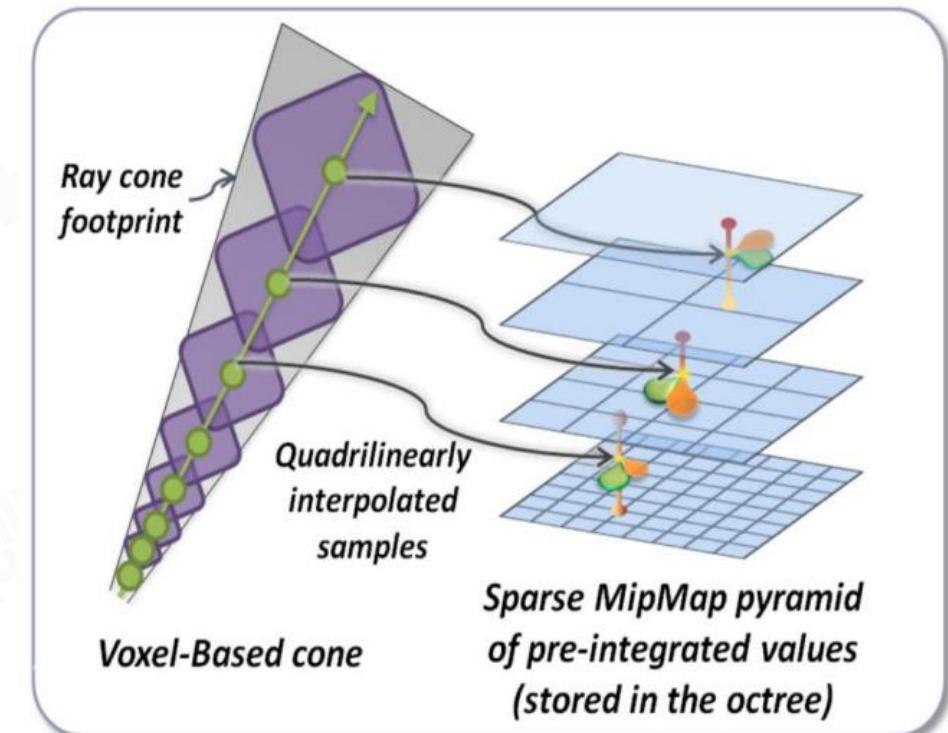
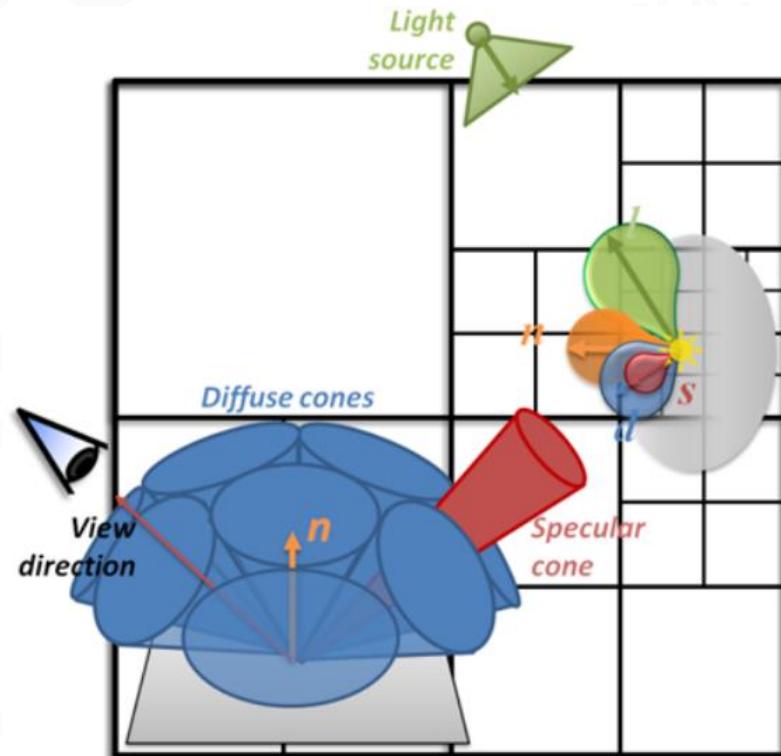




# Shading with Cone Tracing in Voxel Tree

Pass 2 from the camera

- Emit some cones based on diffuse+specular BRDF
- Query in octree based on the (growing) size of the cone

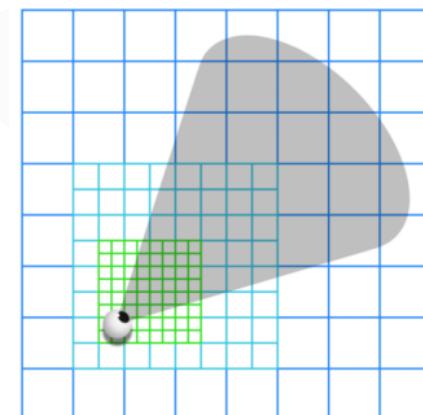
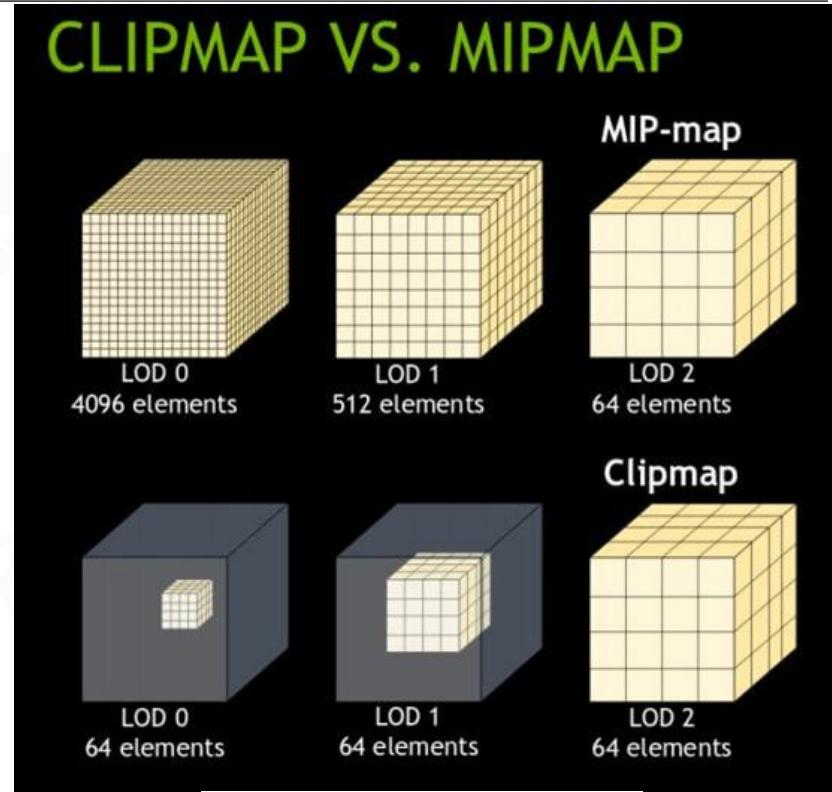




## Voxelization Based Global Illumination (VXGI)



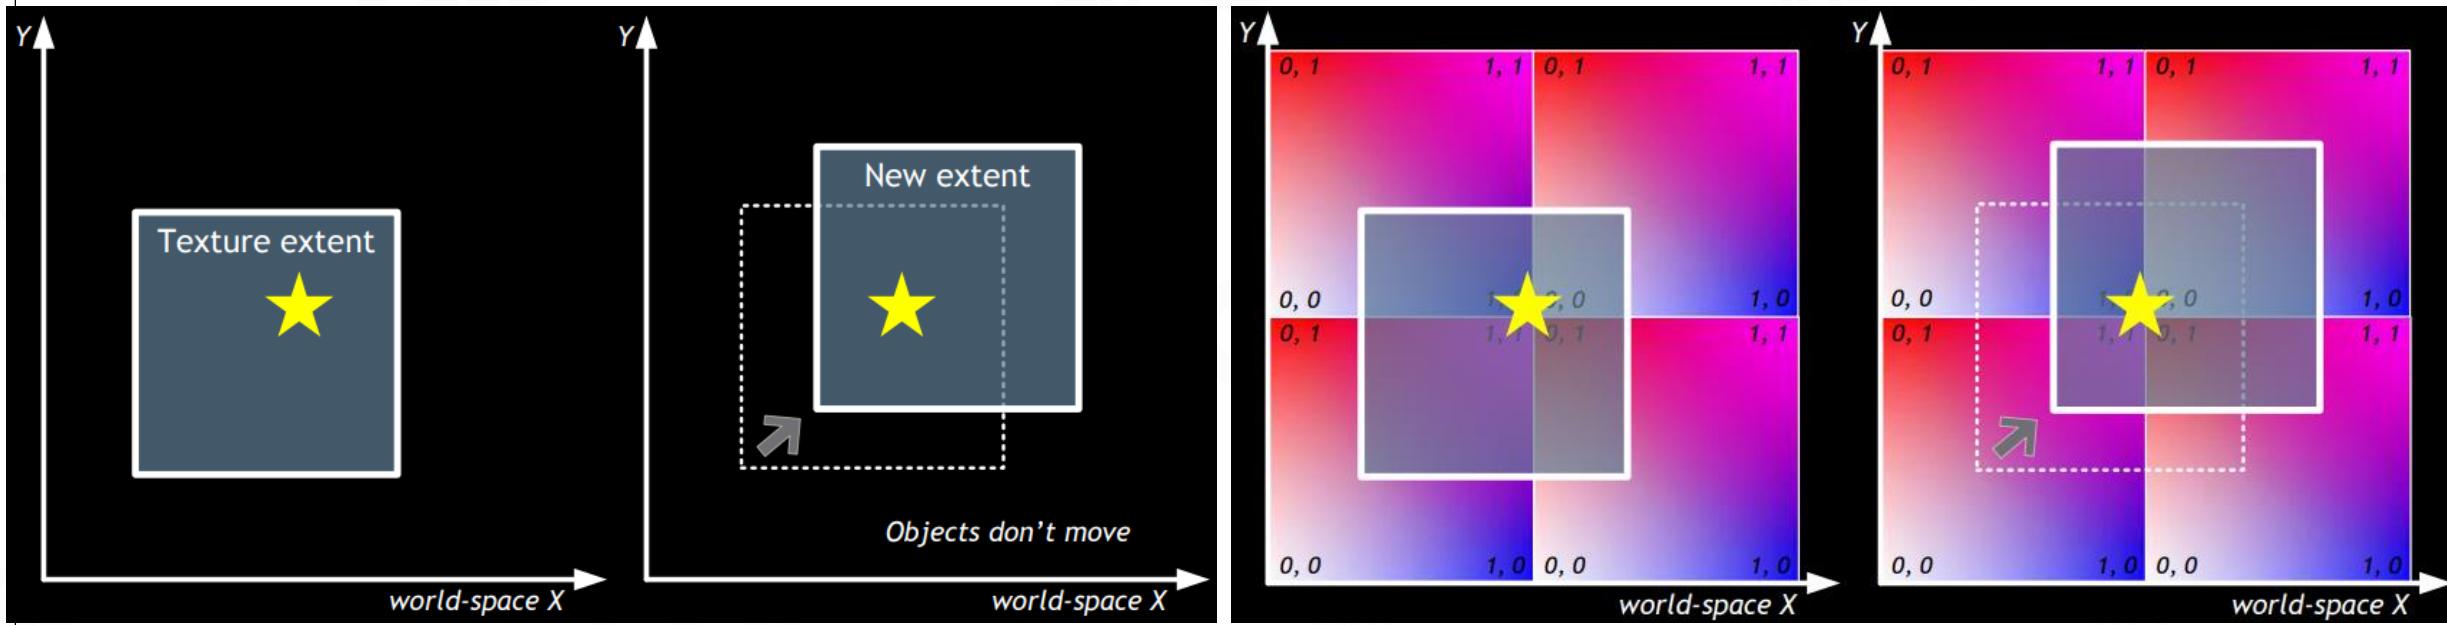
- Store the voxel data in clipmaps
  - Multi-resolution texture
  - Regions near the center have higher spatial resolution
  - Seems to map naturally to cone tracing needs
- A clipmap is easier to build than SVO
  - No nodes, pointers etc., handled by hardware
- A clipmap is easier to read from
- Clipmap size is  $(64\dots256)^3$  with 3\dots5 levels of detail
  - 16\dots32 bytes per voxel => 12 MB \dots 2.5 GB of video memory required

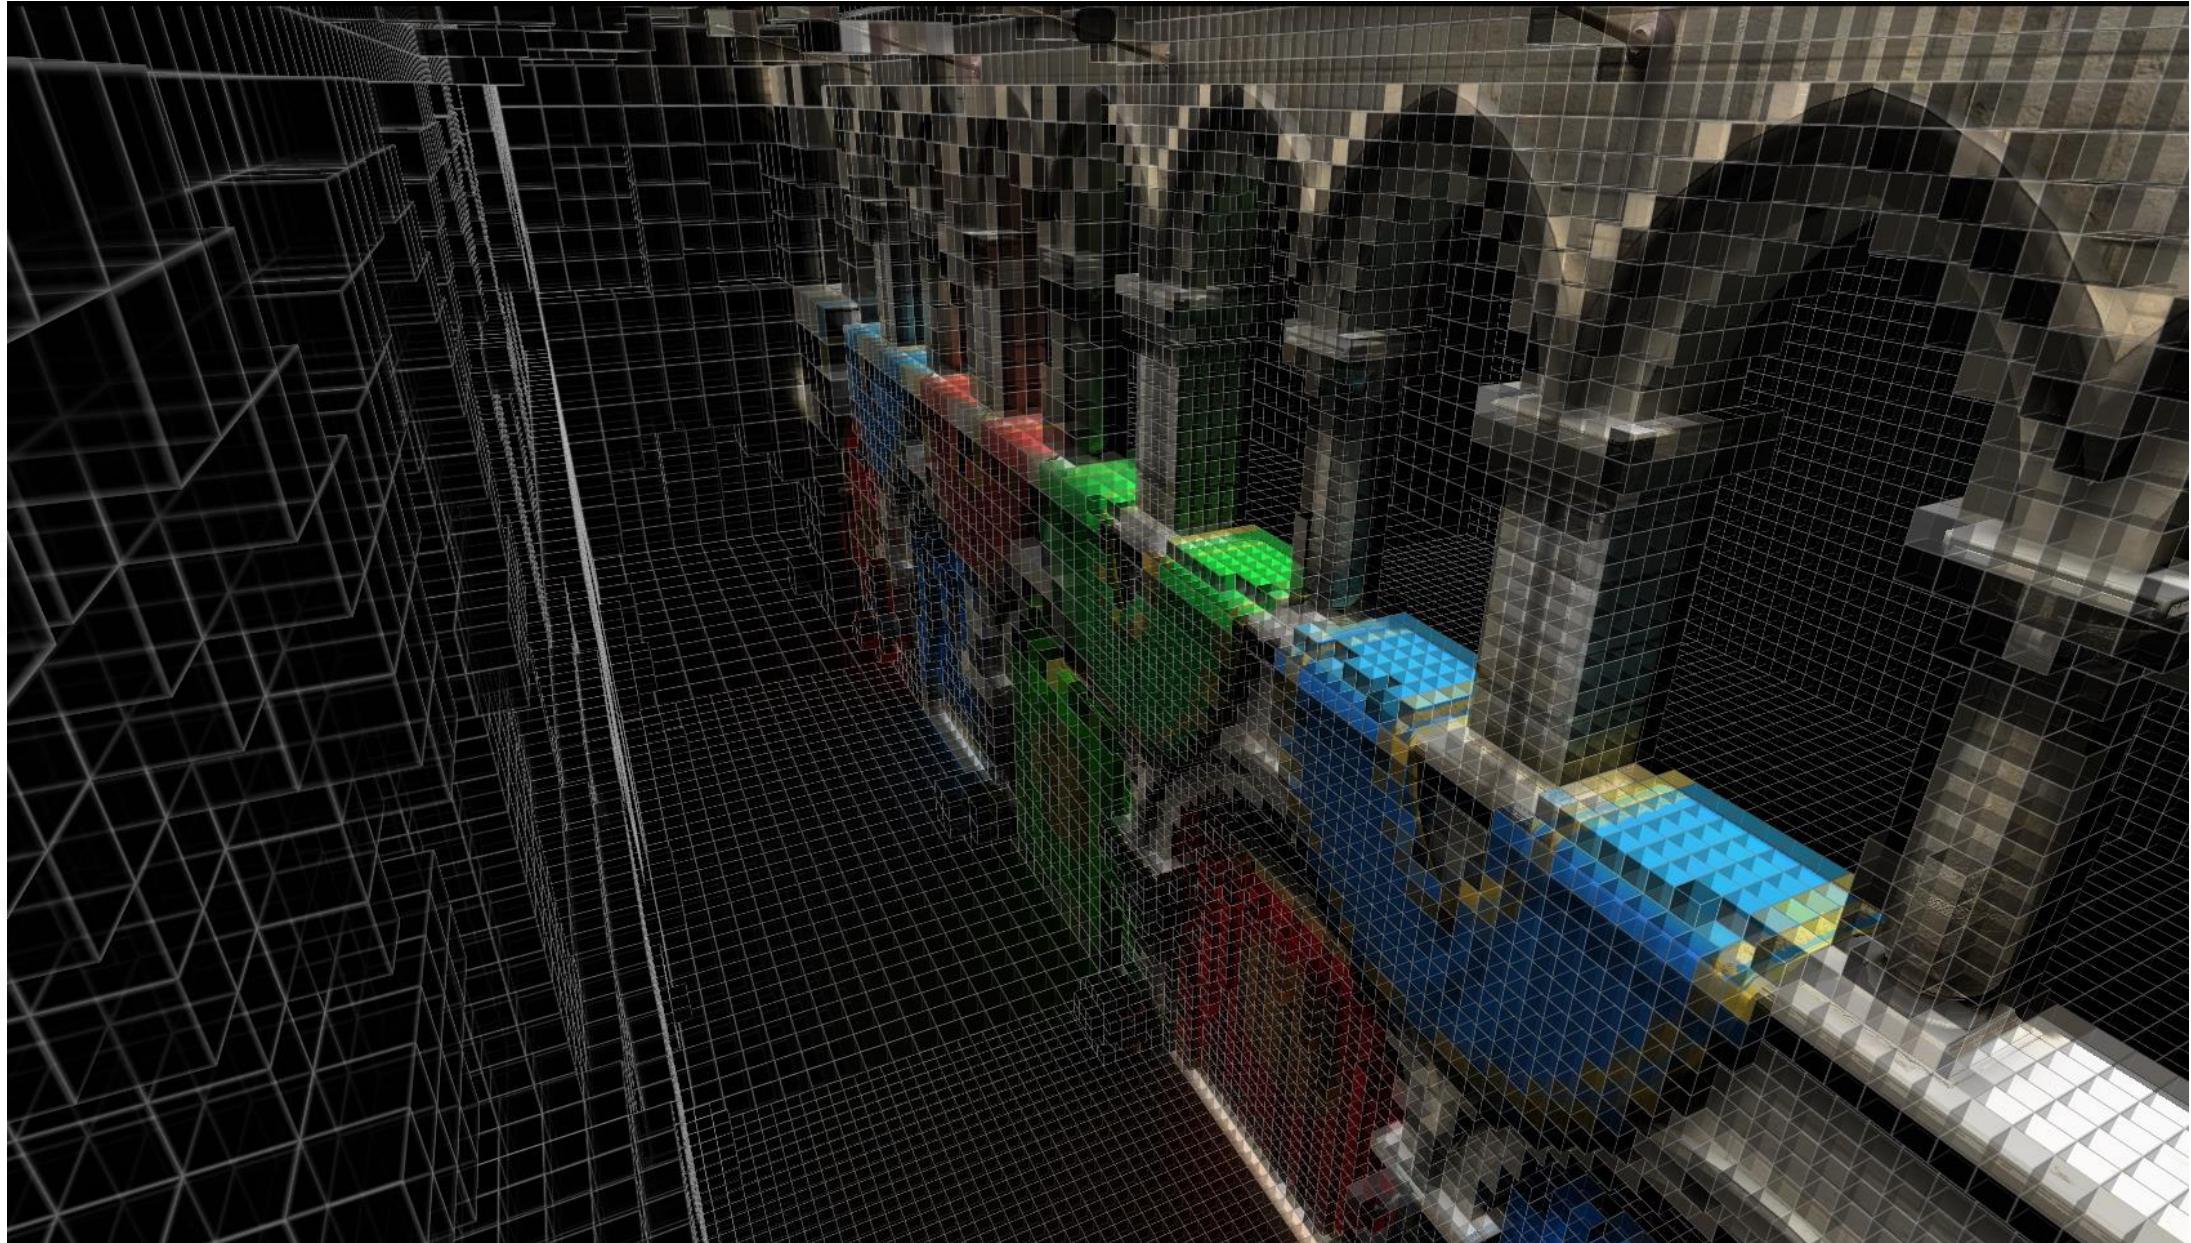




# Voxel Update and Toroidal Addressing

- A fixed point in space always maps to the same address in the clipmap
- The background shows texture addresses:  $\text{frac}(\text{worldPos.xy} / \text{clipmapSize.xy})$

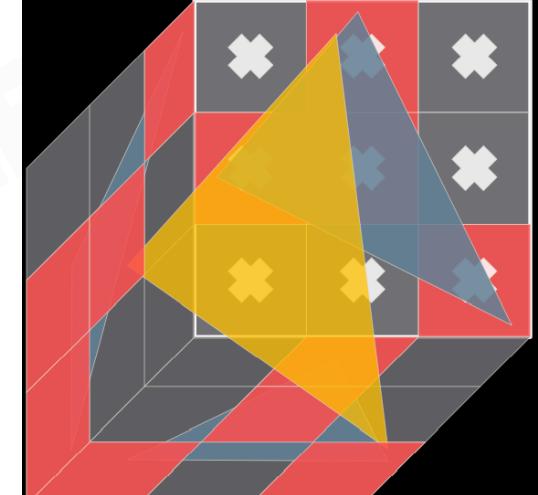
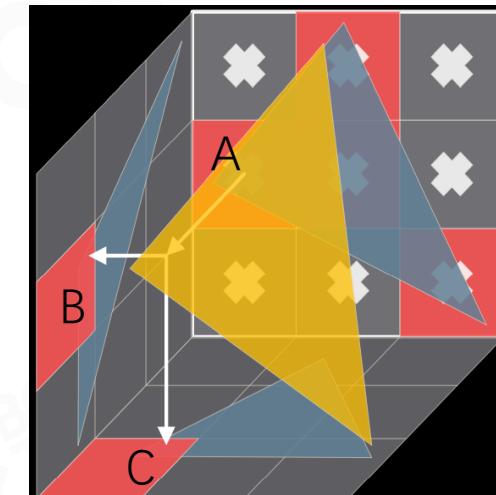
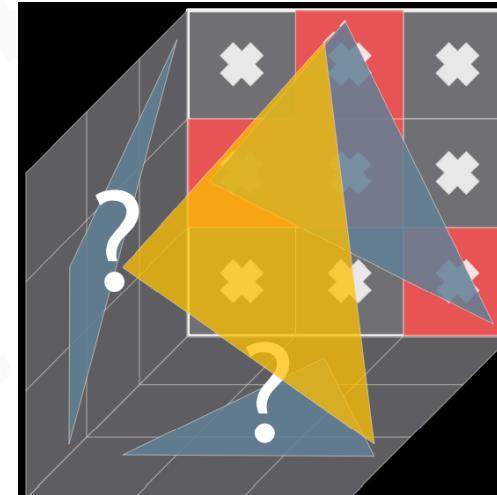
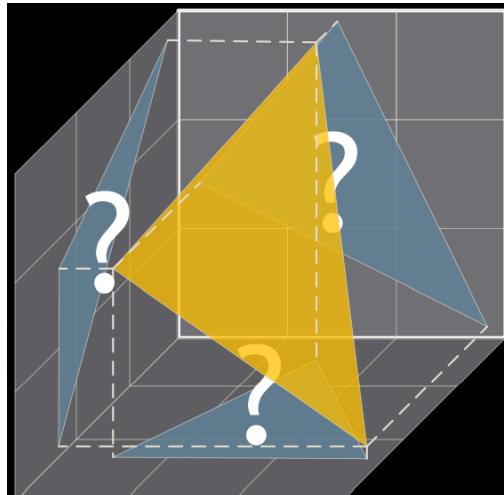






## Voxelization for Opacity

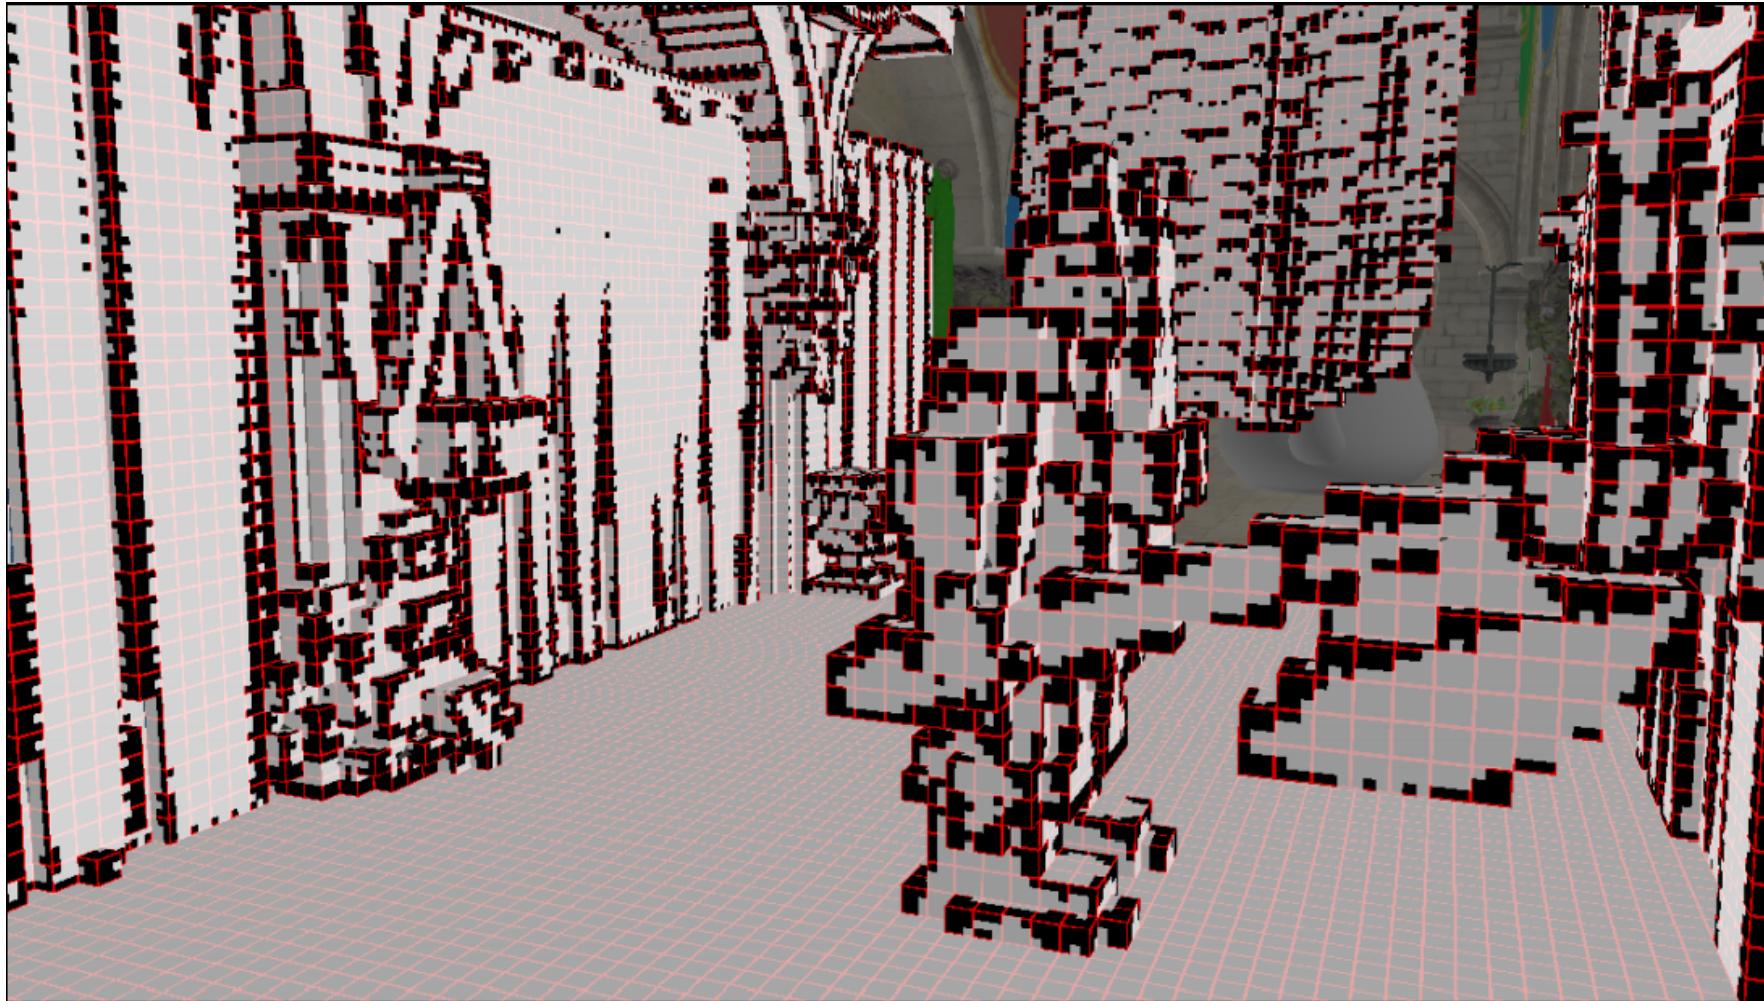
- We have a triangle and a voxel
- Select the projection plane that yields the biggest projection area
- Rasterize the triangle using MSAA to compute one coverage mask per pixel
- Take the MSAA samples and reproject them onto other planes
- Repeat that process for all covered samples
- Thicken the result by blurring all the reprojected samples



Opacity = (number of the covered MSAA samples) / MSAA\_Resolution<sup>2</sup>



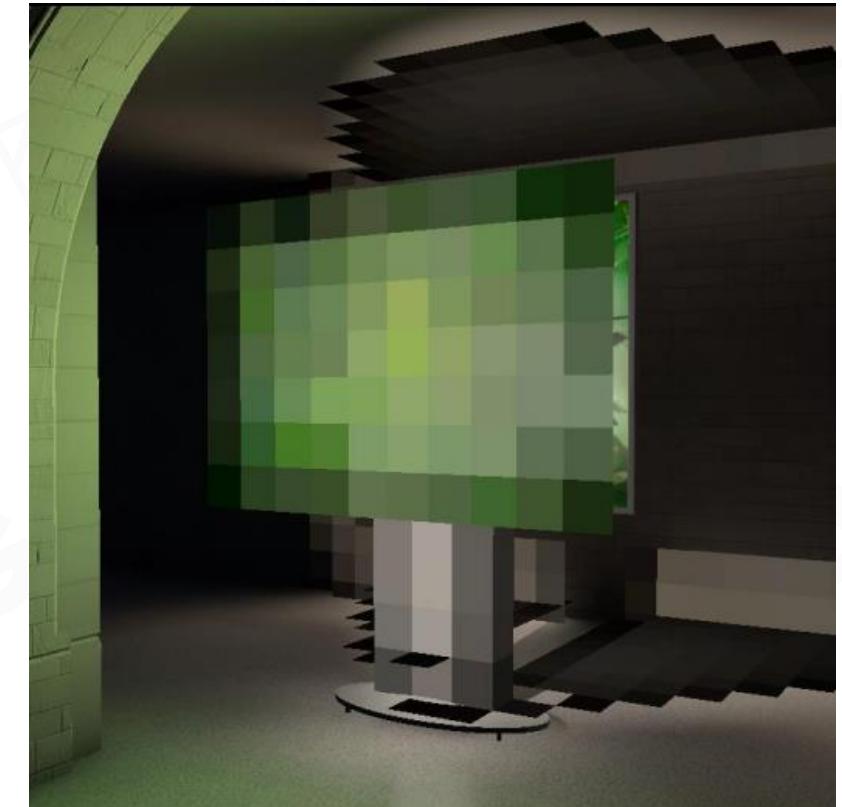
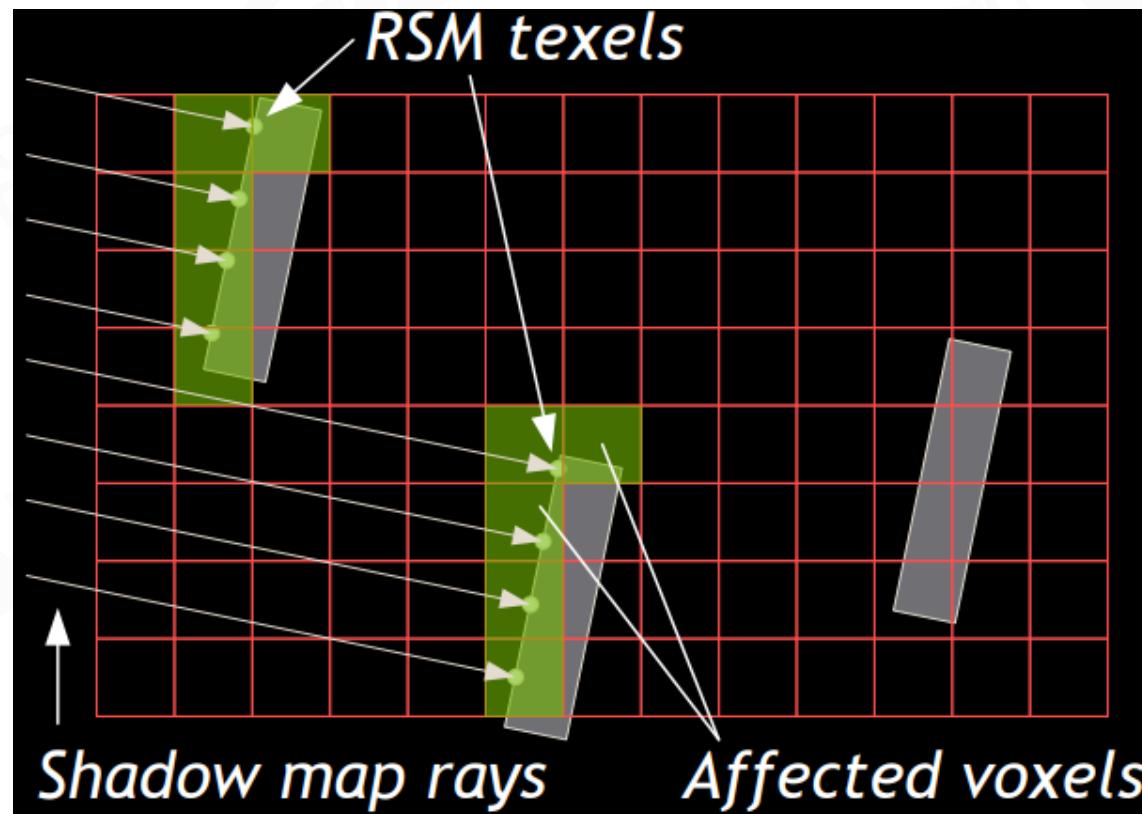
## Voxelization: Directional Coverage

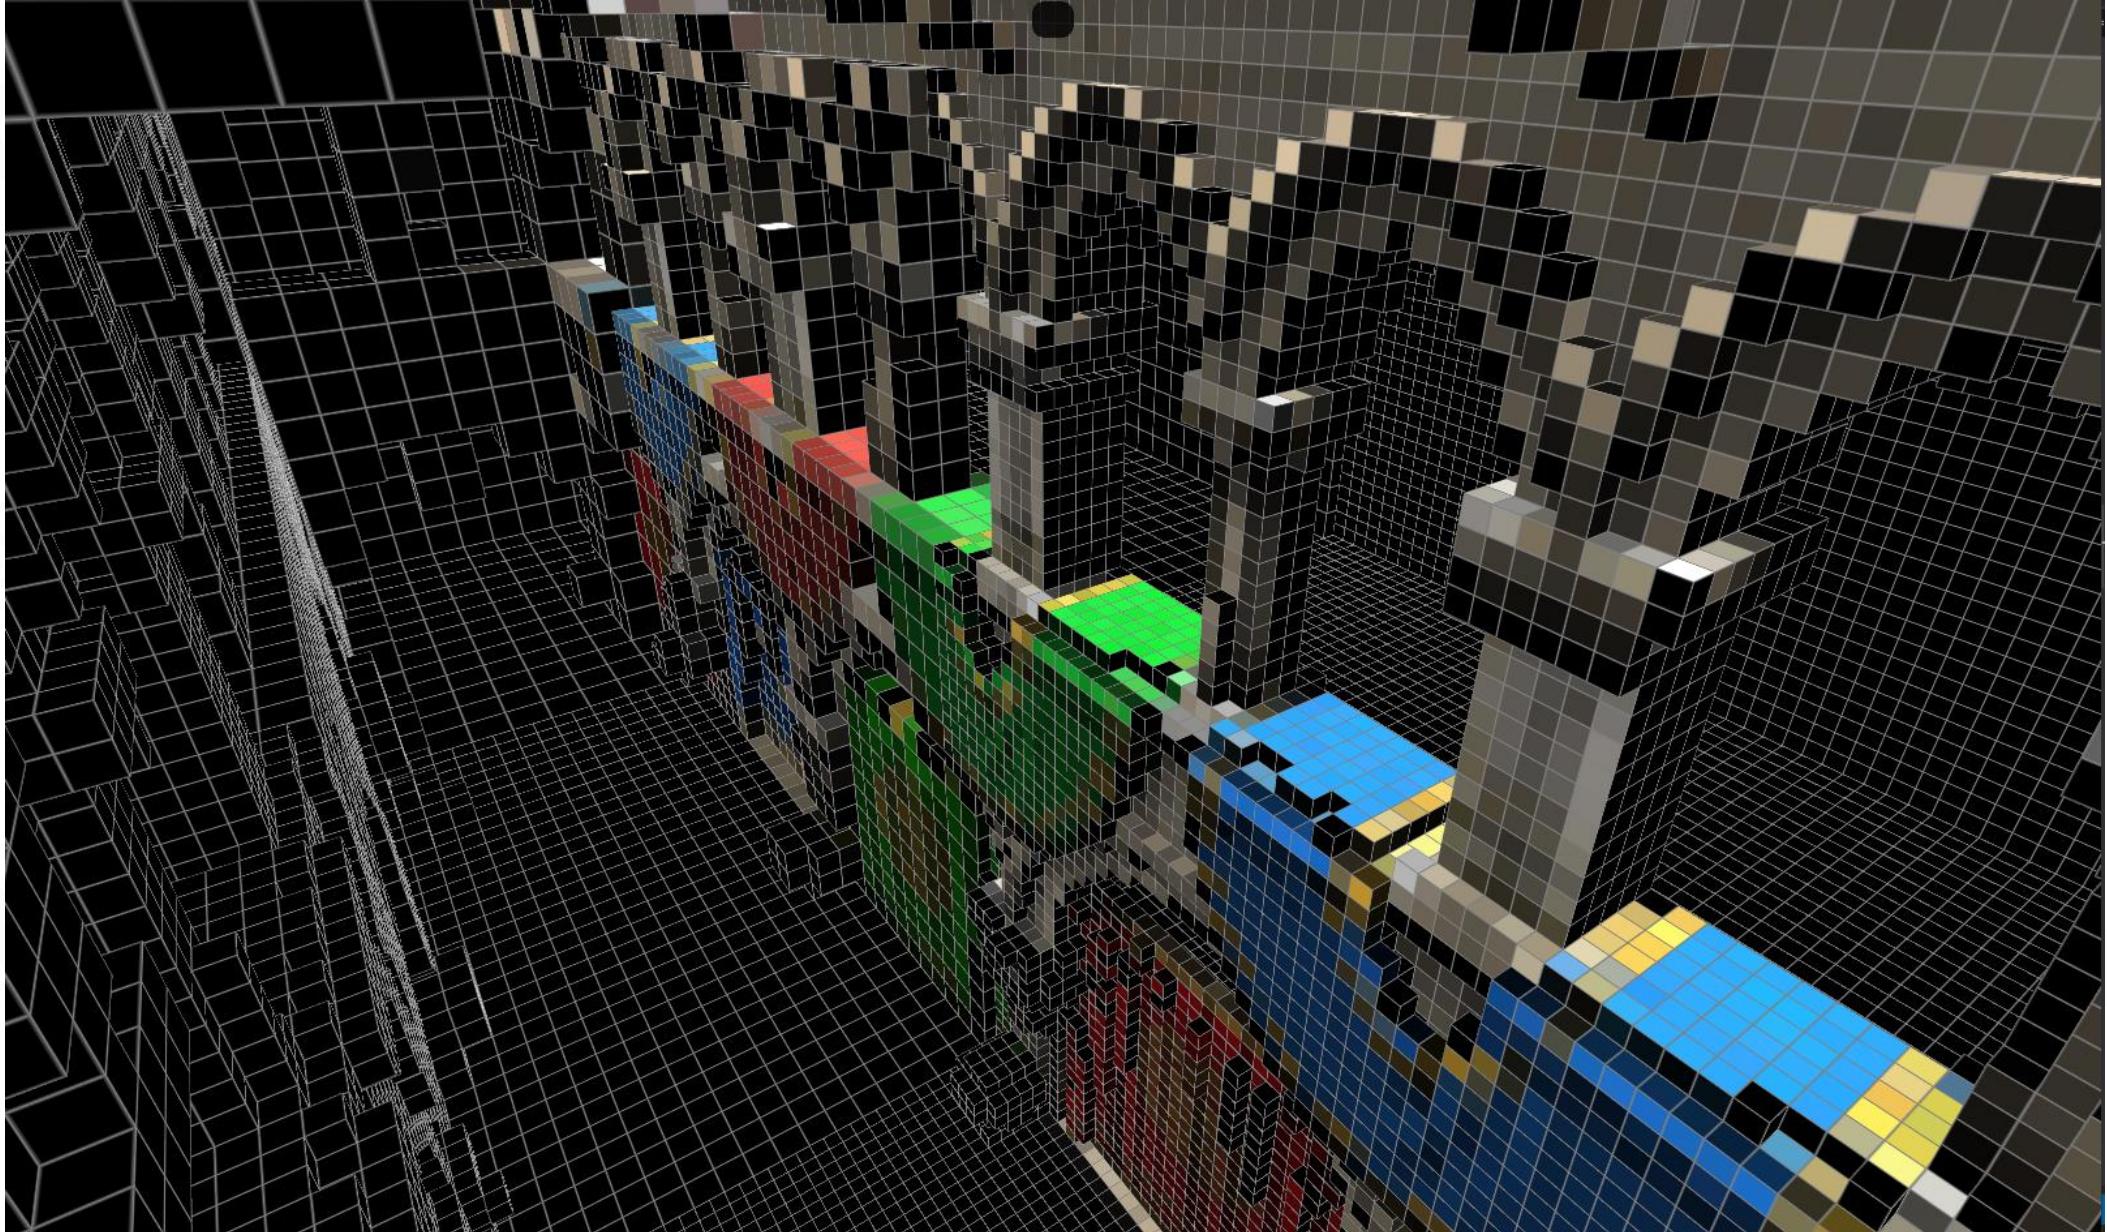




## Light Injection

- Calculate emittance of voxels that contain surfaces lit by direct lights
- Take information from reflective shadow maps (RSM)



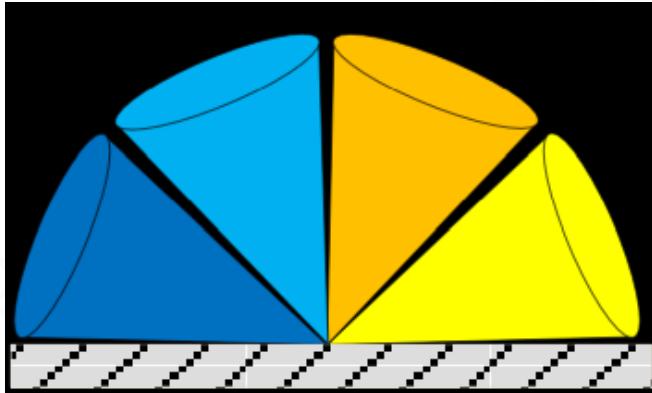




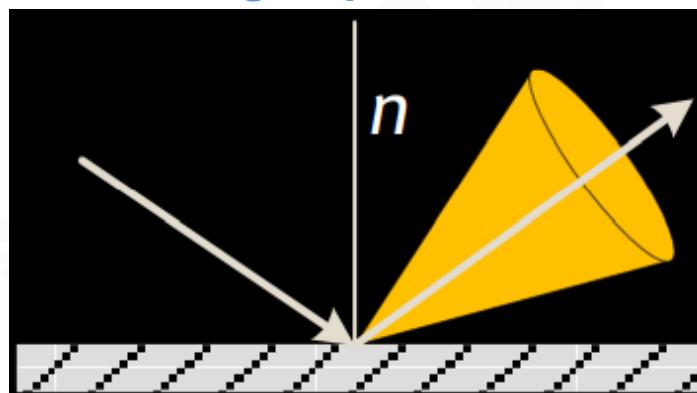
## Shading with Cone Tracing

- generate several cones based on BRDF

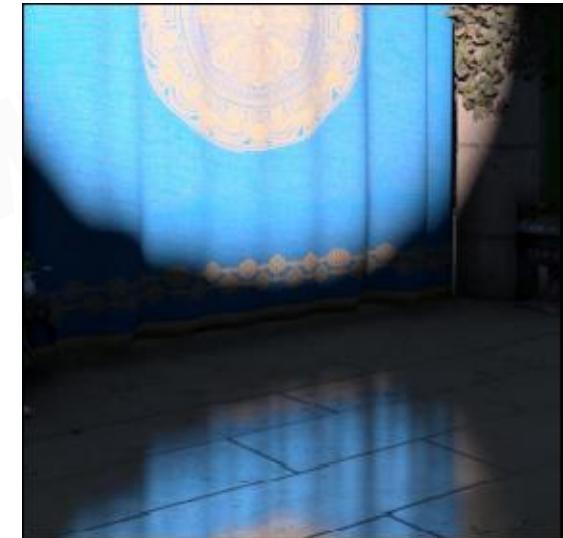
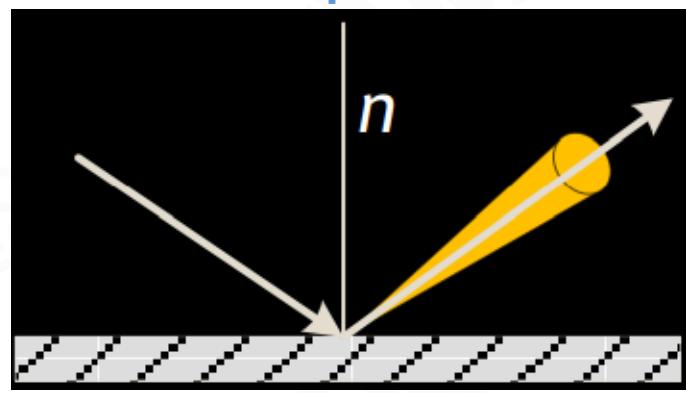
Diffuse



Rough Specular



Fine Specular





## Accumulate Voxel Radiance and Opacity along the Path

$$C_{\text{dst}} \leftarrow C_{\text{dst}} + (1 - \alpha_{\text{dst}})C_{\text{src}}$$
$$\alpha_{\text{dst}} \leftarrow \alpha_{\text{dst}} + (1 - \alpha_{\text{dst}})\alpha_{\text{src}}$$







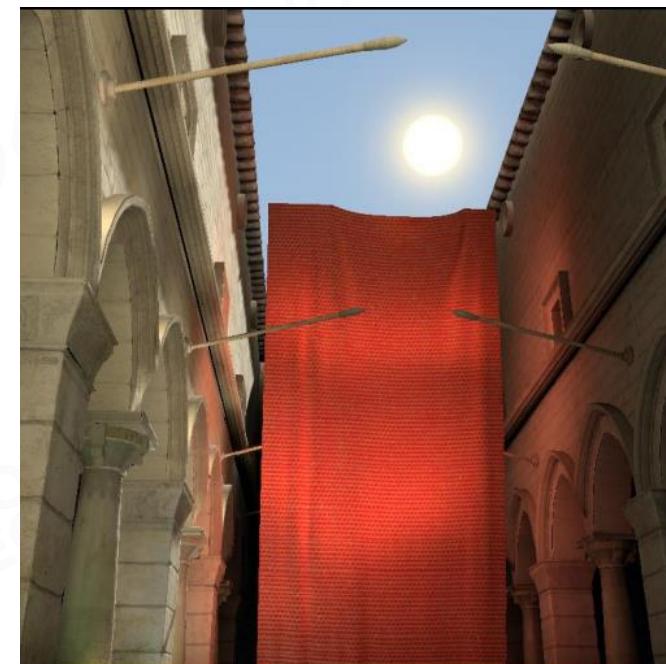
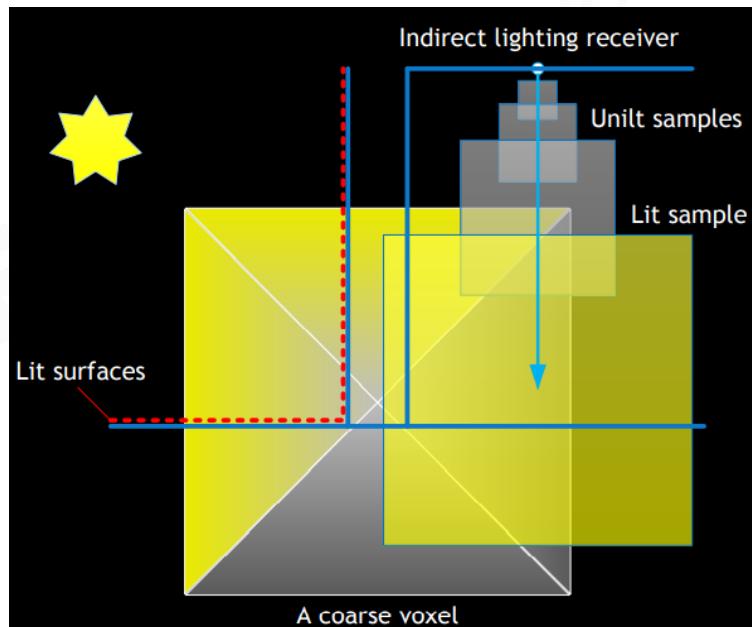
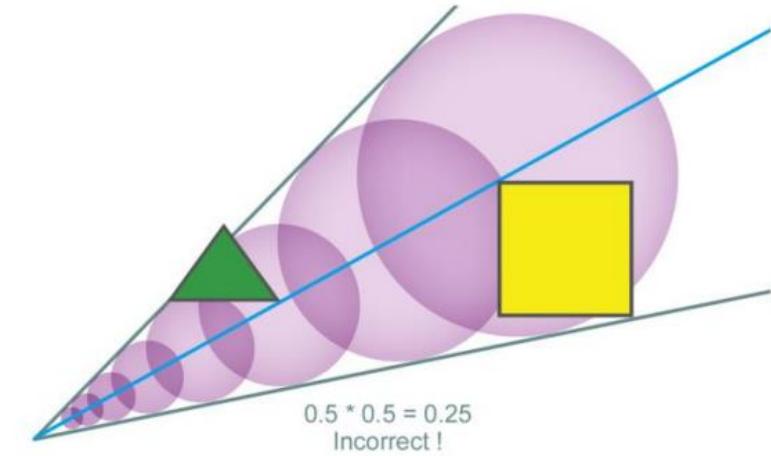
## Problems in VXGI

### Incorrect Occlusion(opacity)

- naively combine the opacity with alpha blending.

### Light Leaking

- when occlusion wall is much smaller than voxel size





## Screen Space Global Illumination (SSGI)



# FROSTBITE™

empowers game creators to shape the future of gaming  
SIGGRAPH 2015: Advances in Real-Time Rendering course





## General Idea

- Reuse screen-space data

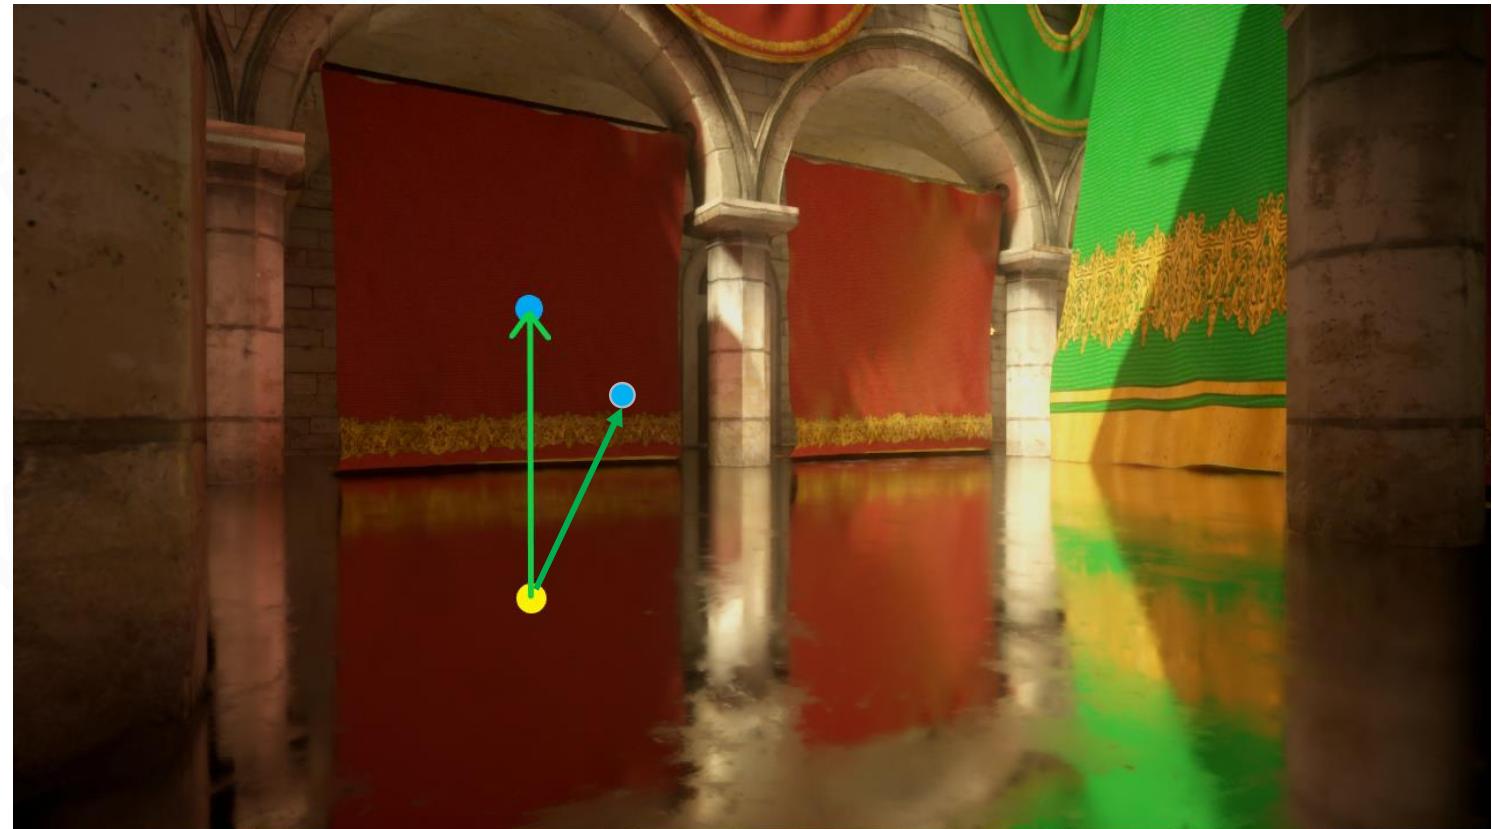




## Radiance Sampling in Screen Space

For each fragment:

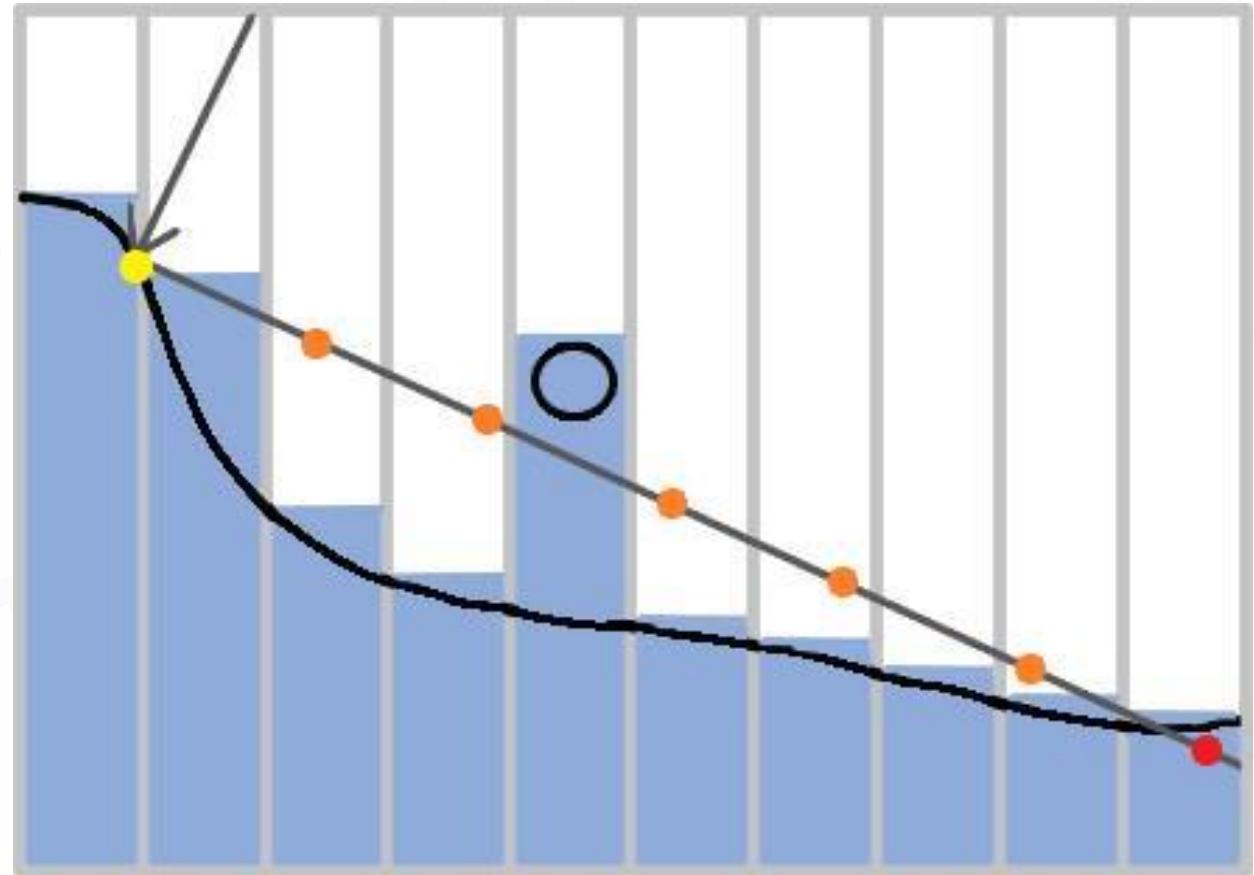
- **Step 1:** compute many reflection rays
- **Step 2:** march along ray direction (in depth gbuffer)
- **Step 3:** use color of hit point as indirect lighting





## Linear Raymarching

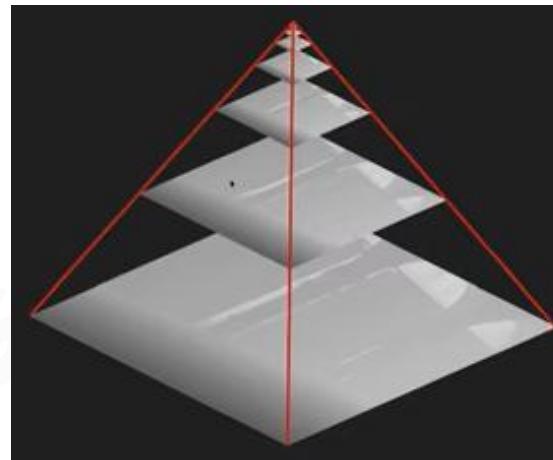
- **General Steps**
  - Step forward at a fixed step size
  - At each step, check depth value
- **Features**
  - Fast
  - May skip thin objects



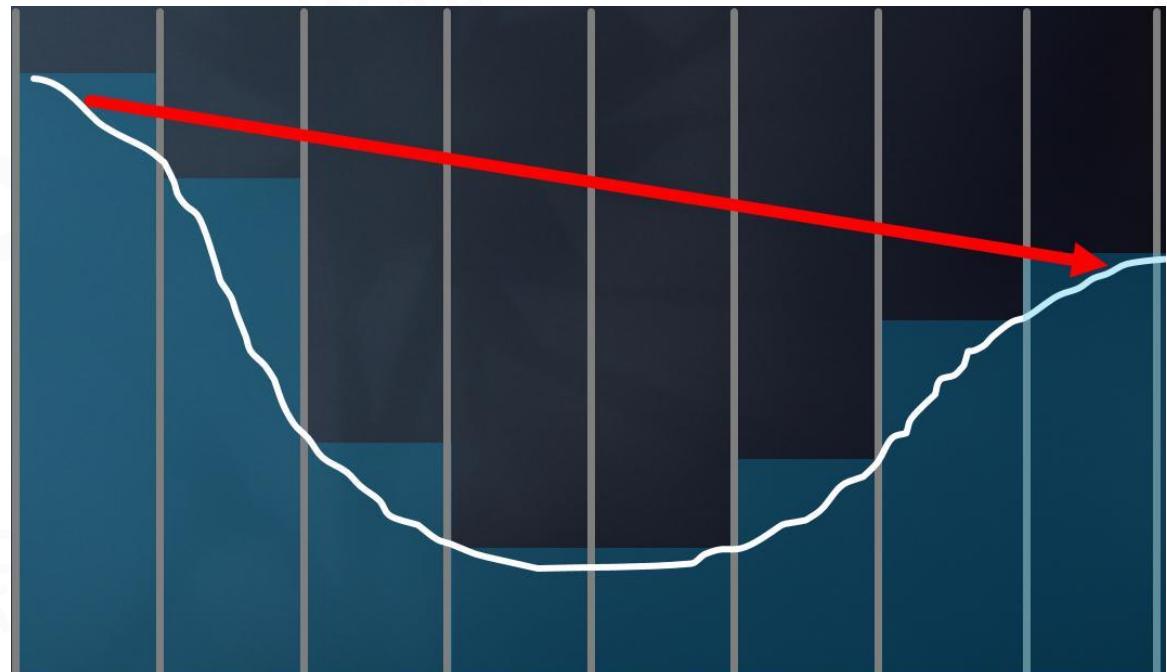


## Hierachical Tracing

- Generate min-depth mipmap (pyramid)
- Stackless ray walk of min-depth mipmap



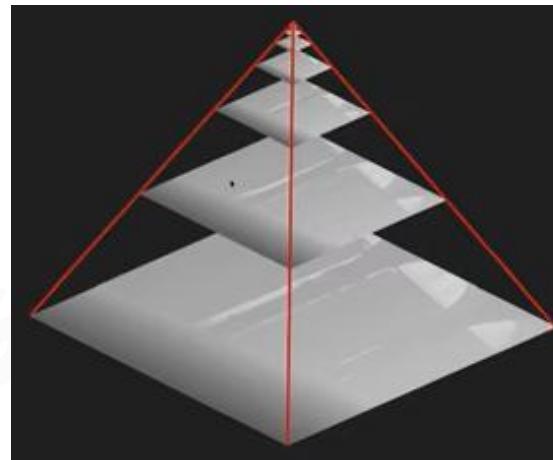
```
level = 0;  
while (level > -1)  
{  
    stepCurrentCell();  
    if (above Z plane) level++;  
    if (below Z plane) level--;  
}
```



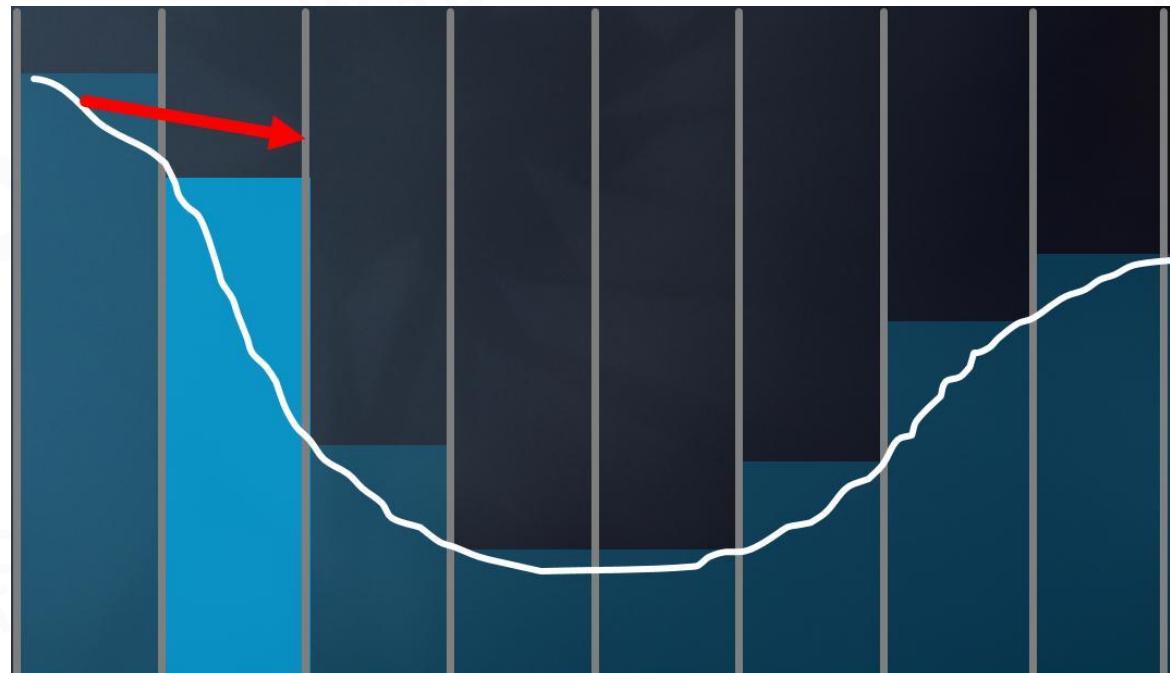


## Hierachical Tracing

- Generate min-depth mipmap (pyramid)
- Stackless ray walk of min-depth mipmap



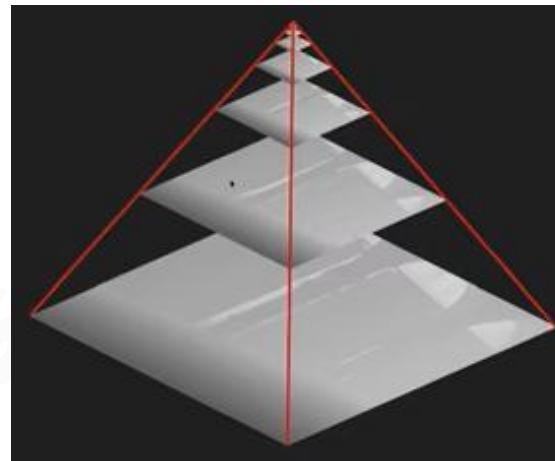
```
level = 0;  
while (level > -1)  
{  
    stepCurrentCell();  
    if (above Z plane) level++;  
    if (below Z plane) level--;  
}
```



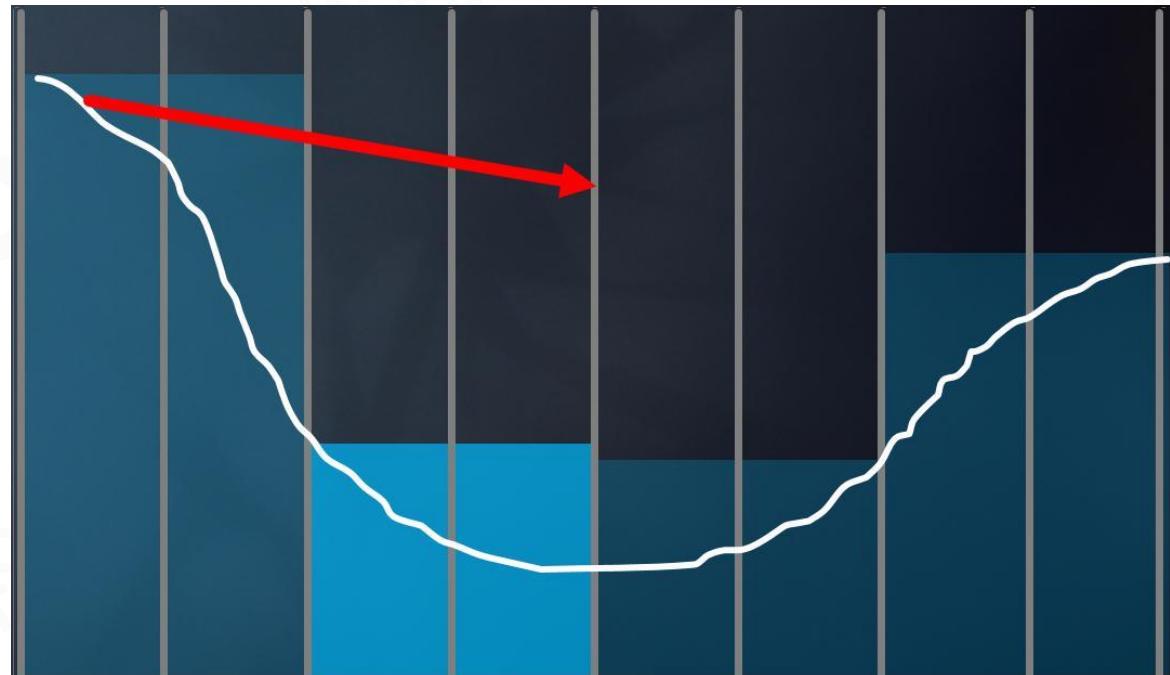


## Hierachical Tracing

- Generate min-depth mipmap (pyramid)
- Stackless ray walk of min-depth mipmap



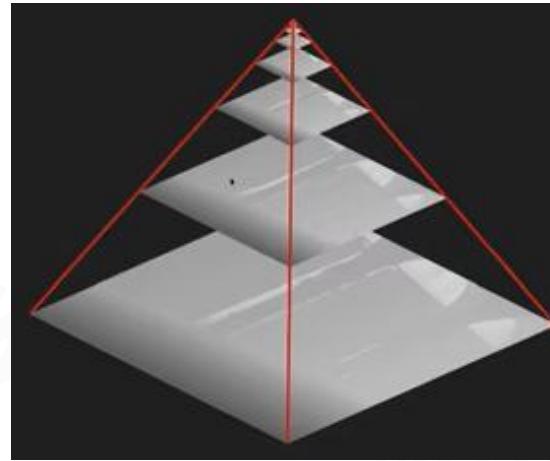
```
level = 0;  
while (level > -1)  
{  
    stepCurrentCell();  
    if (above Z plane) level++;  
    if (below Z plane) level--;  
}
```



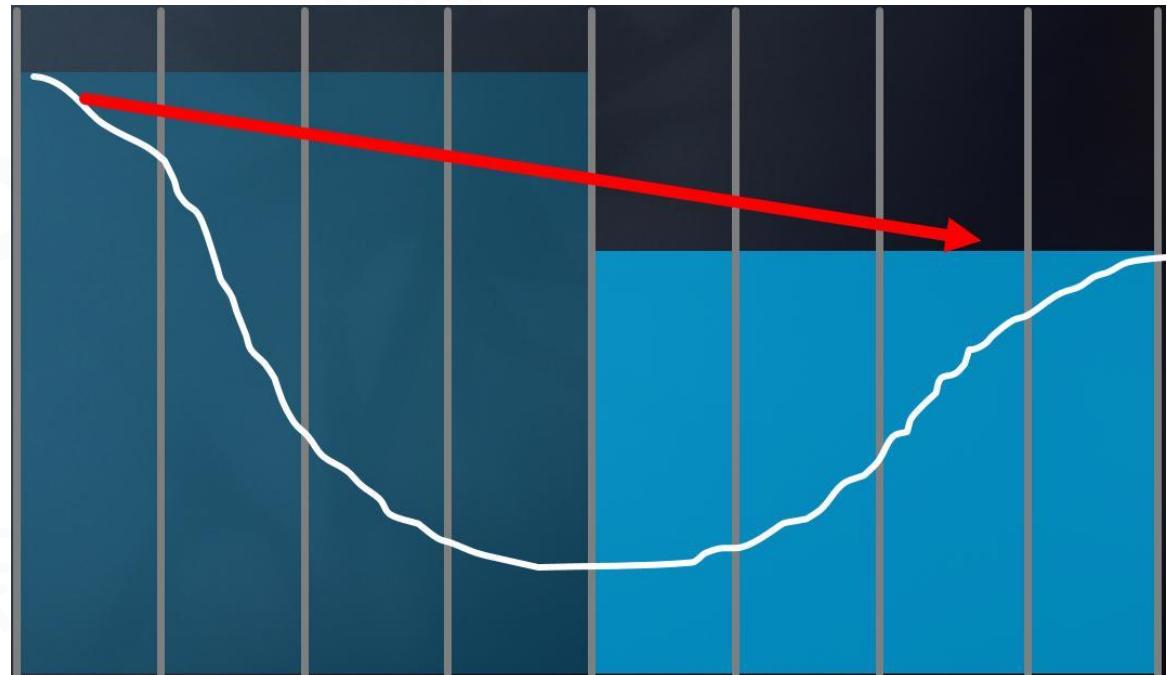


## Hierachical Tracing

- Generate min-depth mipmap (pyramid)
- Stackless ray walk of min-depth mipmap



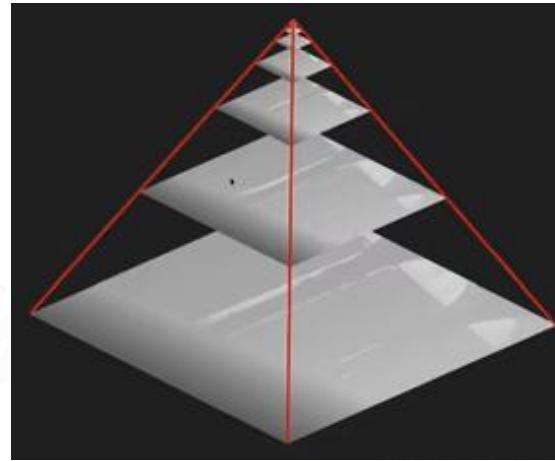
```
level = 0;  
while (level > -1)  
{  
    stepCurrentCell();  
    if (above Z plane) level++;  
    if (below Z plane) level--;  
}
```



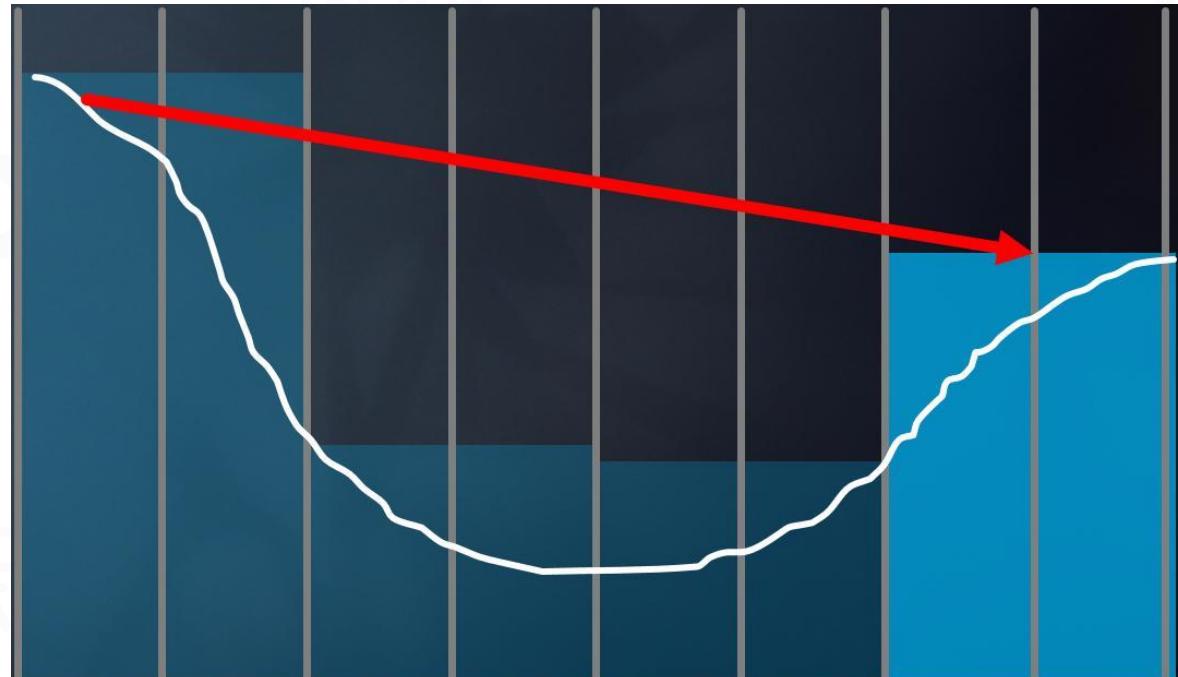


## Hierachical Tracing

- Generate min-depth mipmap (pyramid)
- Stackless ray walk of min-depth mipmap



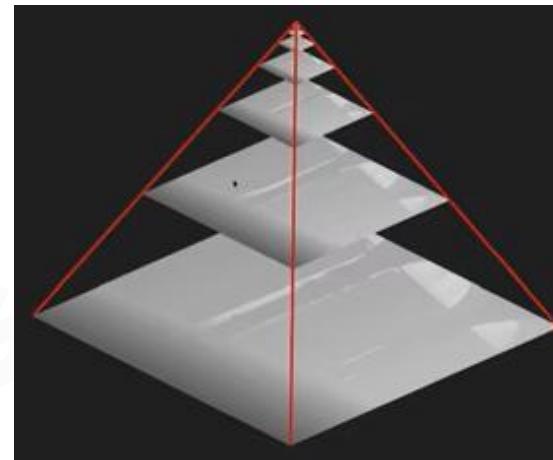
```
level = 0;  
while (level > -1)  
{  
    stepCurrentCell();  
    if (above Z plane) level++;  
    if (below Z plane) level--;  
}
```



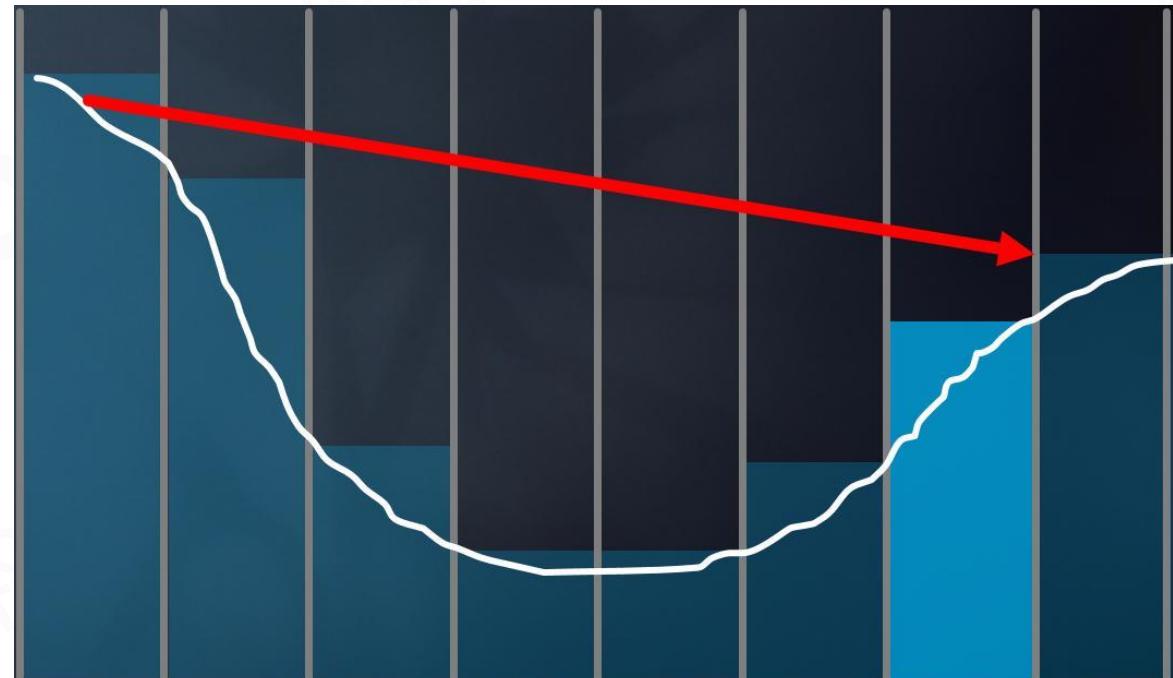


## Hierachical Tracing

- Generate min-depth mipmap (pyramid)
- Stackless ray walk of min-depth mipmap



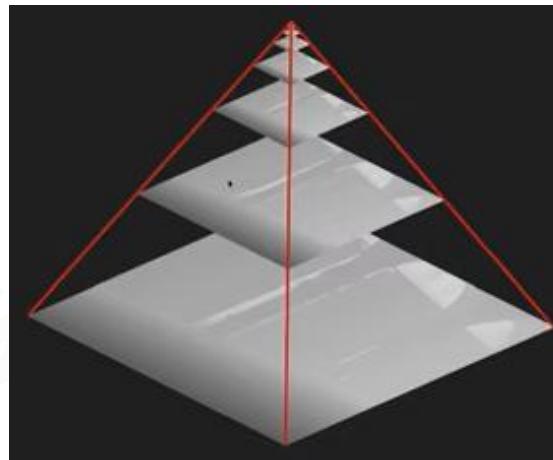
```
level = 0;  
while (level > -1)  
{  
    stepCurrentCell();  
    if (above Z plane) level++;  
    if (below Z plane) level--;  
}
```



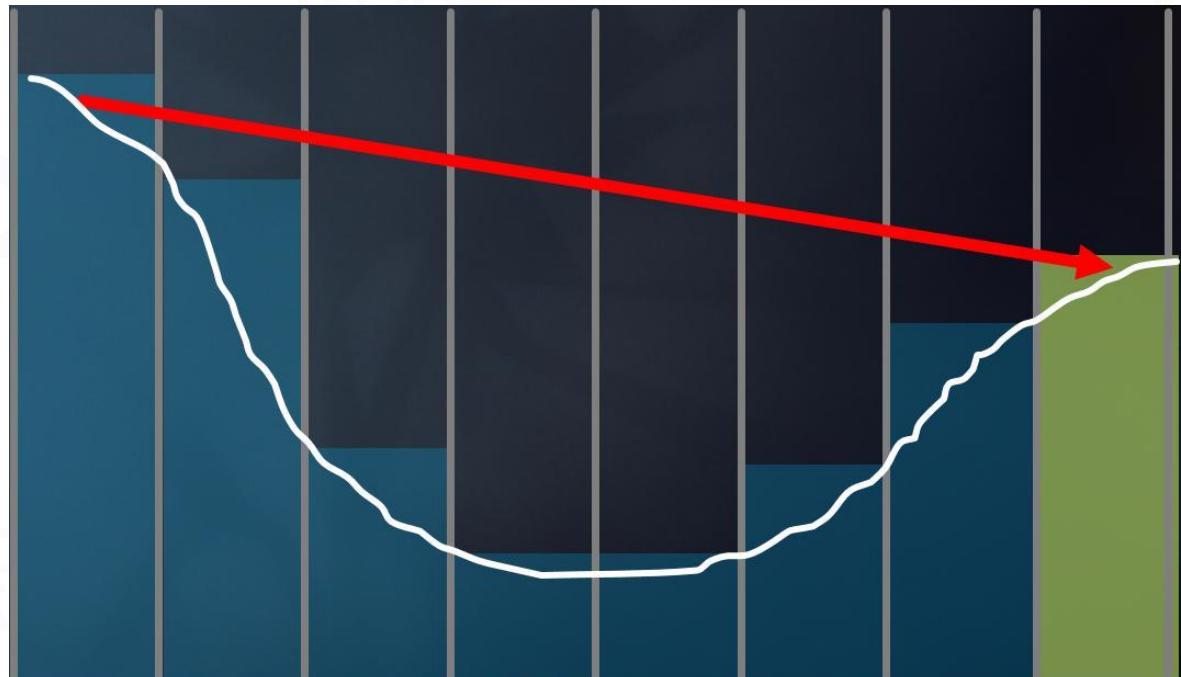


## Hierachical Tracing

- Generate min-depth mipmap (pyramid)
- Stackless ray walk of min-depth mipmap



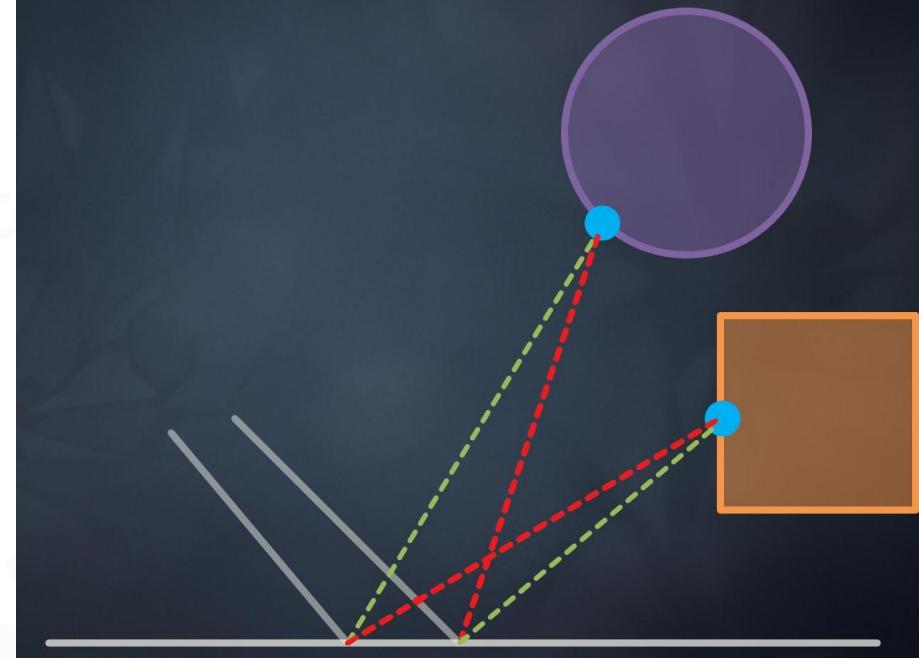
```
level = 0;  
while (level > -1)  
{  
    stepCurrentCell();  
    if (above Z plane) level++;  
    if (below Z plane) level--;  
}
```





## Ray Reuse among Neighbor Pixels

- Store **hitpoint data**
- Assume visibility is the same between neighbors
- Regard **ray to neighbor's hitpoint** as valid





## Cone Tracing with Mipmap Filtering

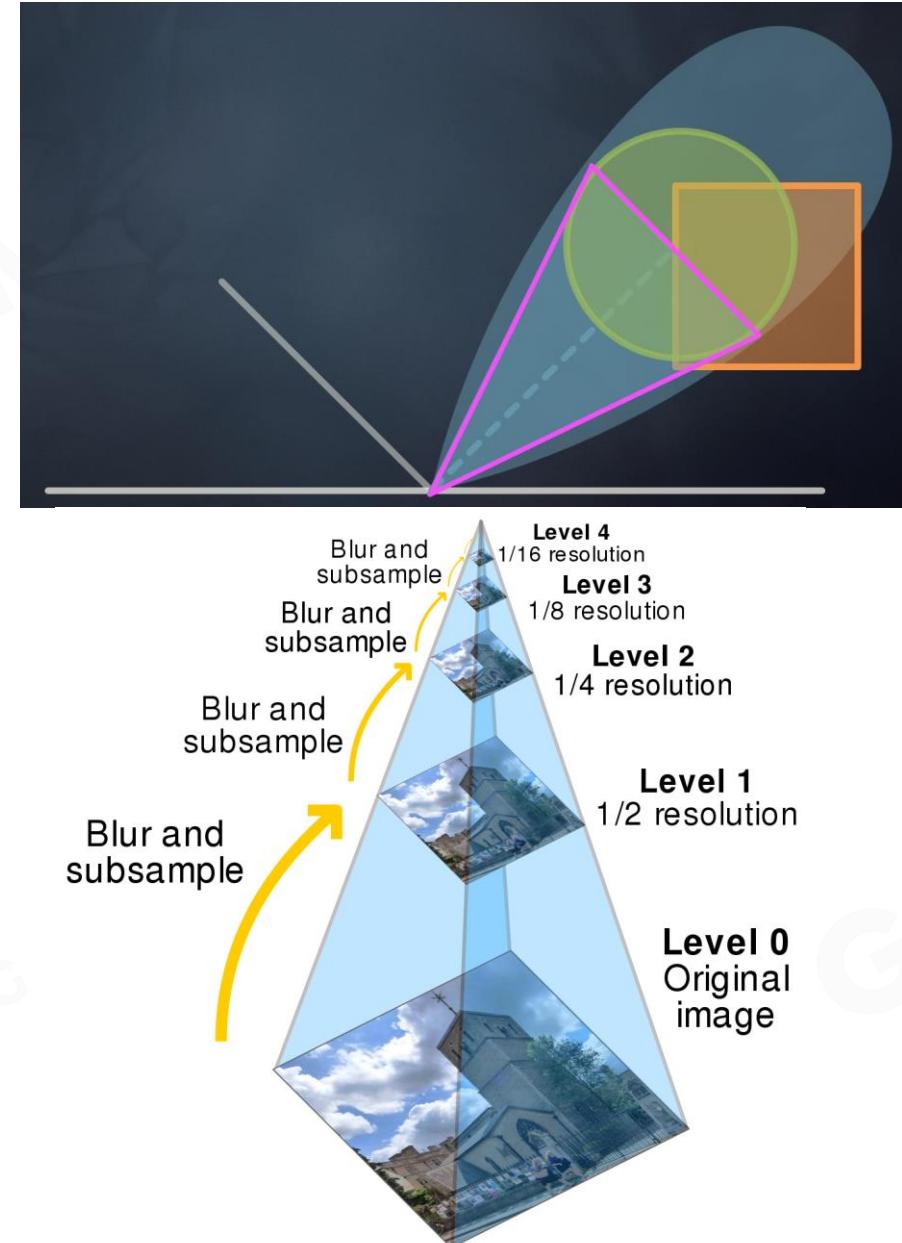
Estimate **footprint** of a **cone** at hit point

- roughness
- distance to hit

Sample the color mipmap

- mip level is determined by **footprint**

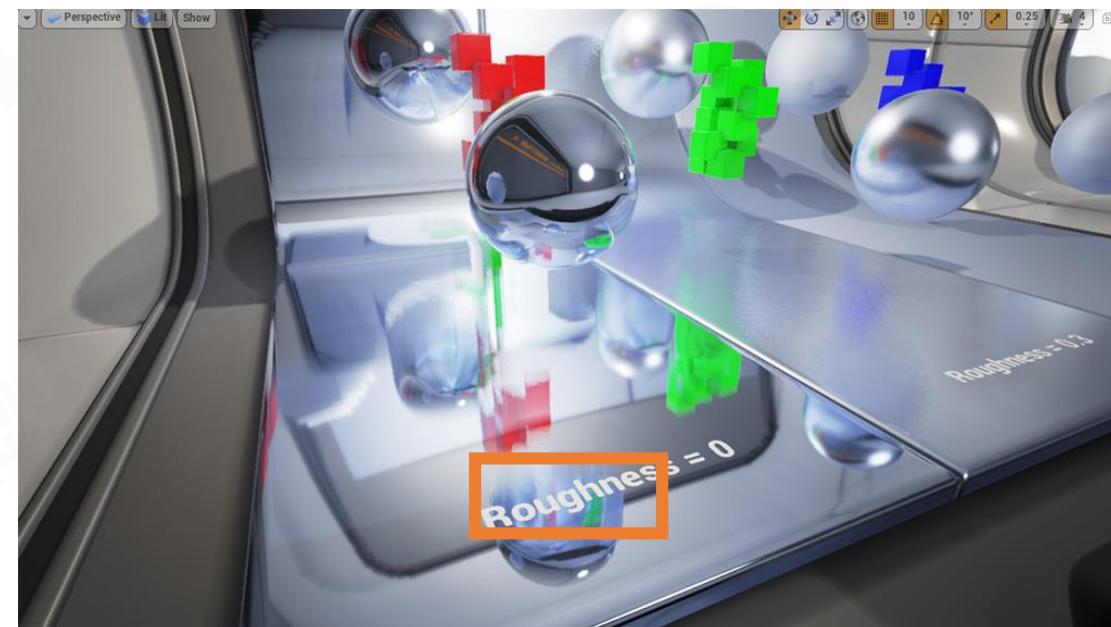
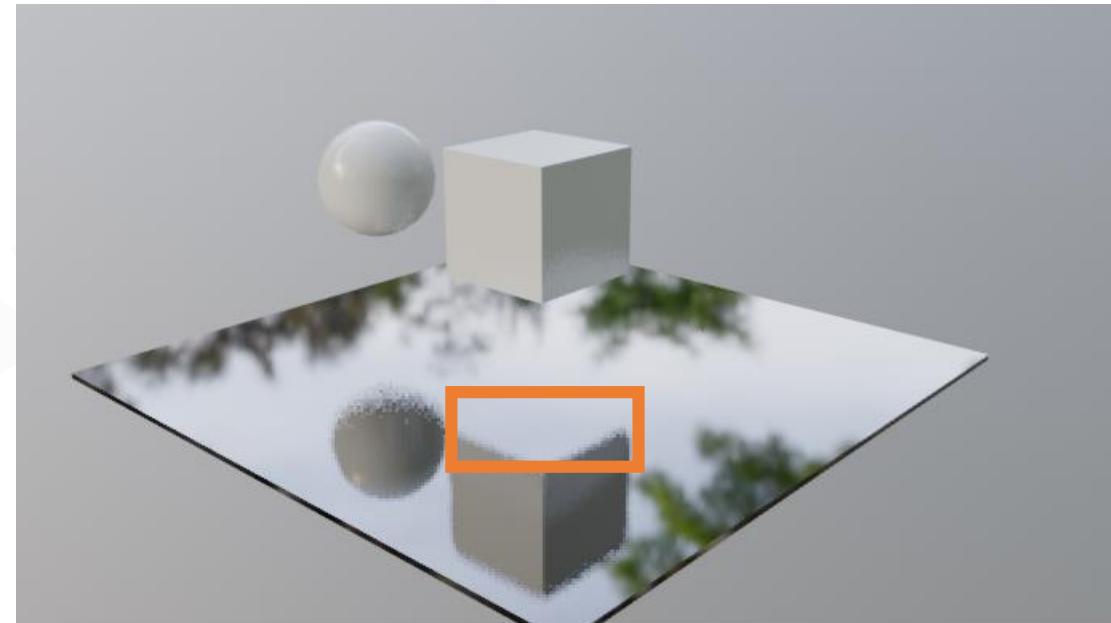
Pre-filter color mipmap (pyramid)





## SSGI Summary

- Pros:
  - Fast for glossy and specular reflections
  - Good quality
  - No occlusion issues
- Cons:
  - Missing information outside screen
  - Affects of incorrect visibility of neighbor ray reuse





## Unique Advantages of SSGI

- Easy to handle close contact shadow
- Precise hit point calculation
- Decouple from scene complexity
- Handle dynamic objects

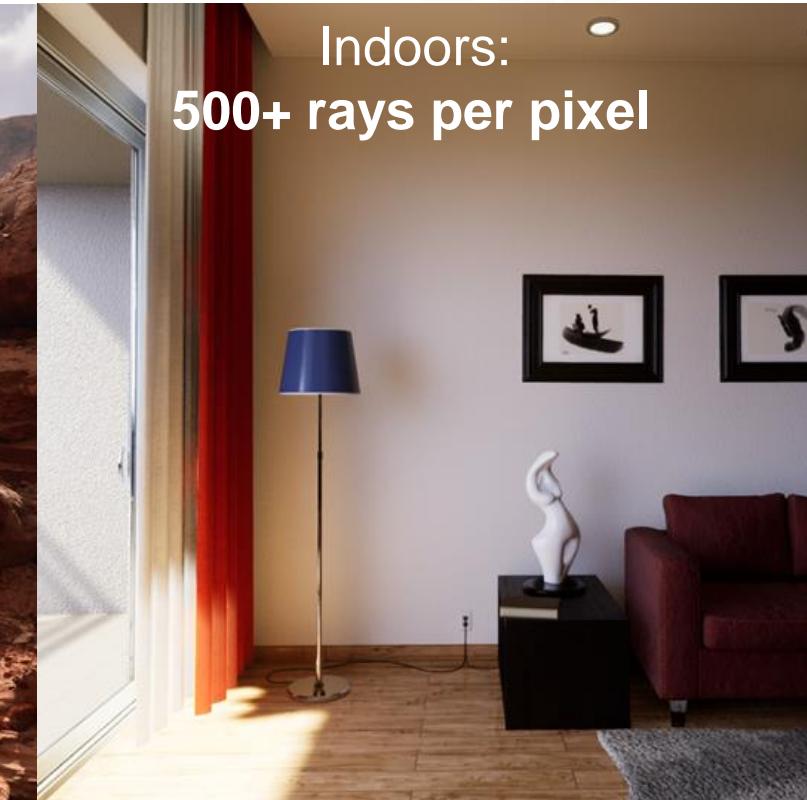
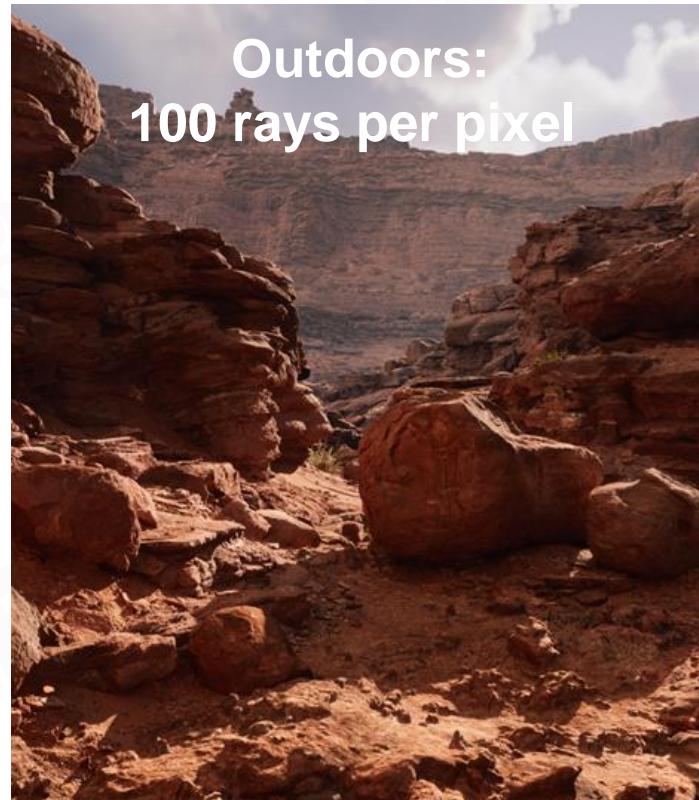
# Lumen





## Ray Traces are slow

- Can only afford 1/2 ray per pixel
- But quality GI needs hundreds





## Sampling is hard

Previous real-time work: Irradiance Fields

- Problems:
  - Leaking and over-occlusion
  - Probe placement
  - Slow lighting update
  - Distinctive flat look

Near bright window



Far bright window



Previous real-time work: Screen Space

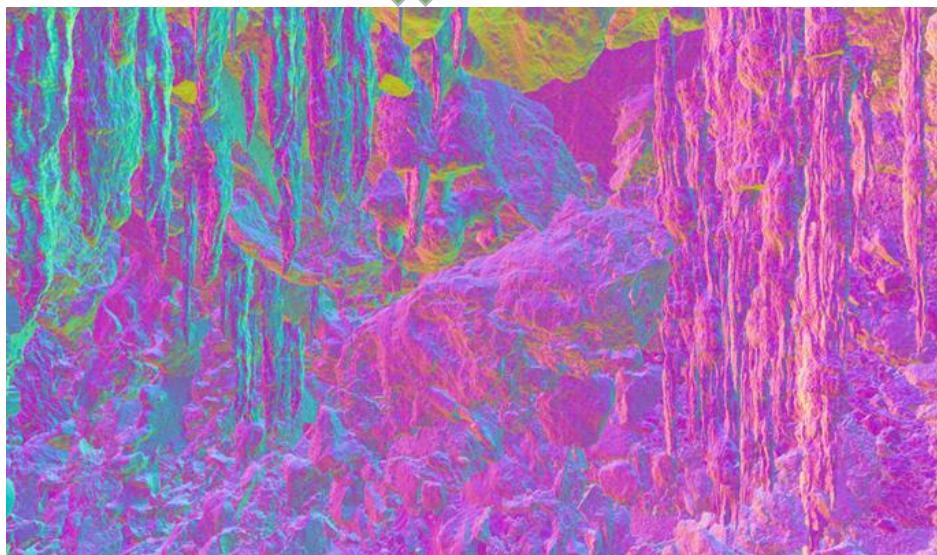
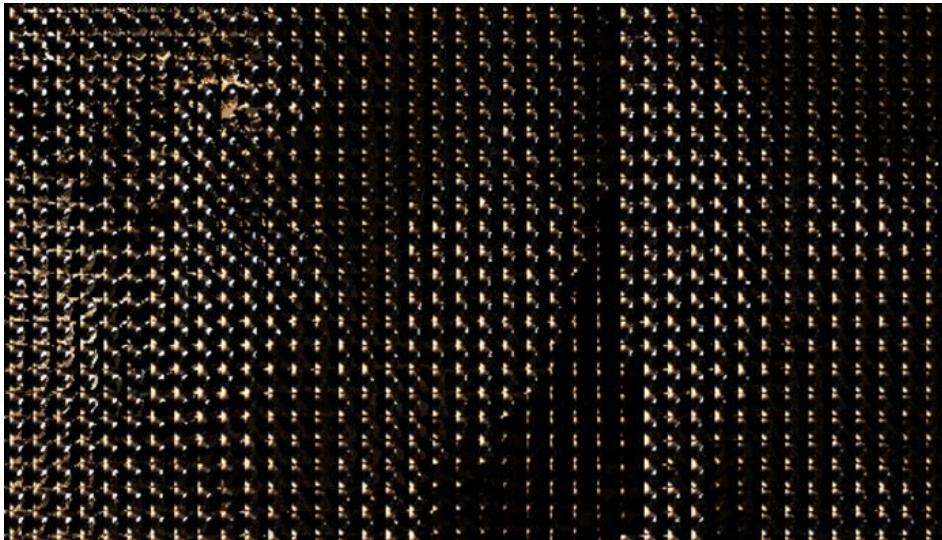
Denoiser

- Problems:
  - Too noisy in many difficult indoor cases
  - Noise is not constant.





## Low-res filtered scene space probes lit full pixels





## Phase 1 : Fast Ray Trace in Any Hardware

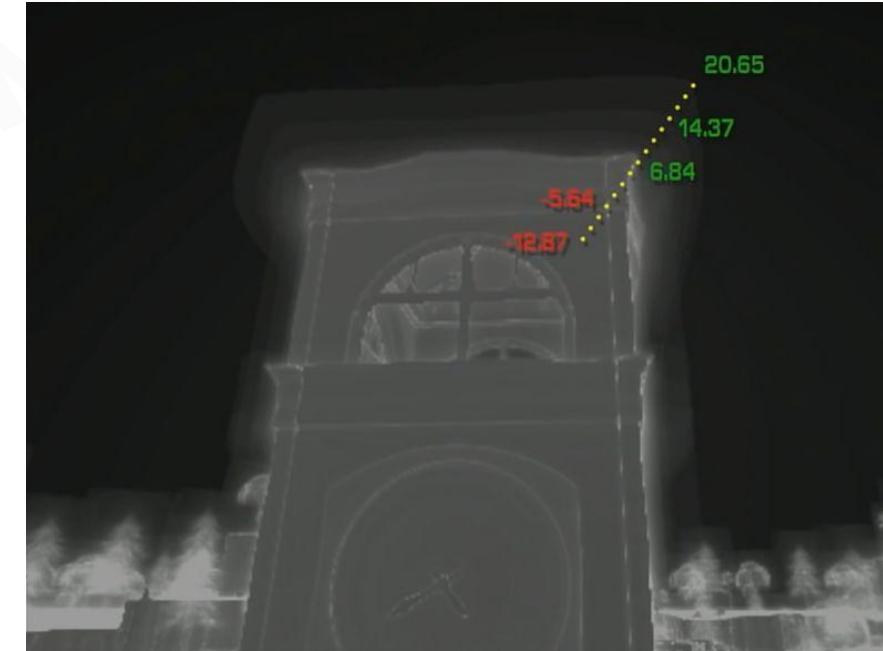


## Signed Distance Field (SDF)



## What is SDF

- The distance to the nearest surface at every point
- Inside regions store negative distance (signed)
- Distance = 0 is the surface





## Per-Mesh SDF

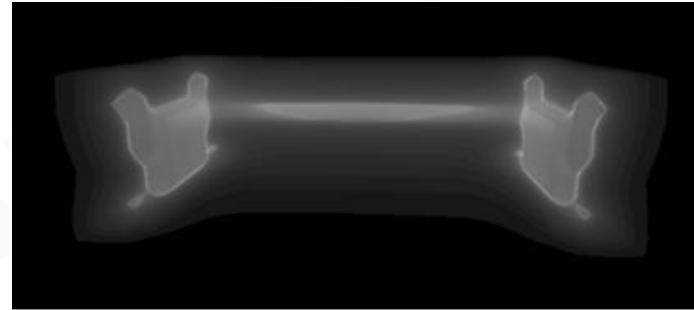
**Store SDF of the whole scene is expensive**

**Generated for each mesh**

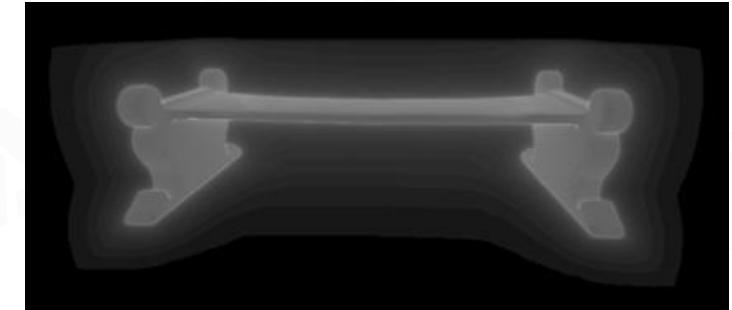
- Resolution based on mesh size
- Embree point query
- Trace rays and count triangle back faces for sign ( more than 25% hit back is negative)



Original Mesh



Resolution is too low,  
important features are lost

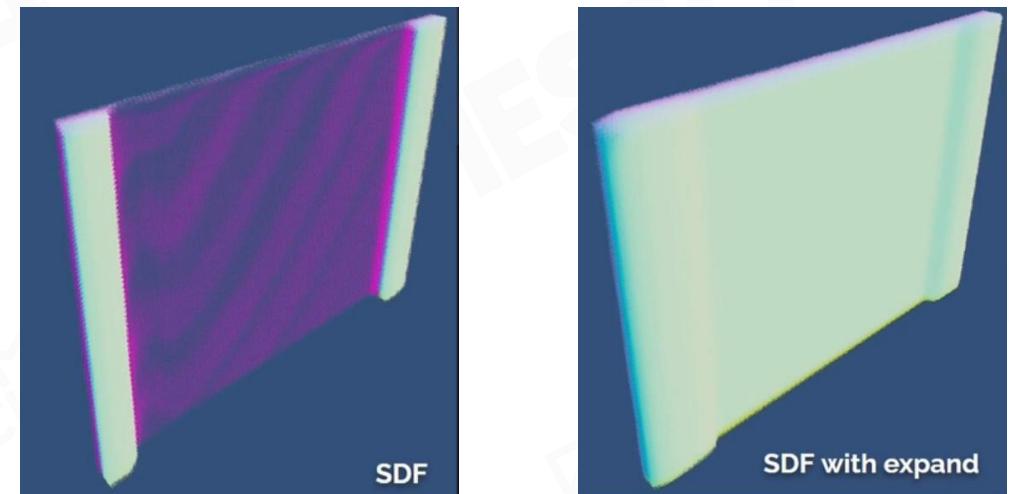
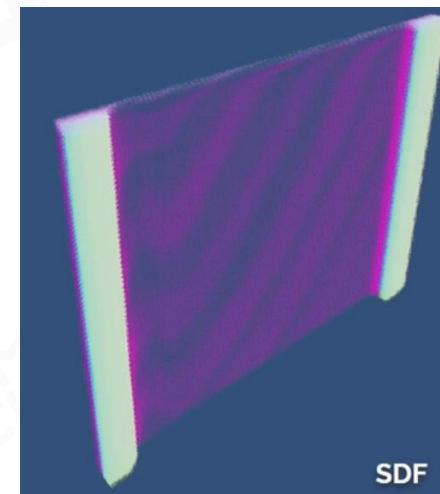
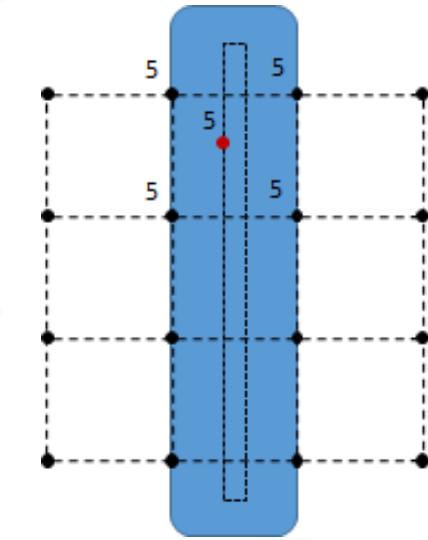
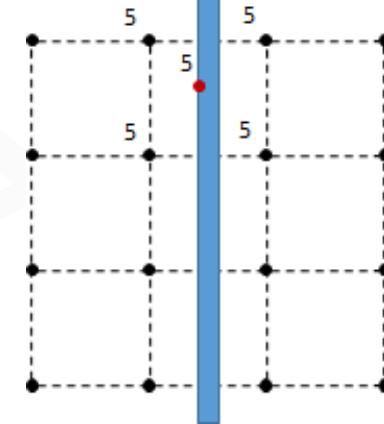


Resolution has been increased,  
important features represented



## SDF for Thin meshes

- Half voxel expand to fix leaking
- Lost contact shadows due to surface bias
  - Over occlusion better than leaking

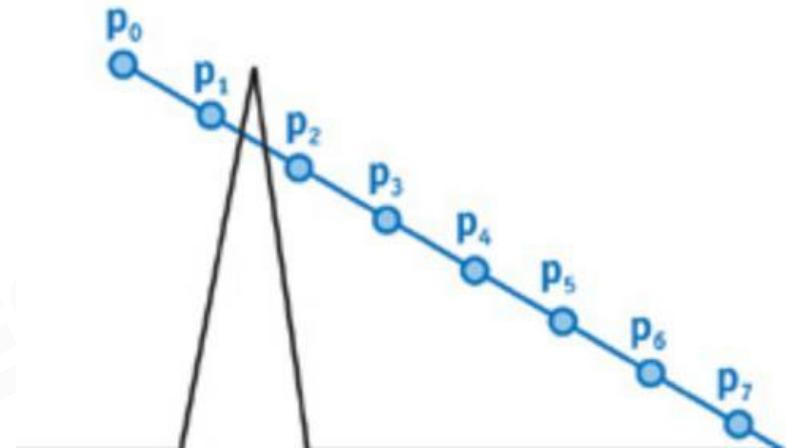




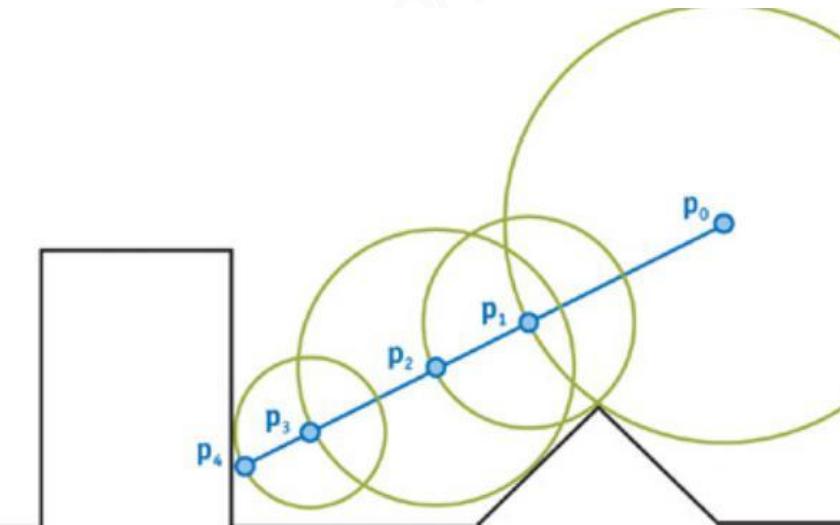
## Ray Tracing with SDF

**Ray intersection skips through empty space based on distance to surface**

- Safe and fast
- Each time at  $p$ , just travel  $\text{SDF}(p)$  distance



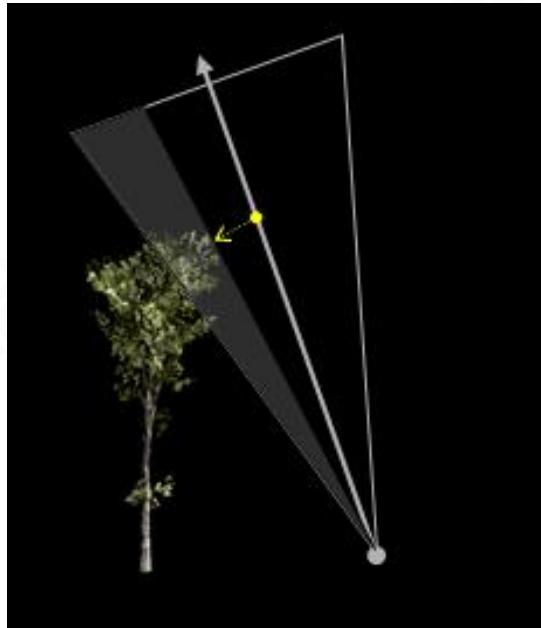
Fixed steps tracing



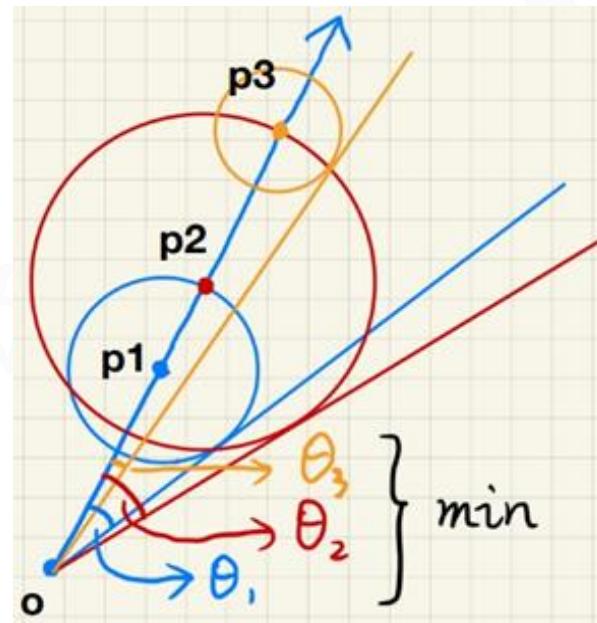
Sphere tracing



# Cone Tracing with SDF(ie. Soft Shadow)

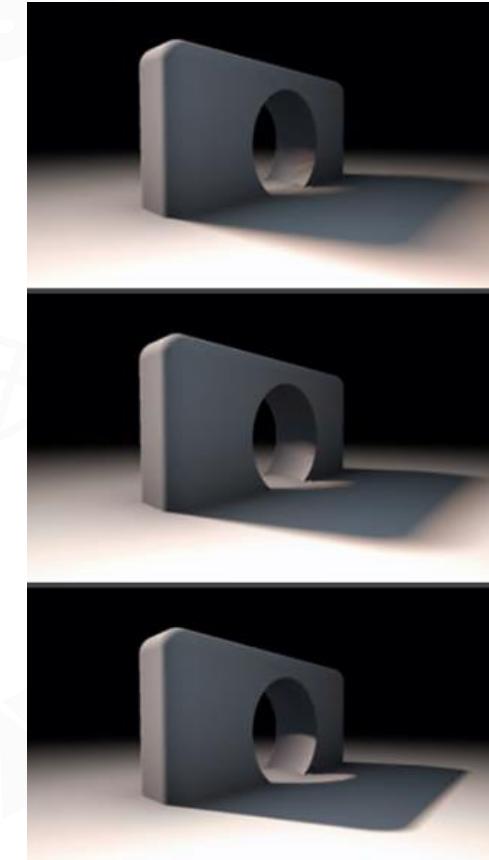


Cone intersection



$$\theta = \arcsin \frac{\text{SDF}(p)}{\|p - o\|}$$

$$\min \theta \approx \min \left\{ \frac{k \cdot \text{SDF}(p)}{\|p - o\|}, 1.0 \right\}$$



$k = 2$

$k = 8$

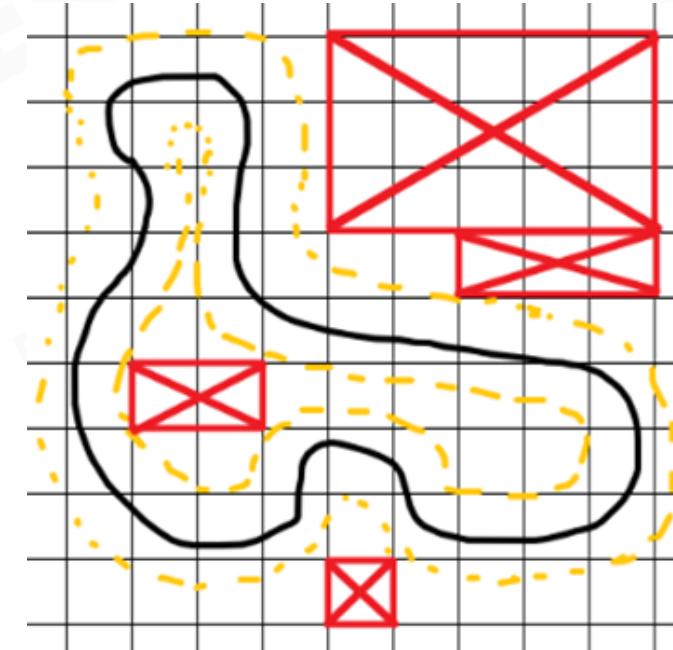
$k = 32$



## Sparse Mesh SDF

**Divides the Mesh SDF into bricks**

- Define a `max_encode_distance`
  - Invalid if  $\forall \text{sdf(brick)} > \text{max\_encode\_distance}$
- `IndirectionTable` store the index of each brick





# Sparse Mesh SDF

Divides the Mesh SDF into bricks

- Define a max\_encode\_distance
  - Invalid if  $\forall \text{sdf(brick)} > \text{max\_encode\_distance}$
- IndirectionTable store the index of each brick

Example:

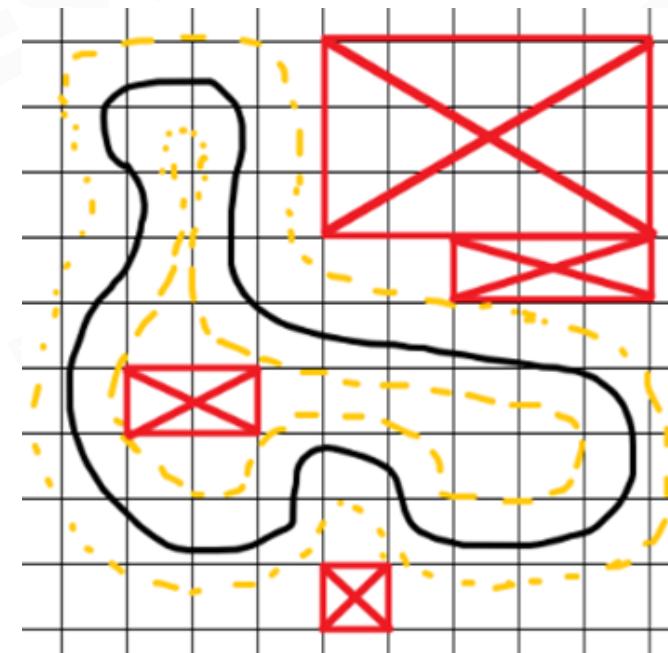
0	1	2
3	4	5
6	7	8

IndirectionTable:

0	1	X	2	X	3	X	4	5
---	---	---	---	---	---	---	---	---

BrickData:

Brick 0	Brick 1	Brick 3	Brick 5	Brick 7	Brick 8
---------	---------	---------	---------	---------	---------





## Mesh SDF LoD

- Every frame GPU gathers requests
- CPU download requests and streams pages in/out
- 3 mips are generated
  - Lowest resolution always loaded and the other 2 streamed



Mip 0



Mip 1



Mip 2



## Sparse Mesh SDF

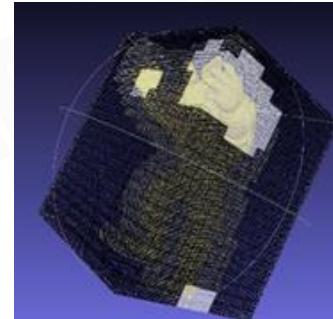
model size = 6.35m x 7.57m x 5.77m

**mip0**

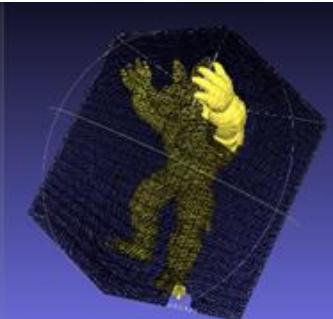
126x154x112

5cm

vaild



invaild



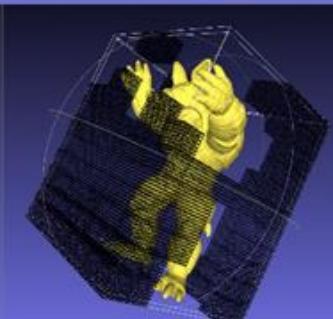
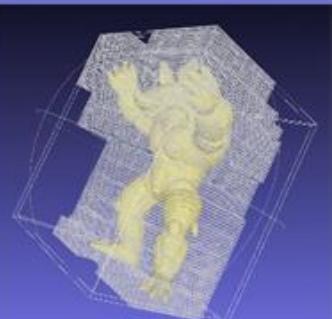
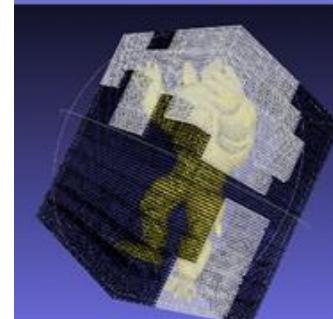
**brick:2310/6336**

1.15MB/3.09MB

**mip1**

63x77x56

10cm



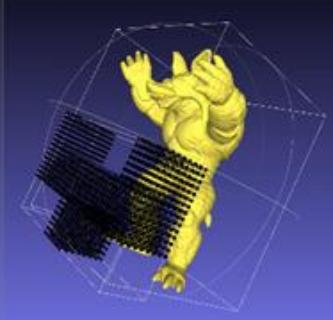
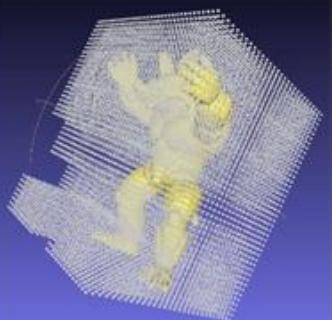
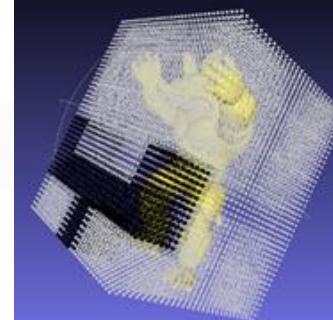
**brick:532/792**

0.26MB/0.39MB

**mip2**

35x42x28

20cm

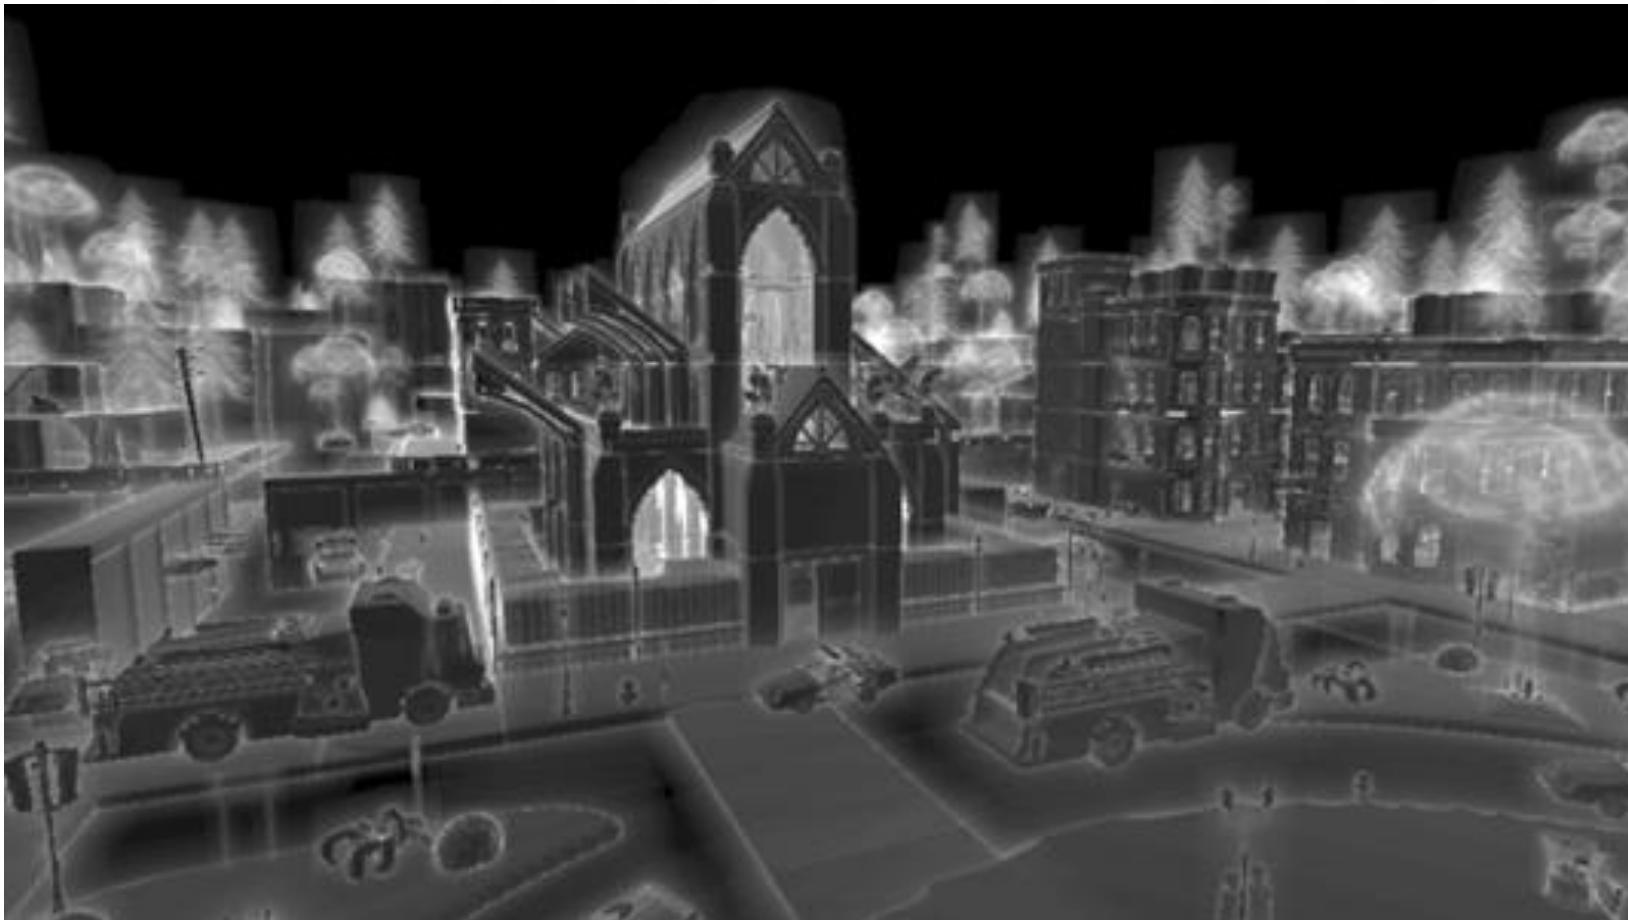


**brick:112/120**

56.5KB/60.0KB



## Ray Tracing Cost in Real Scene



Trace camera rays and visualize the number of steps



## Many Objects along Each Ray

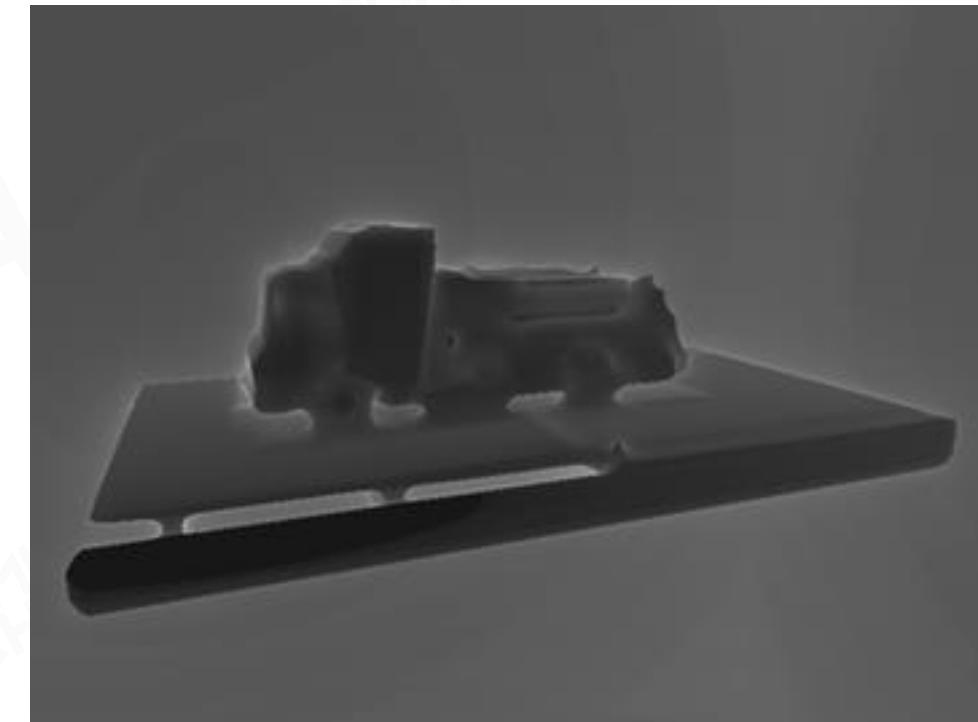
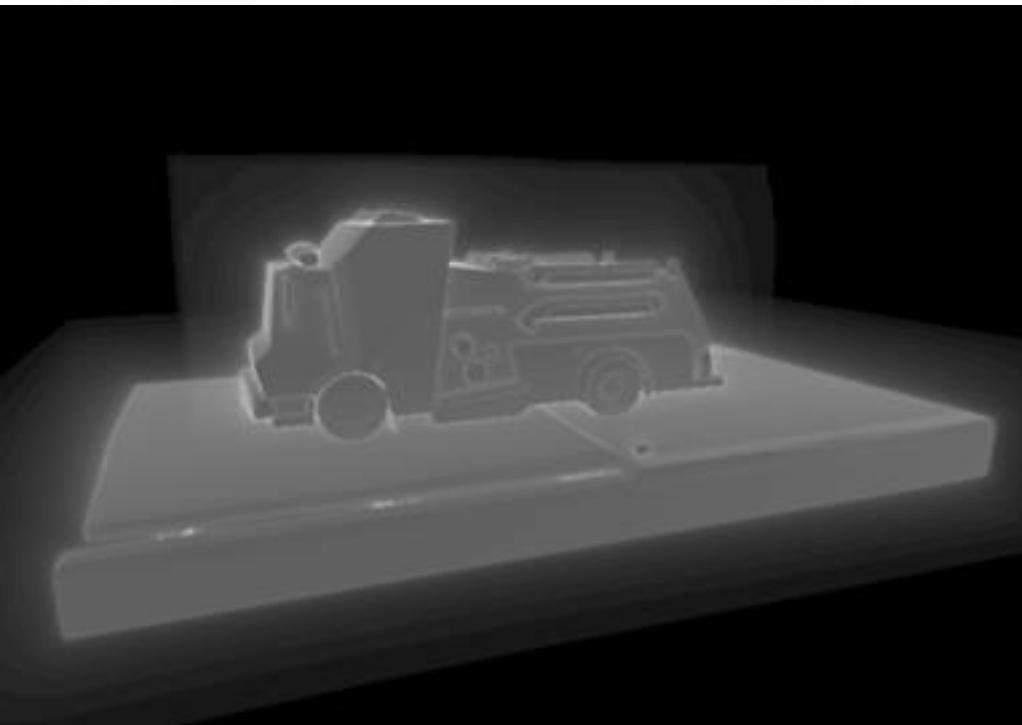


Number of hit objects along each ray



## Global SDF

- Global SDF is inaccurate near surface
- Sample object SDFs near start of cone, global SDF for the rest





## Ray Tracing with Global SDF

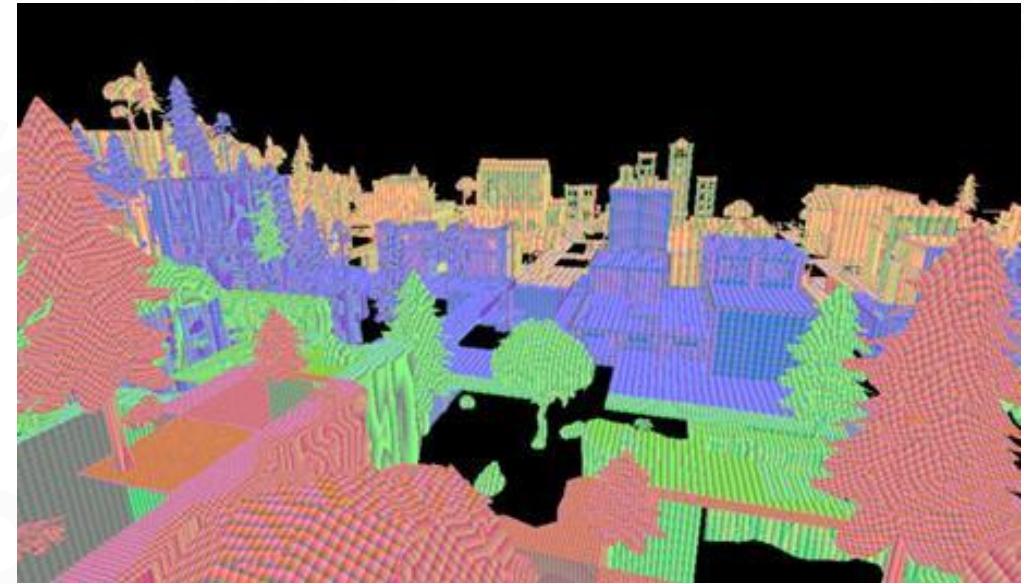
**Massively reduces tracing cost on overlapping objects**





## Cache Global SDF around Camera

- 4 clipmaps centered around camera
- Clipmaps are scrolled with movement
- Distant clipmaps updated less frequently
- Also sparsely stored (~16x memory saving)





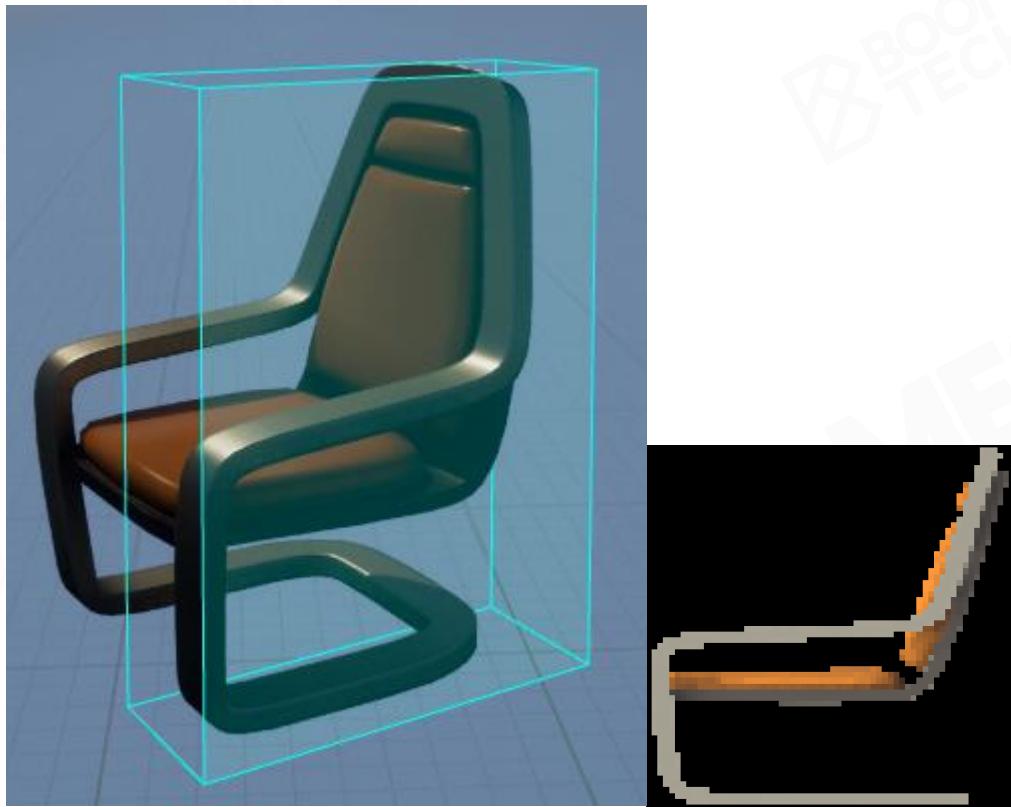
## Phase 2 : Radiance Injection and Caching



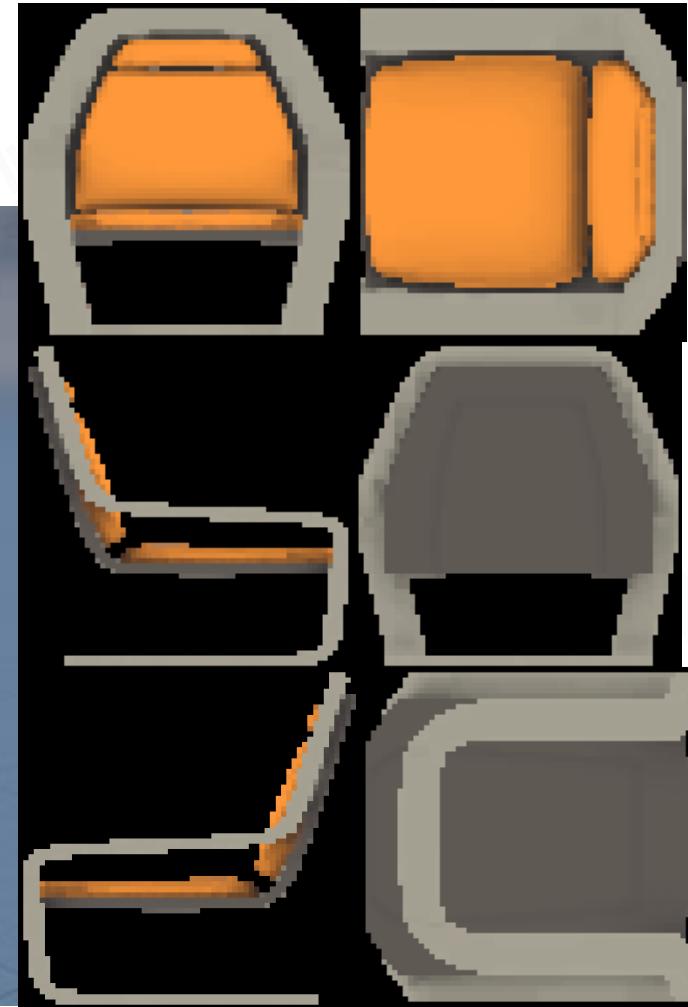
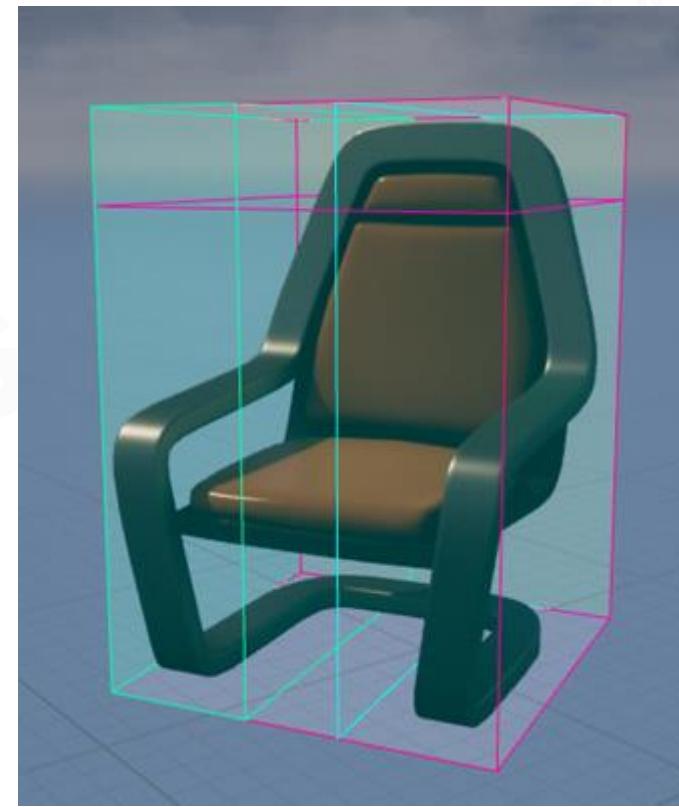
## Mesh card – orthogonal camera on 6-Axis Aligned directions

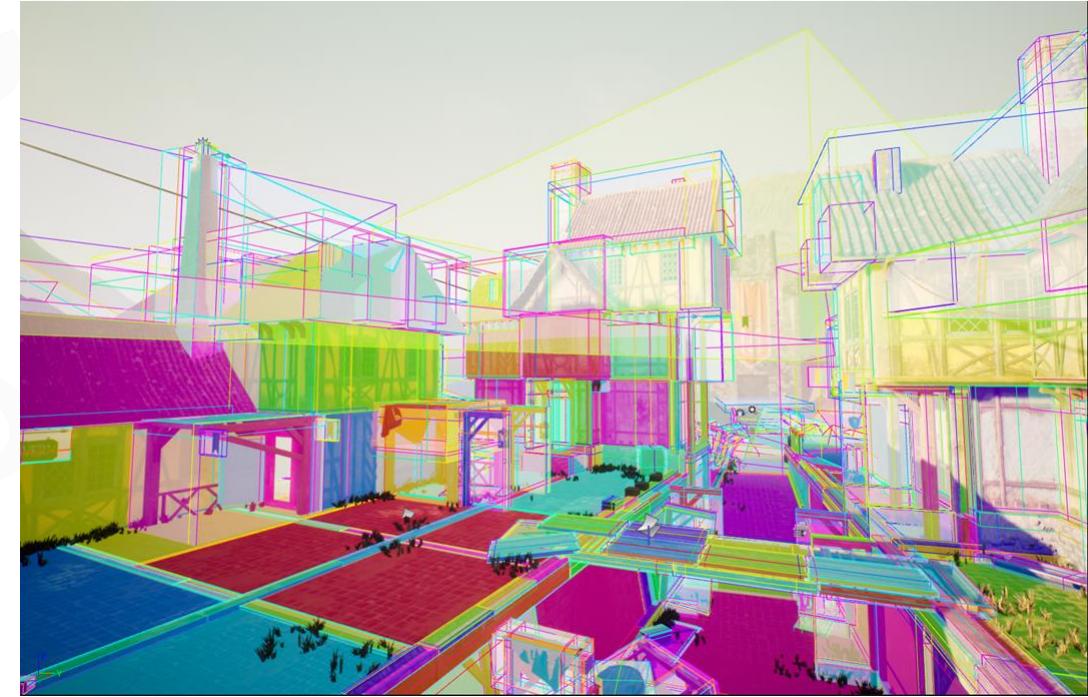
```
class ·FLumenCard
{
    ... ·FLumenCardOBB ·LocalOBB;
    ... ·FLumenCardOBB ·WorldOBB;
    ... uint8 ·AxisAlignedDirectionIndex;
};
```

-Y direction



All 6 directions







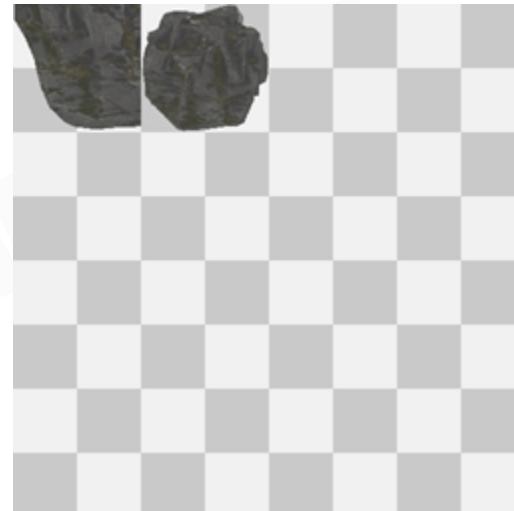
## Generate Surface Cache

### Two Passes

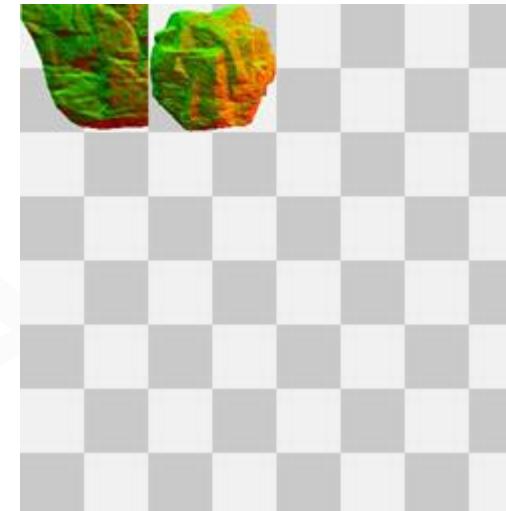
Pass 1: Card capture

- Fix texel budget per frame (512x512)
- Sort by distance to camera and GPU feedback
- Capture resolution depends on card projection on screen

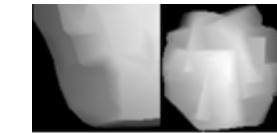
**Albedo**



**Normal**



**Depth**



**Emissive**

512x512 RGBA8

512x512 R8G8

512x512 D32S8

512x512 R11G11B10

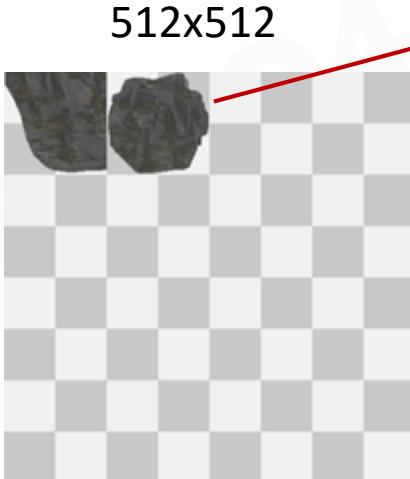


## Generate Surface Cache

### Two Passes

Pass 1: Card capture

Pass 2: Copy cards to surface cache  
and compress

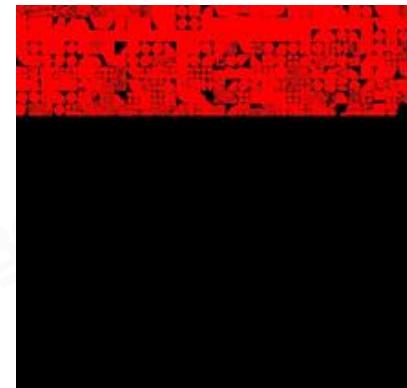
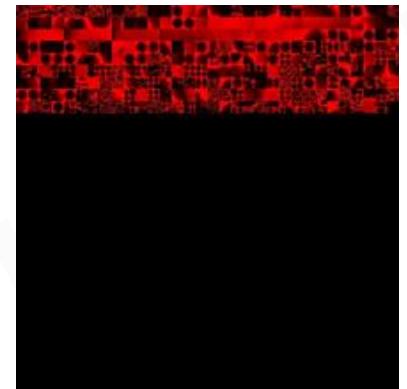
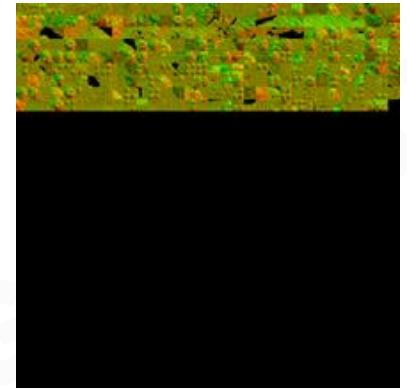


4096x4096



Lumen.CardCaptureAlbedoAltas

Lumens.SceneAlbedo





# Generate Surface Cache

## Two Passes

Pass 1: Card capture

Pass 2: Copy cards to surface cache  
and compress

4096x4096 Surface Cache Atlas			
Albedo	RGB8	BC7	16mb
Opacity	R8	BC4	8mb
Depth	R16	-	32mb
Normal	Hemisphere RG8	BC4	16mb
Emissive	RGB Float16	BC6H	16mb

compress from 320mb to 88mb

```
static FLumenSurfaceLayerConfig Configs[(uint32)ELumenSurfaceCacheLayer::MAX] =  
{  
    { TEXT("Depth"), PF_G16, PF_Unknown, PF_Unknown, FVector(1.0f, 0.0f, 0.0f) },  
    { TEXT("Albedo"), PF_R8G8B8A8, PF_BC7, PF_R32G32B32A32_UINT, FVector(0.0f, 0.0f, 0.0f) },  
    { TEXT("Opacity"), PF_G8, PF_BC4, PF_R32G32_UINT, FVector(1.0f, 0.0f, 0.0f) },  
    { TEXT("Normal"), PF_R8G8, PF_BC5, PF_R32G32B32A32_UINT, FVector(0.0f, 0.0f, 0.0f) },  
    { TEXT("Emissive"), PF_FloatR11G11B10, PF_BC6H, PF_R32G32B32A32_UINT, FVector(0.0f, 0.0f, 0.0f) }  
};
```



## View Dependent Per-Object Card Resolution

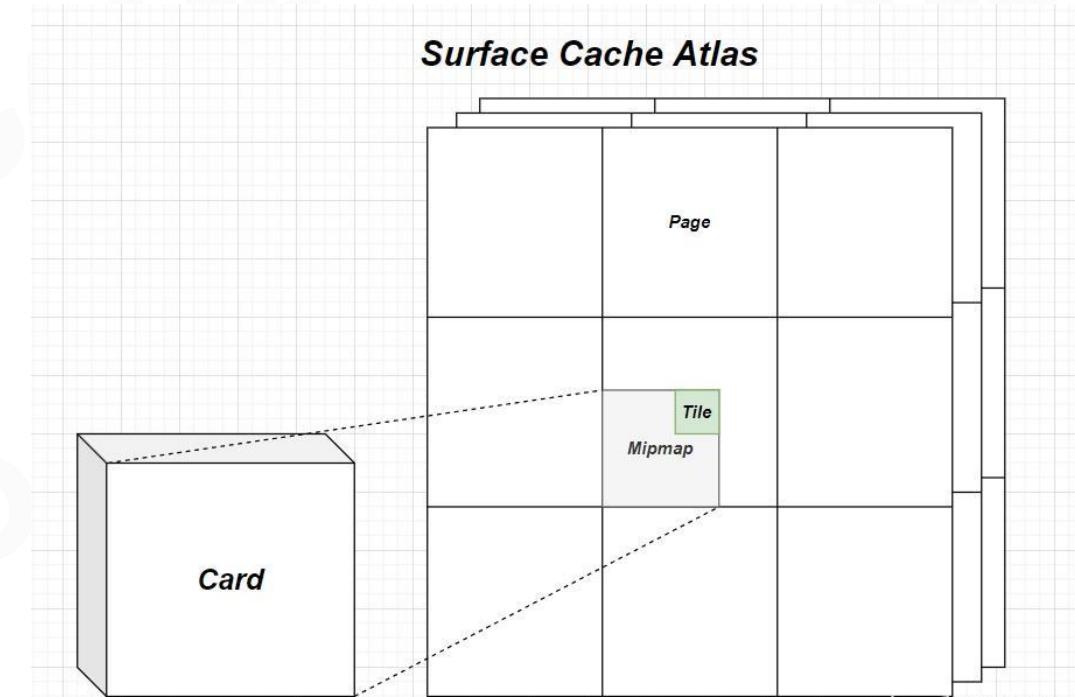
**128x128 physical pages in a 4096x4096 atlas**

**Card capture res  $\geq 128 \times 128$**

- Split into multiple 128x128 physical pages

**Card capture res  $< 128 \times 128$**

- Sub-allocate from a 128x128 physical page





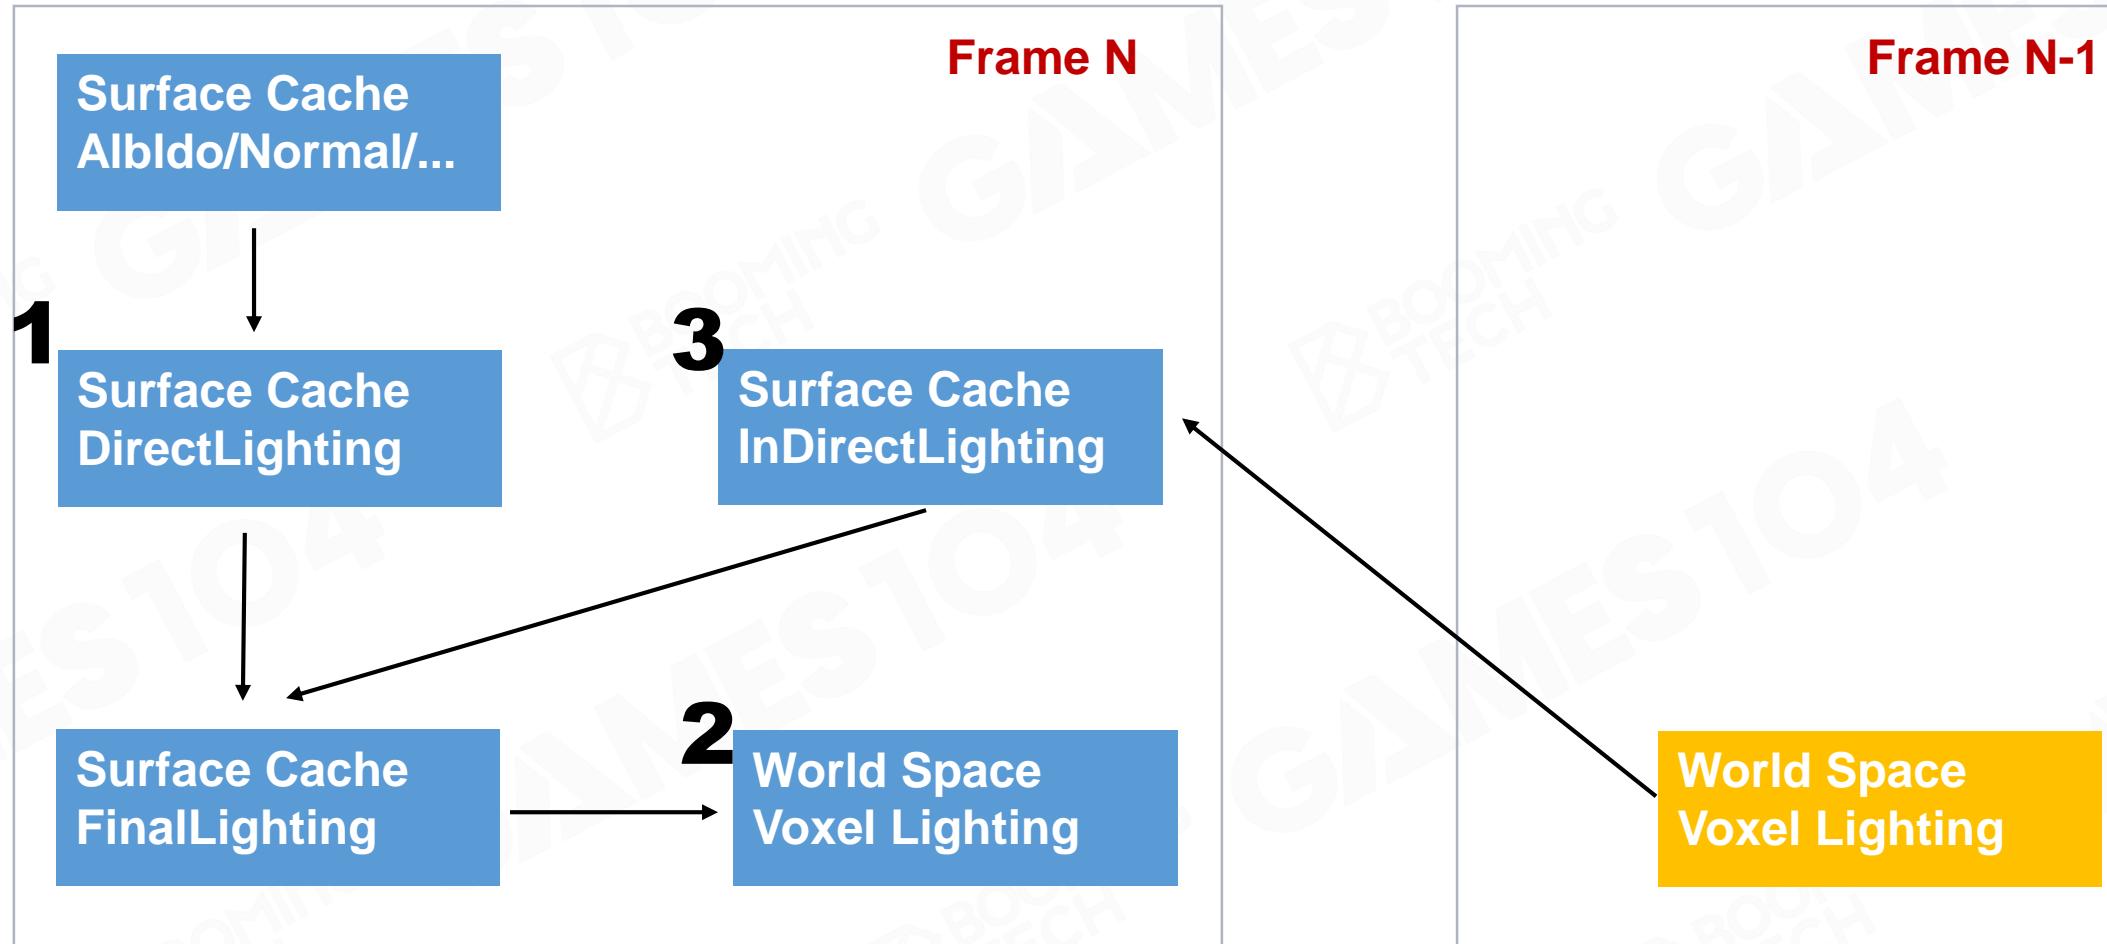
## How can we “freeze” lighting on Surface Cache

### How to compute lighting on hit?

- Is the pixel under the shadow
- How can we handle multi-bounce



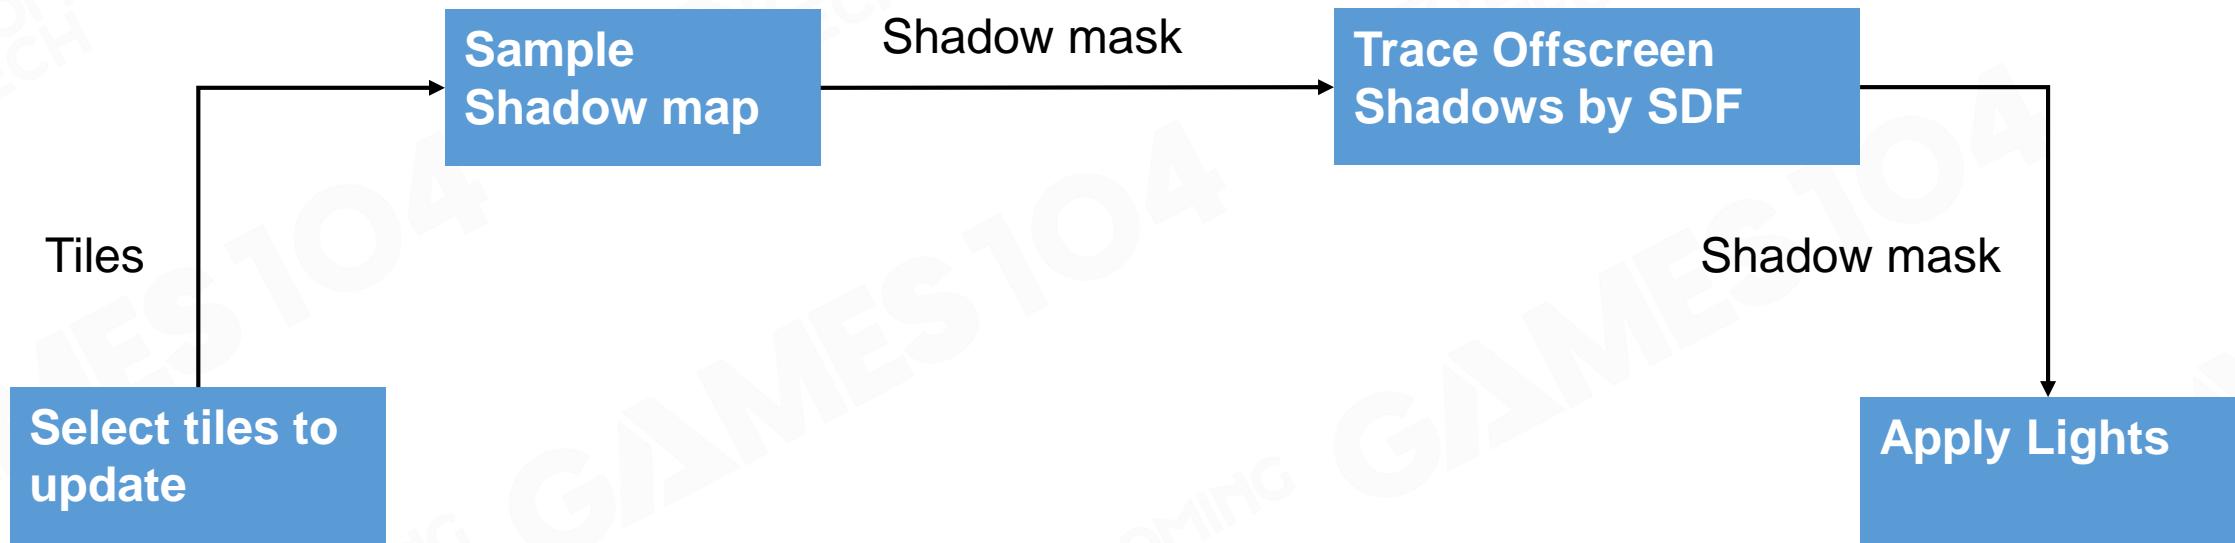
## Lighting Cache Pipeline





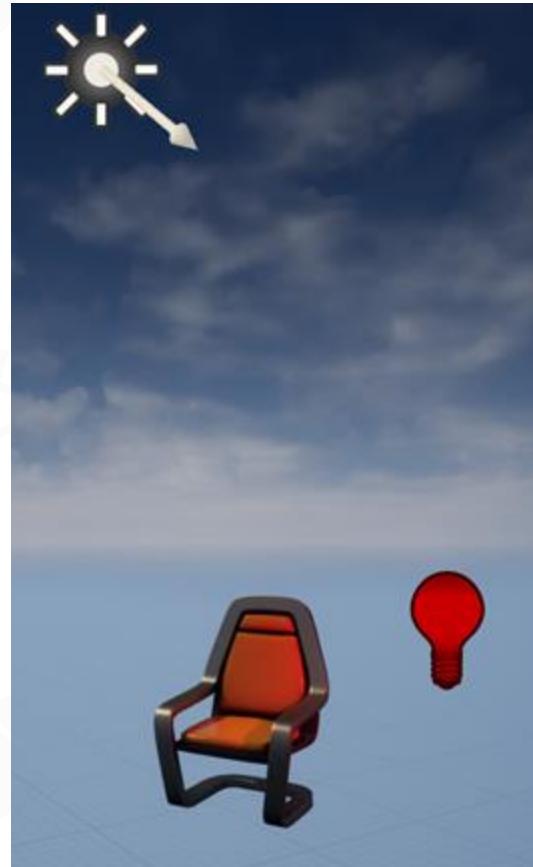
## Direct Lighting

- Divide 128x128 page into 8x8 tiles
- Cull lights with 8x8 tile
- Select first 8 lights per tile
- 1 bit shadow mask

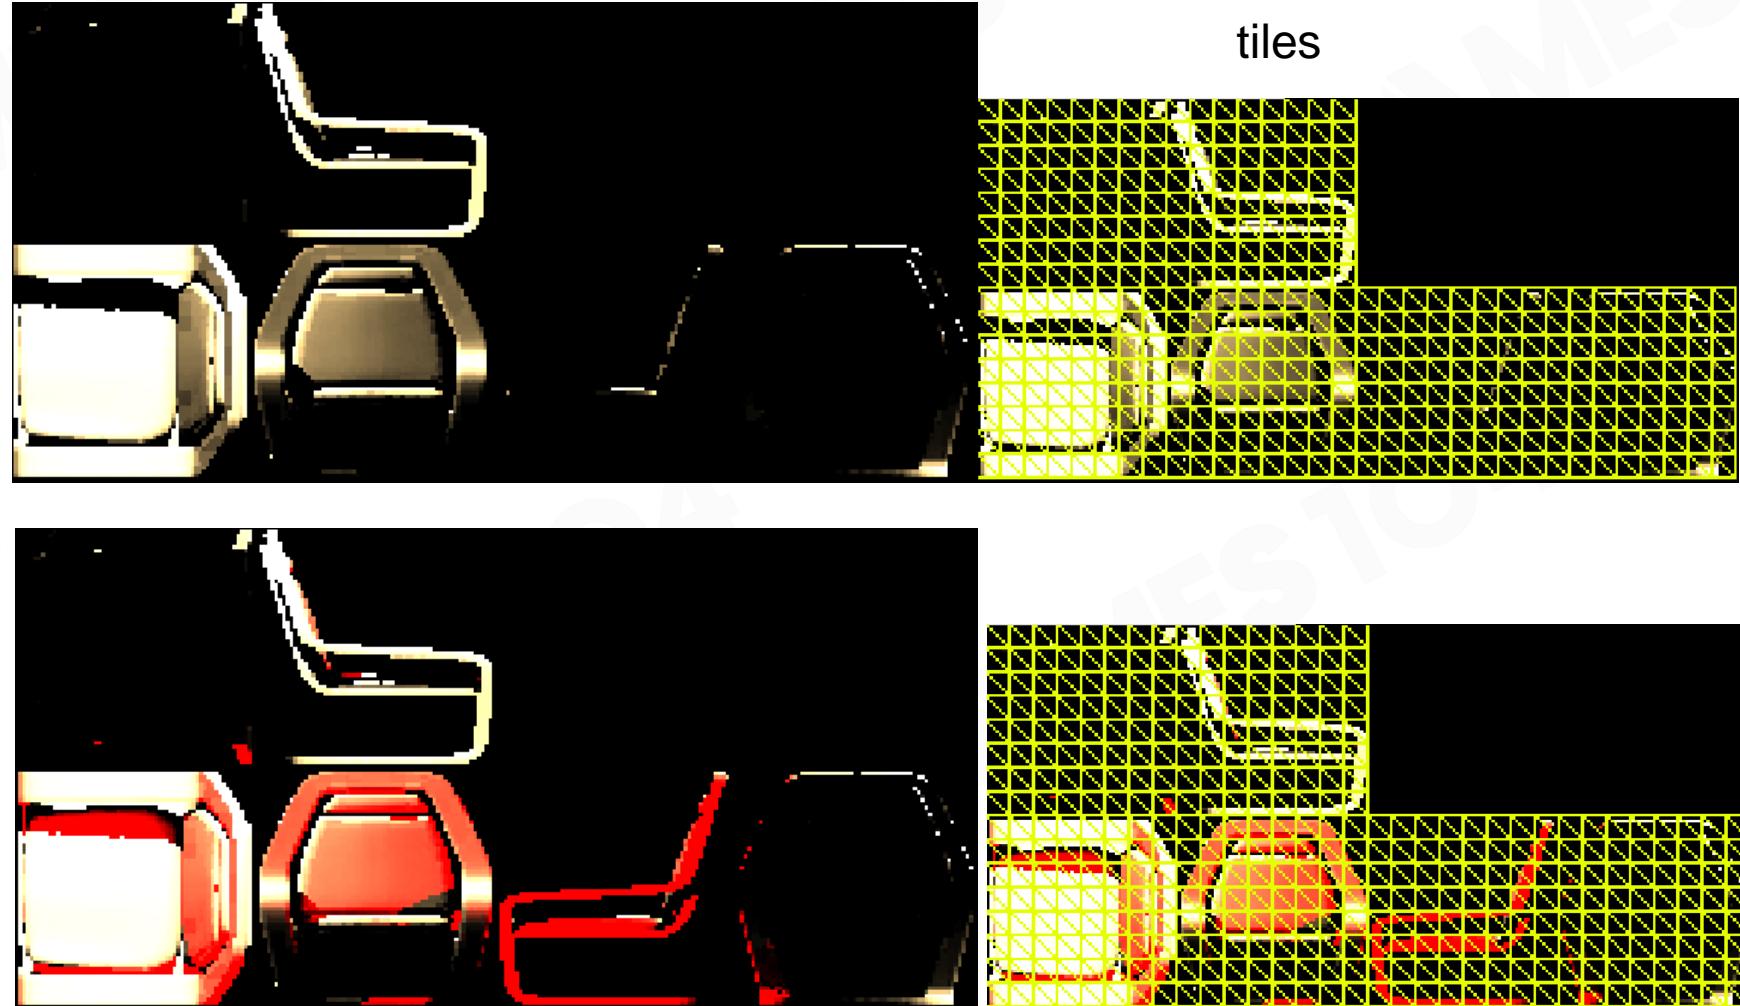




One tile can be lited by multi lights, the result will be accumulated



▼ Lights  
  > Minimal\_Default.Light Source  
  > Minimal\_Default.PointLight

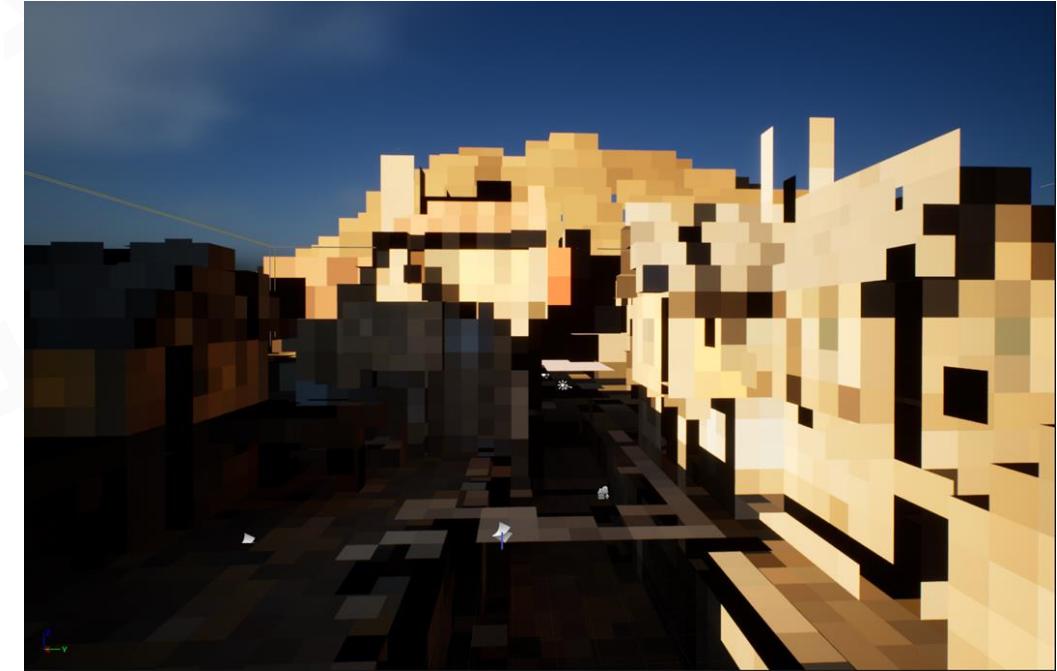




## Global SDF can't sample surface cache

- no per mesh information, only hit position and normal

Use voxel lighting to sample





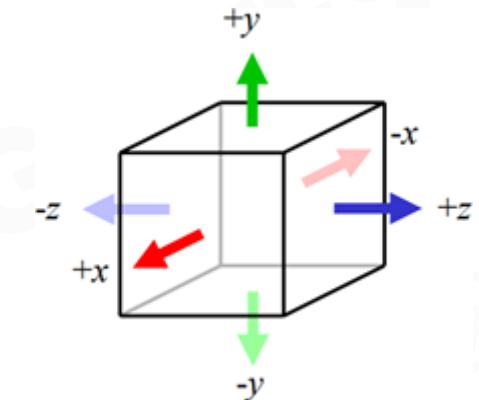
# Voxel Clipmap for Radiance Caching of the Whole Scene

## 4 level clipmaps of 64x64x64 voxels

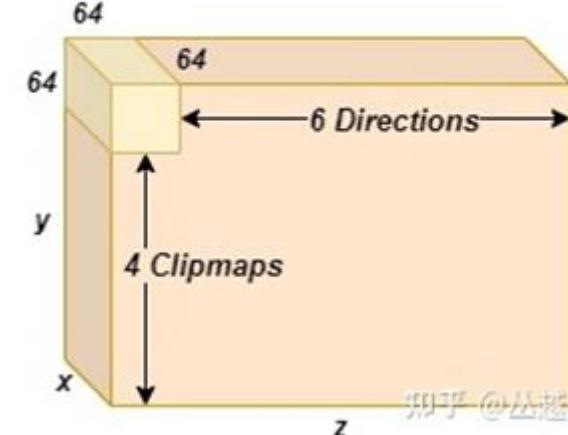
- Radiance per 6 directions per voxel
- Sample and interpolate 3 directions by normal
- Clipmap0 cover  $50m^3$ , voxel size is 0.78m
- Store in 3D texture

## Clipmap update frequency rules

	Clipmap 0	Clipmap 1	Clipmap 2	Clipmap3
Start_Frame	0	1	3	7
Update_interval	2	4	8	8



## Voxel Lighting

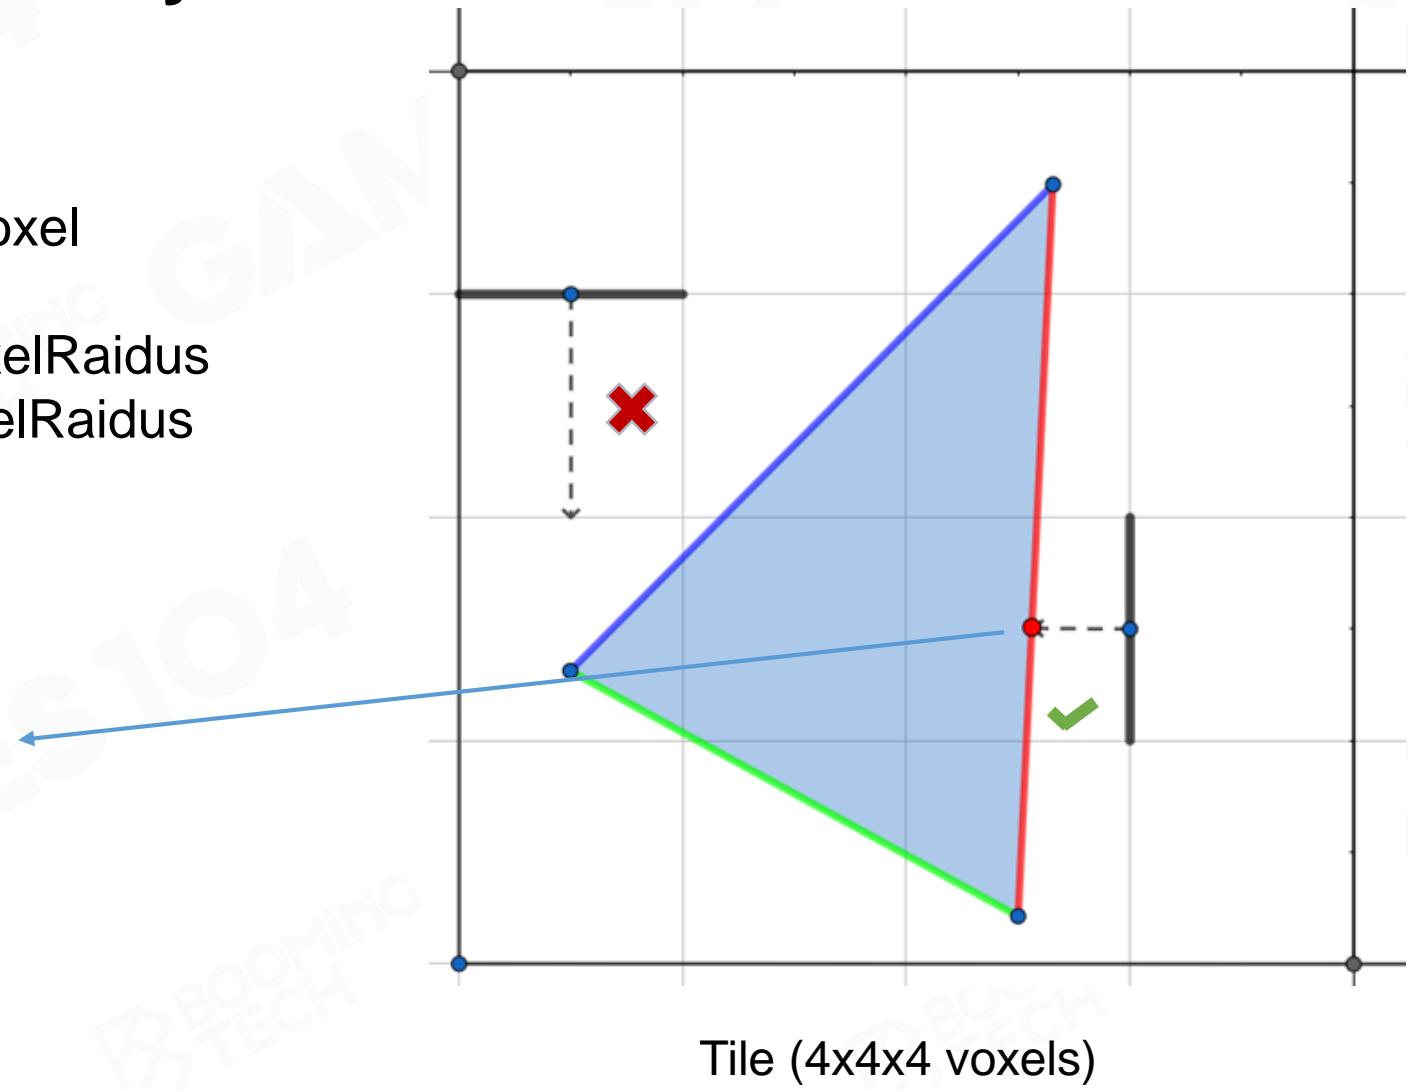




## Build Voxel Faces by Short Ray cast

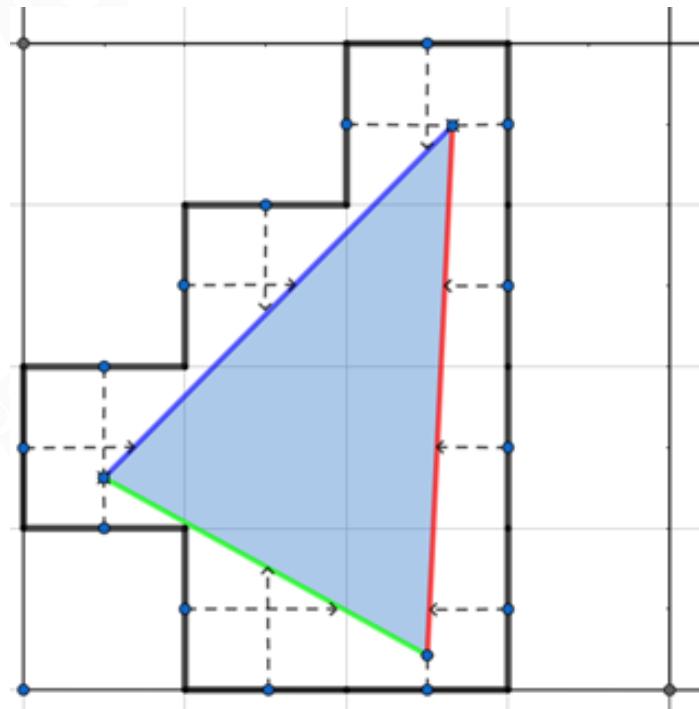
- Trace mesh DF on 6 directions per voxel
- Hit mesh id and hit distance
- $\text{RayStart} = \text{VoxelCenter} - \text{AxisDir} * \text{VoxelRadius}$
- $\text{RayEnd} = \text{VoxelCenter} + \text{AxisDir} * \text{VoxelRadius}$

store hit info into visibility buffer  
`uint32 [Hit distance| Hit object id]`





## Filter Most Object Out by 4x4x4 Tiles

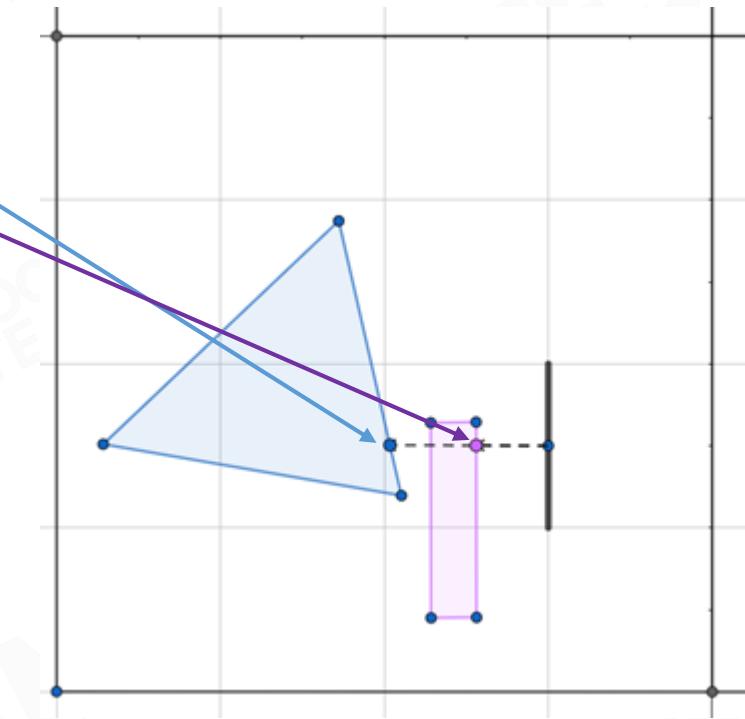


Tile (4x4x4 voxels)

[Hit distance| Hit object id]

[Hit distance| Hit object id]

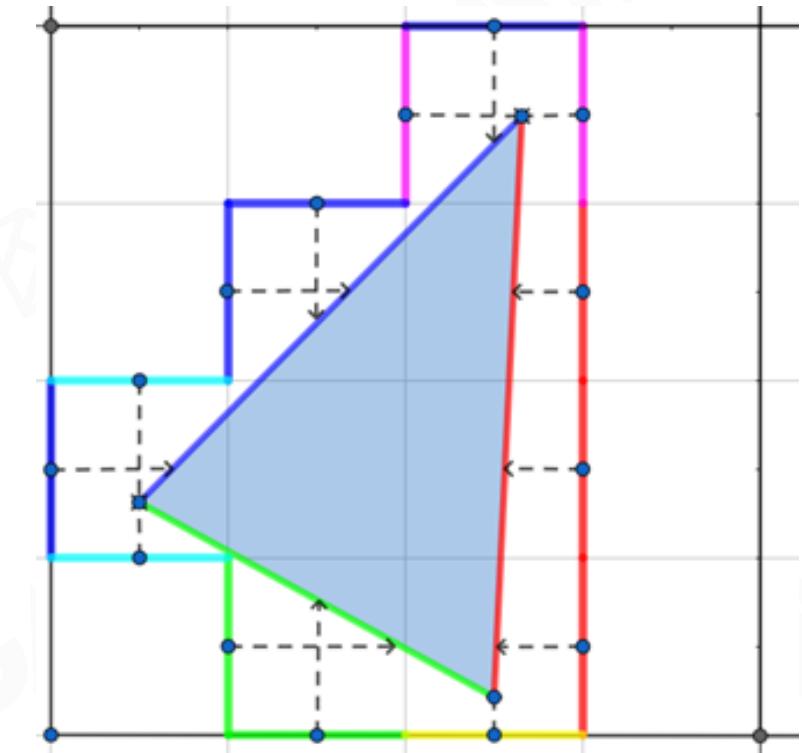
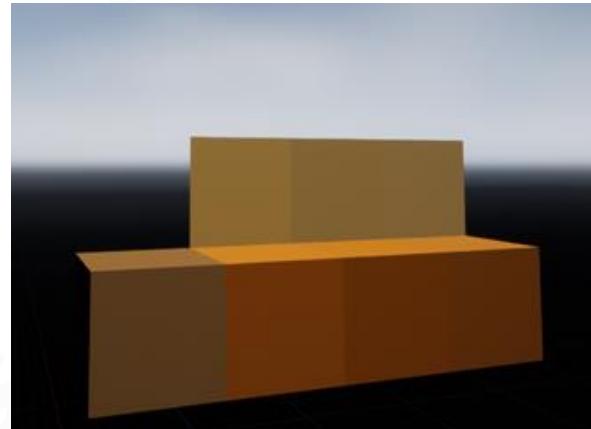
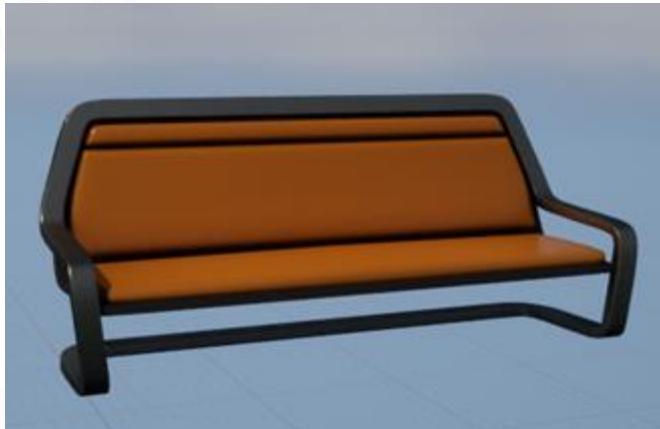
InterlockdMin





## Inject light into clipmap

- Clear all voxel lighting in entire Clipmap
- Compact all valid VisBuffer in Clipmap
- Sampling FinalLighting from VisBuffer and inject lighting



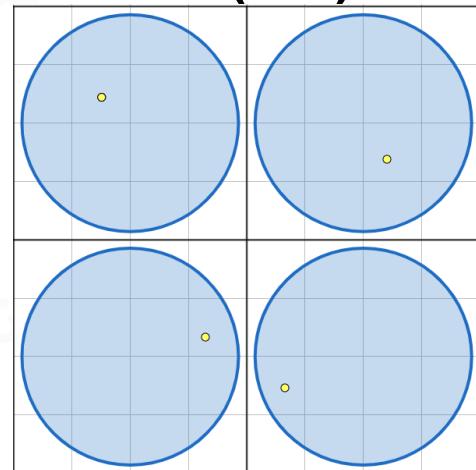


## Indirect Lighting

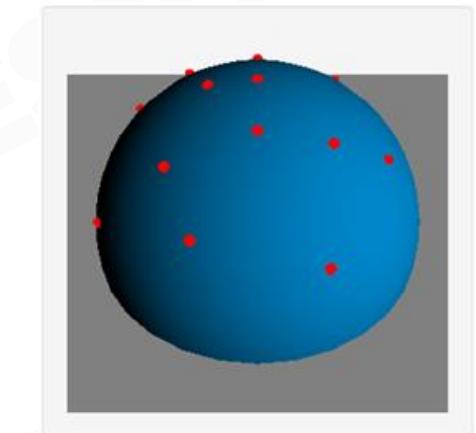
- Place 2x2 probes on each tile - each probe cover 4x4 texels
- Trace 16 rays from hemisphere per probe
- Jitter probe placement and ray directions



Tile(8x8)



probe



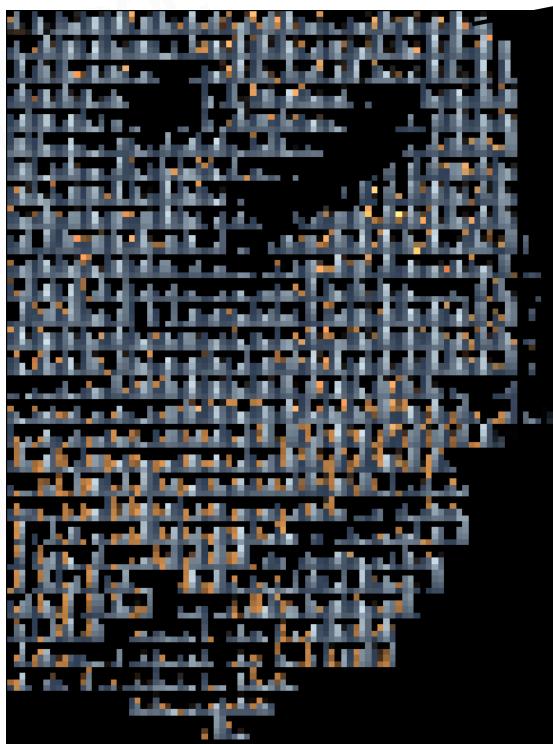
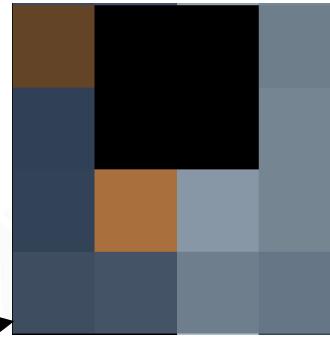
16 rays



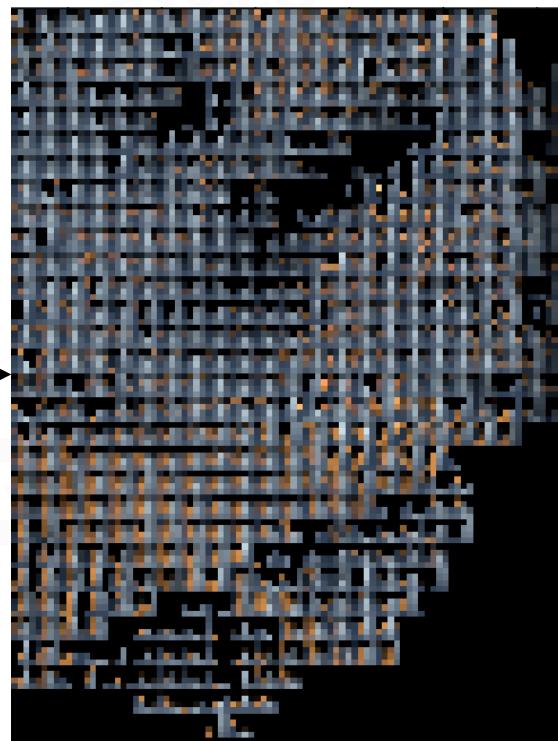
## Indirect Lighting

- Spatial filtering between probes
- Convert to TwoBandSH(store in half4)

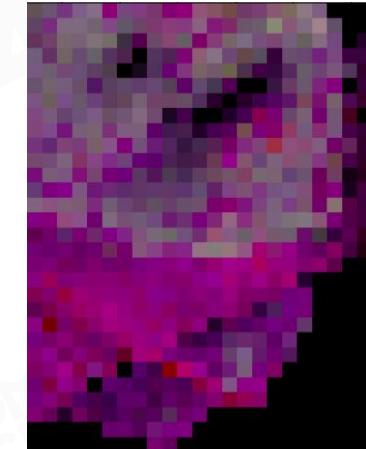
4x4 radiance alphas per probe



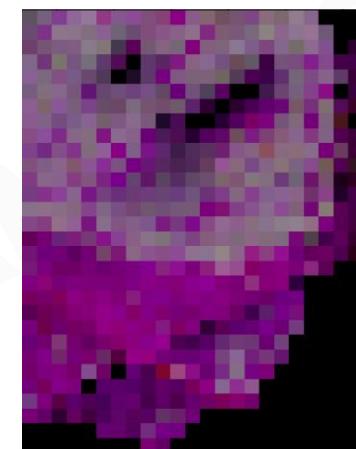
Filter



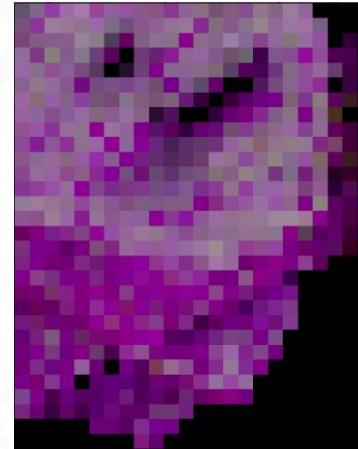
ProbeSHRed



ProbeSHGreen



ProbeSHBlue

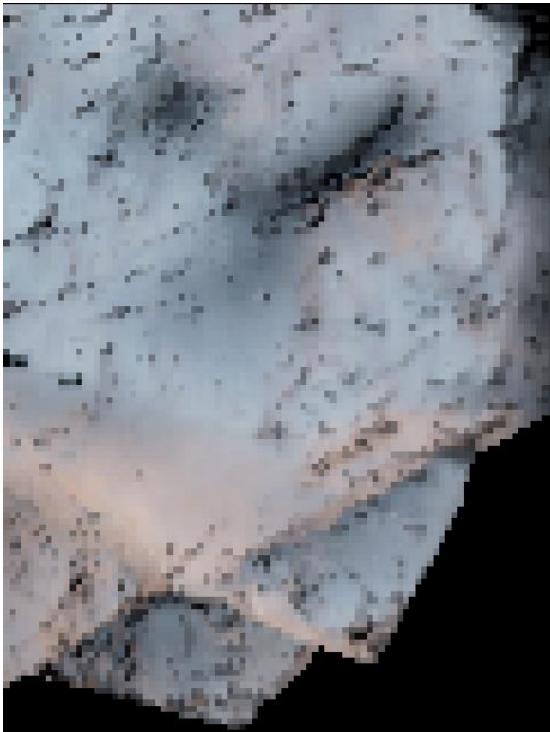




## Per-Pixel Indirect Lighting with 4 Probe Interpolation

- Integrate on pixel - bilinear interpolation of 4 neighbor probes

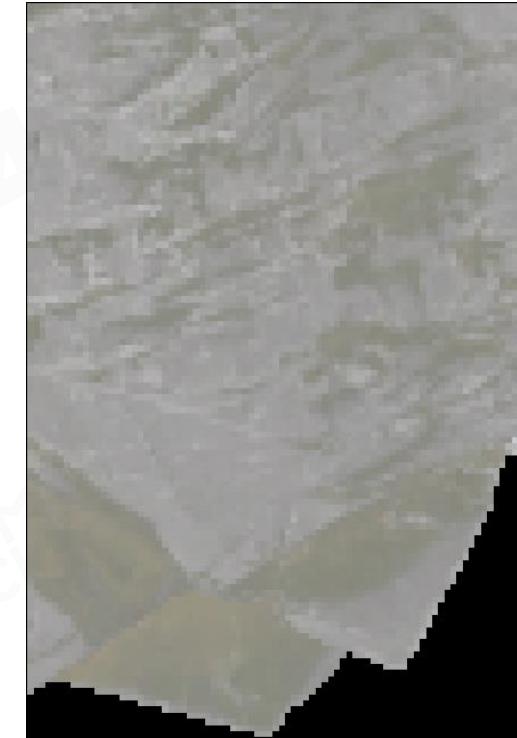
IndirectLighting



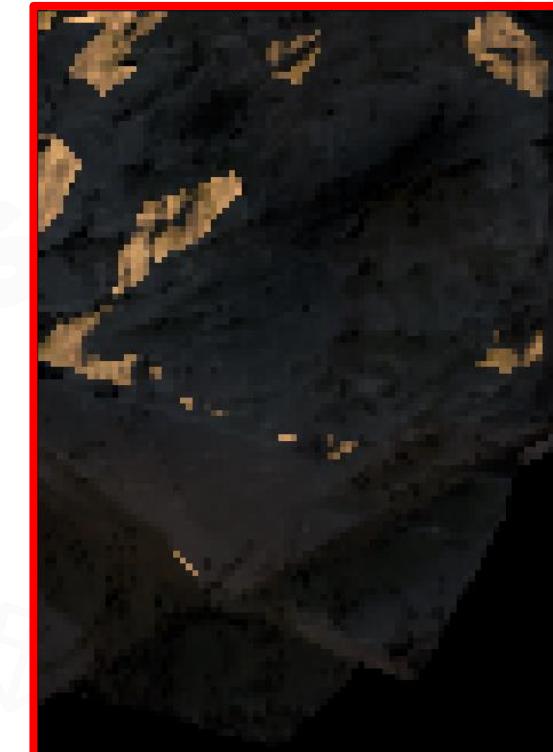
DirectLighting



Albedo



FinalLighting



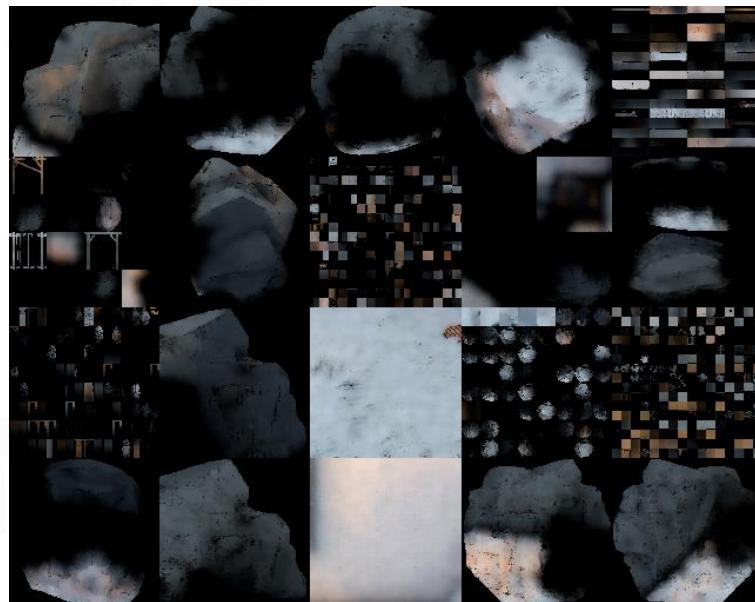


## Combine Lighting

FinalLighting = (DirectLighting + IndirectLighting) \* Diffuse\_Lambert(Albedo)  
+ Emissive;



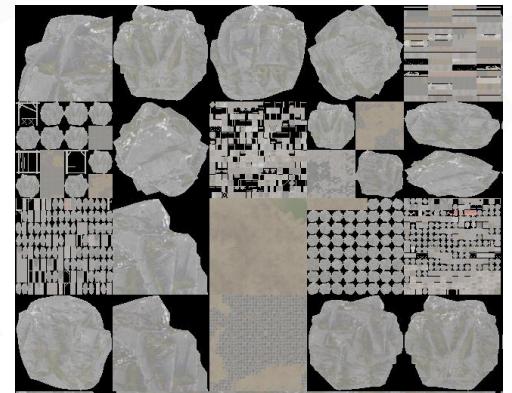
DirectLighting  
(HDR)



IndirectLighting  
(LDR)



FinalLighting



Albedo



## Lighting Update Strategy

### Fix budget

- 1024x1024 texels for direct lighting
- 512x512 texels for indirect lighting
- Select pages to update based on Priority = LastUsed - LastUpdated

### Priority queue using bucket sort

- 128 buckets
- Update buckets with priority until reaching budget



## Phase 3 : Build a lot of Probes with Different Kinds



## Screen Space Probe

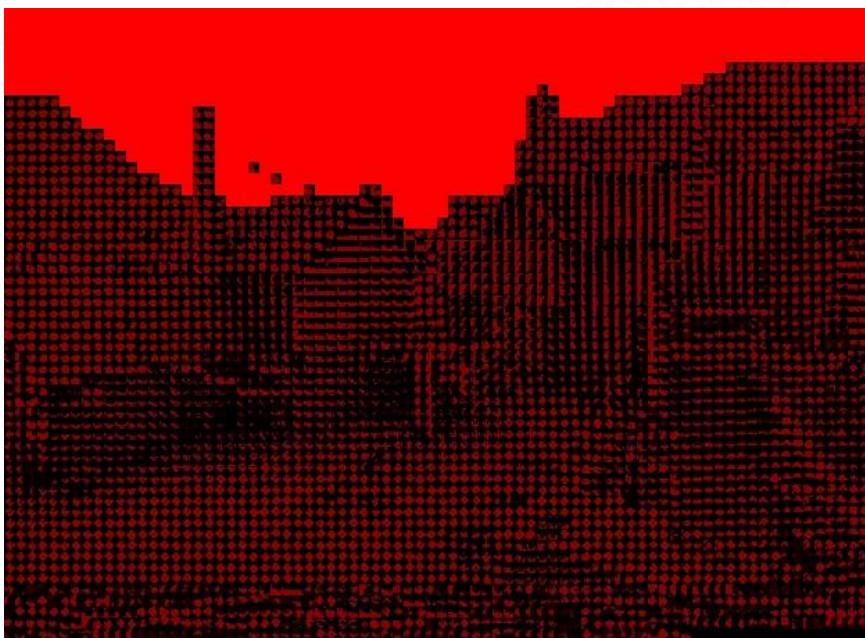


## Screen Probe structure

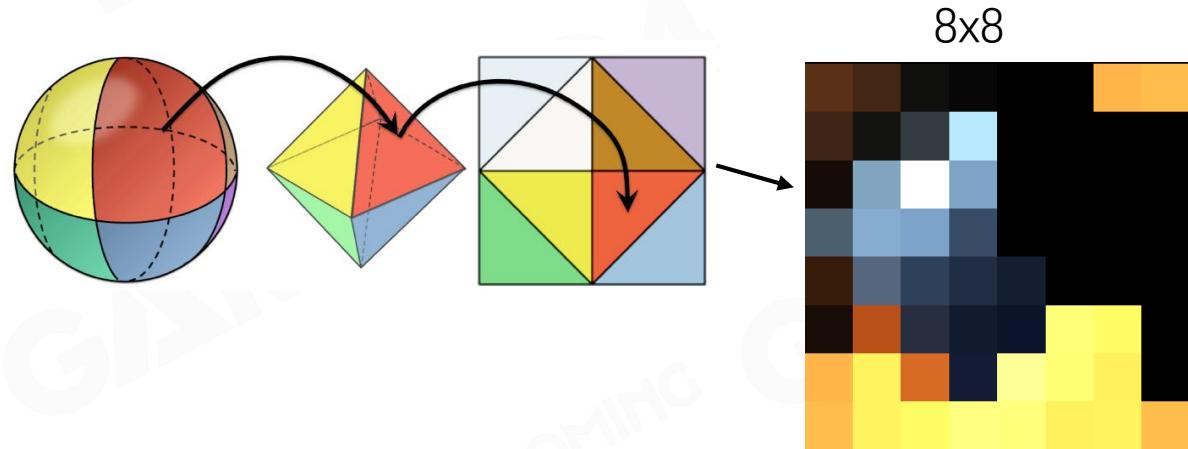
Octahedral atlas with border

- Typically 8x8 per probe
- Uniformly distributed world space directions
- Neighbors have matching directions

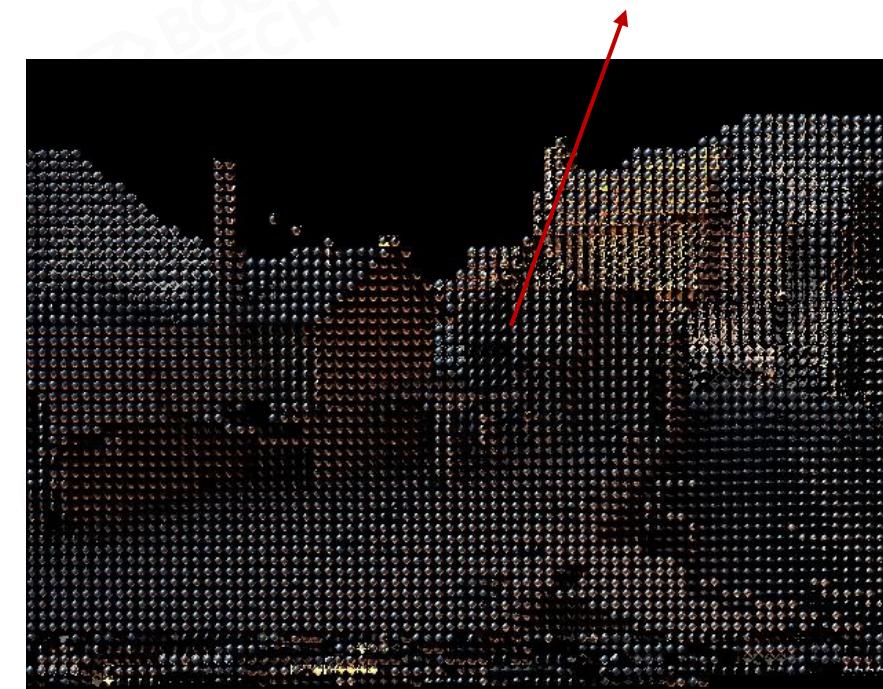
Radiance and HitDistance in 2d atlas



Hit Distance



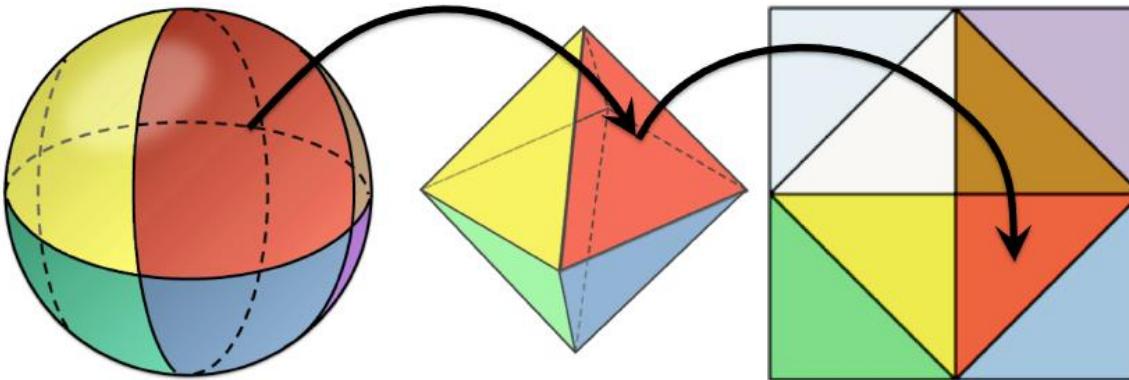
8x8



Radiance Altas



## Octahedron mapping



```
float2 unitVectorToOctahedron(float3 N)
{
    N.xy /= dot(1, abs(N));
    if (N.z <= 0)
    {
        x_factor = N.x >= 0? 1.0 : -1.0;
        y_factor = N.y >= 0 ? 1.0 : -1.0;
        N.xy = (1 - abs(N.yx)) * float2(x_factor, y_factor);
    }
    return float2(N.xy);
```

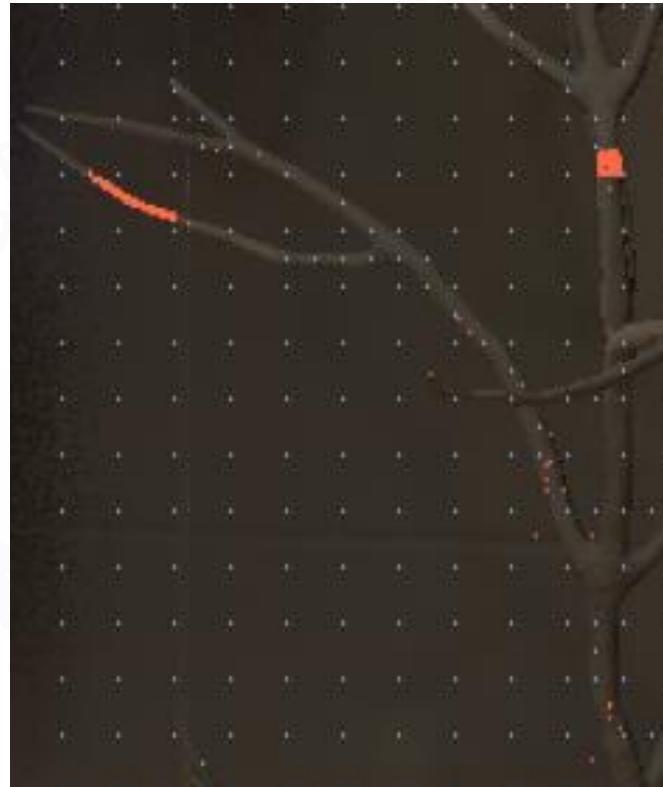


## Screen Probe Placement

- Adaptive placement with Hierarchical Refinement
- Iteratively place where interpolation fails



16x16



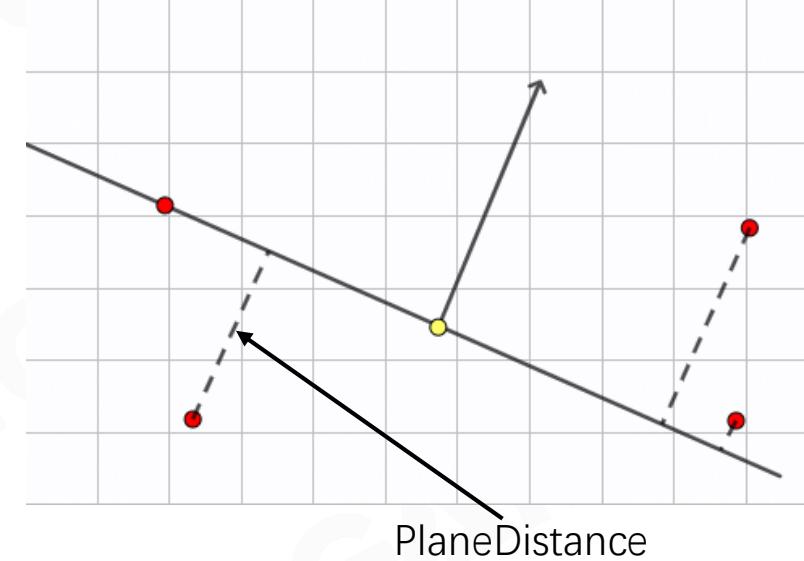
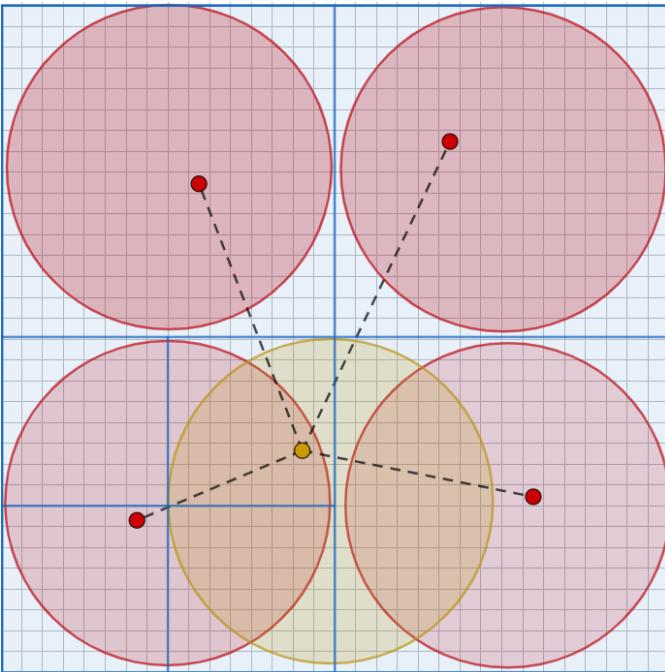
8x8



4x4



## Plane distance weighting of Probe Interpolation





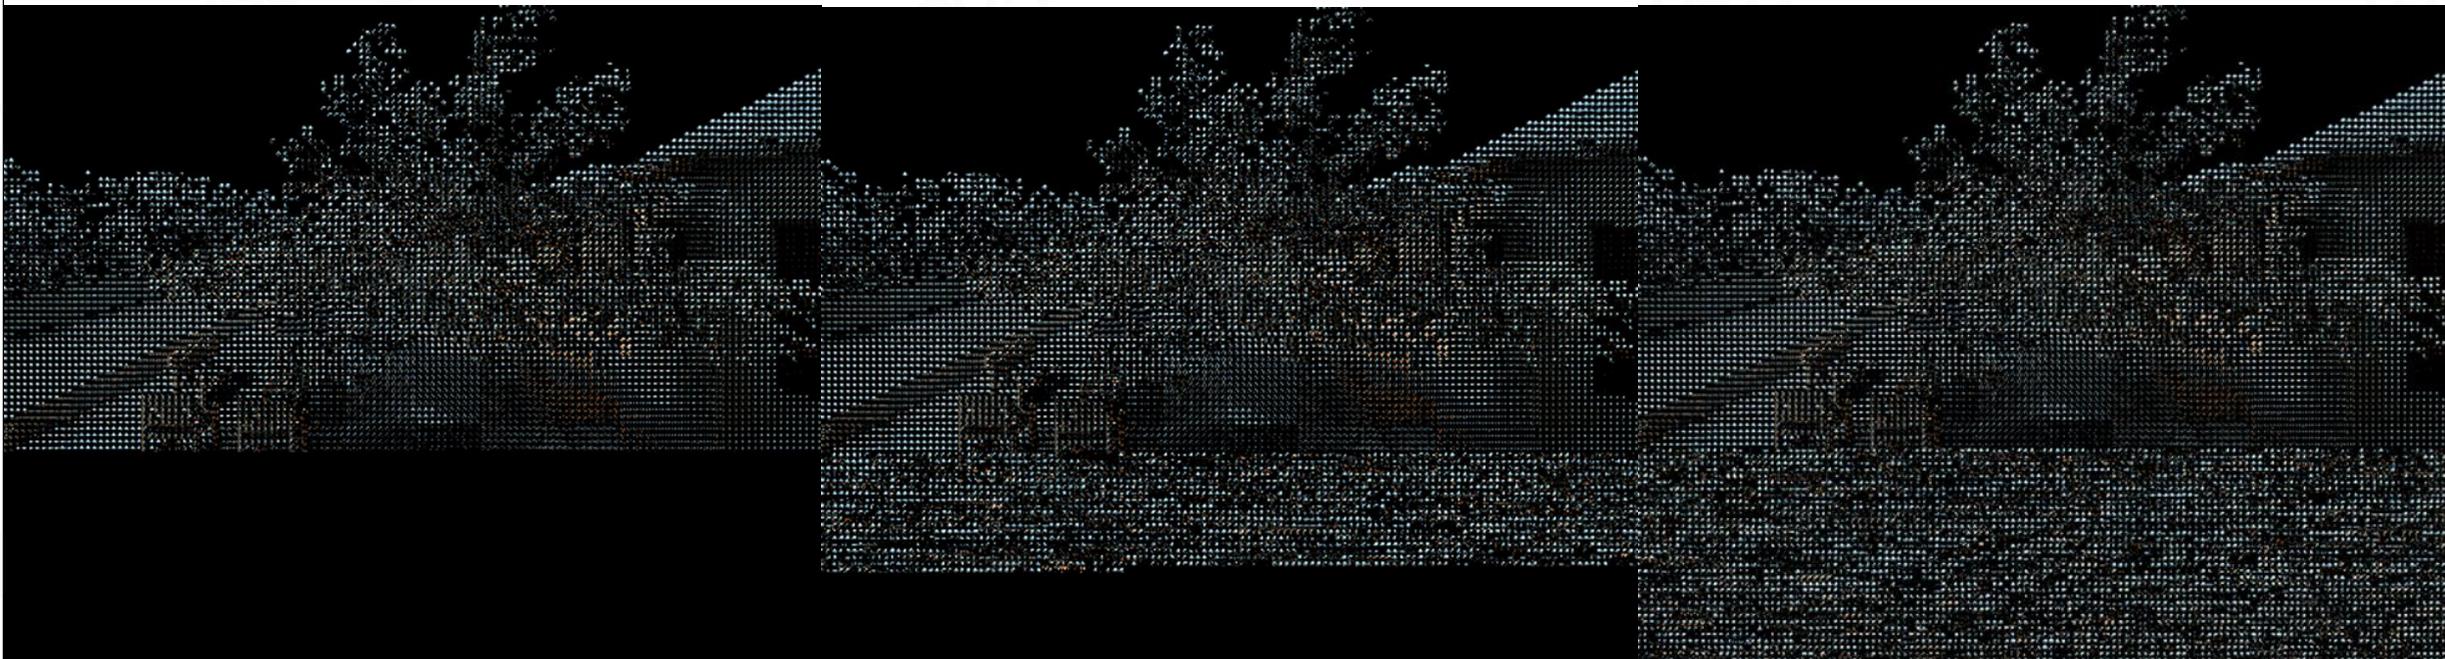
## Detect Non-Interpolatable Cases

```
·float4·PlaneDistances;  
·PlaneDistances. x ·=· abs (dot (float4 (Position00, ·-1), ·ScenePlane));  
·PlaneDistances. y ·=· abs (dot (float4 (Position10, ·-1), ·ScenePlane));  
·PlaneDistances. z ·=· abs (dot (float4 (Position01, ·-1), ·ScenePlane));  
·PlaneDistances. w ·=· abs (dot (float4 (Position11, ·-1), ·ScenePlane));  
  
·float4·RelativeDepthDifference ·=· PlaneDistances ·/· SceneDepth;  
  
·DepthWeights ·=· CornerDepths ·>· 0 ·?· exp2 (-10000. 0f ·*· (RelativeDepthDifference ·*· RelativeDepthDifference)) ·::· 0;  
  
InterpolationWeights ·=· float4 (  
    .... (1 ·--· BilinearWeights. y) ·*· (1 ·--· BilinearWeights. x),  
    .... (1 ·--· BilinearWeights. y) ·*· BilinearWeights. x,  
    .... BilinearWeights. y ·*· (1 ·--· BilinearWeights. x),  
    .... BilinearWeights. y ·*· BilinearWeights. x);  
  
InterpolationWeights ·*=· DepthWeights;  
  
float·Epsilon ·=· .01f;  
ScreenProbeSample. Weights ·/=· max (dot (ScreenProbeSample. Weights, ·1), ·Epsilon);  
  
float·LightingIsValid ·=· (dot (ScreenProbeSample. Weights, ·1) ·<· 1. 0f ·--· Epsilon) ·?· 0. 0f ·::· 1. 0f;
```



## Screen Probe Atlas

- Atlas have upper limit for real-time
- Place adaptive probes at the bottom of the atlas



16x16

8x8

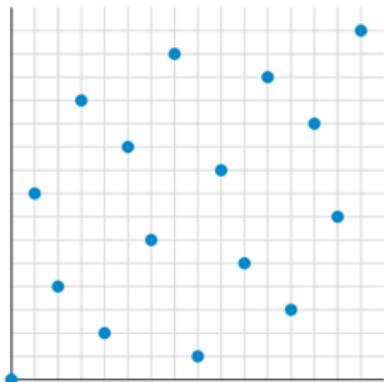
4x4



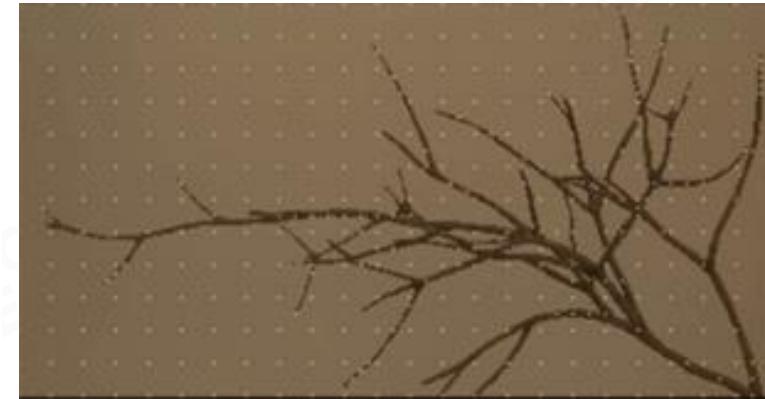
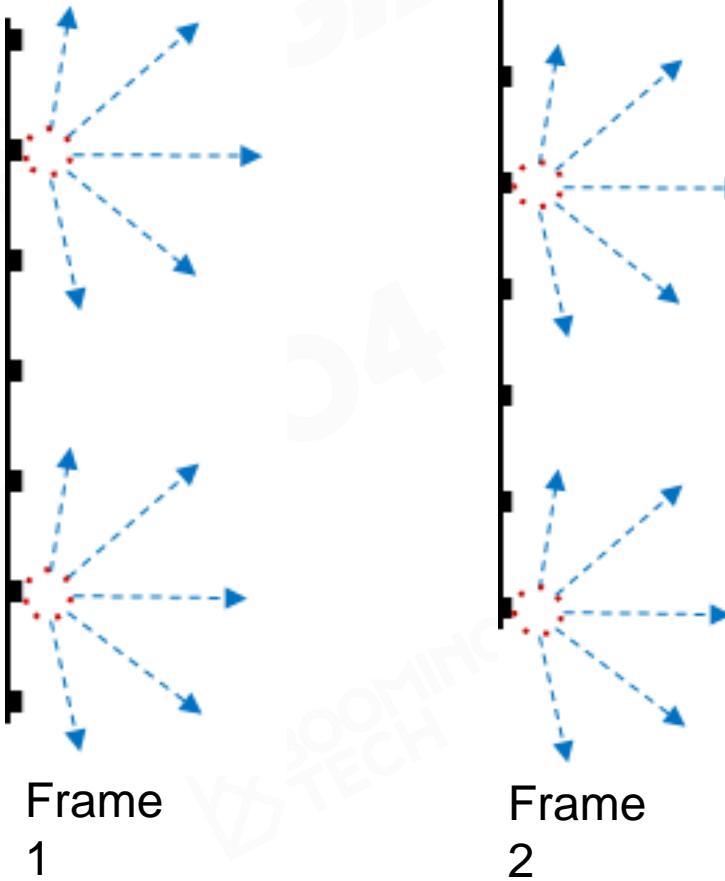


## Screen Probe Jitter

- Place probe directly on pixels
- Temporally jitter placement and direction
- Use Hammersley points in  $[0, 15]$



Hammersley Points in [0-15]





## Importance Sampling

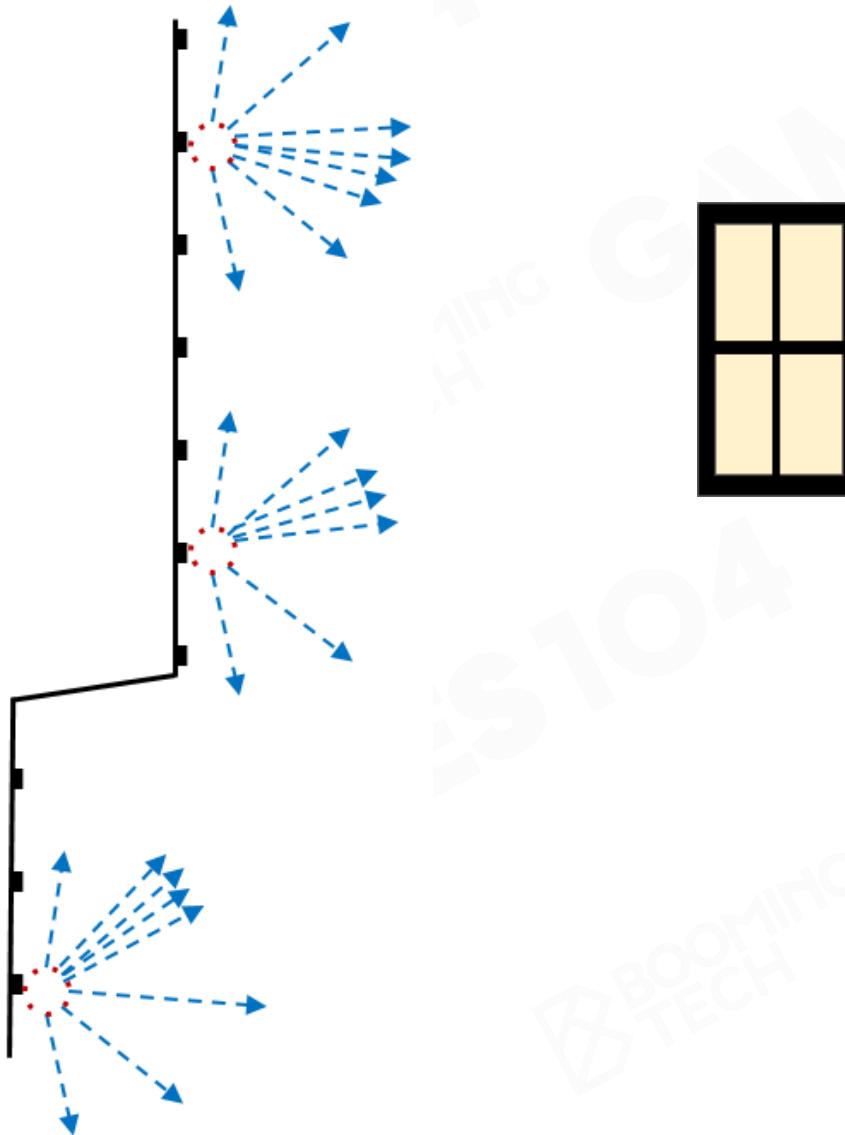


But too much noise at  $\frac{1}{2}$  ray per pixel





Better sampling - importance sample incoming lighting and BRDF





## Importance Sampling

$$\lim_{N \rightarrow \infty} \frac{1}{N} \sum_{k=1}^N \frac{L_i(l) f_s(l \rightarrow v) \cos(\theta l)}{P_k}$$

We would like to distribute rays proportional to the integrand  
How can we estimate these?

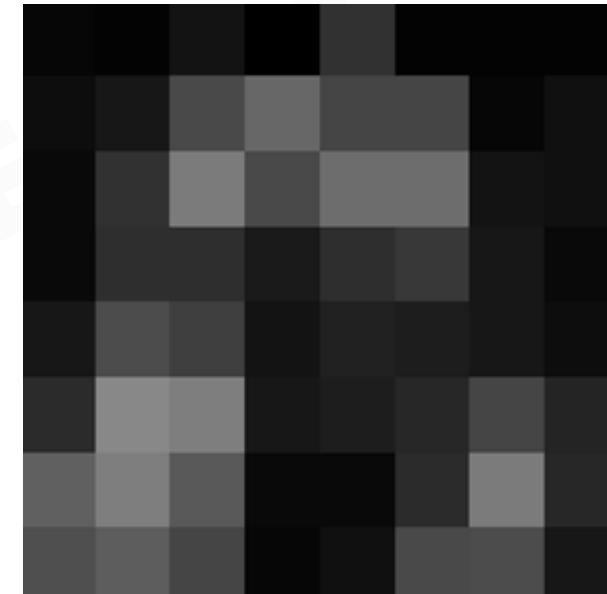


## Approximate Radiance Importance from Last Frame Probes

$$\lim_{N \rightarrow \infty} \frac{1}{N} \sum_{k=1}^N \frac{L_i(l) f_s(l \rightarrow v) \cos(\theta l)}{P_k}$$

### Incoming Radiance:

- Reproject to last frame and average the four neighboring Screen Probes Radiance
- No need to do an expensive search, as rays already indexed in octahedral atlas
- Fallback to World Space Probe Radiance if neighboring probes are occluded



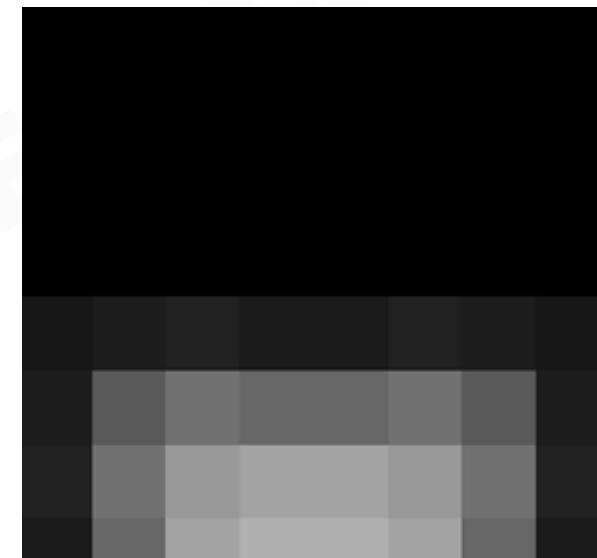


## Accumulate Normal Distribution Nearby

$$\lim_{N \rightarrow \infty} \frac{1}{N} \sum_{k=1}^N \frac{L_i(l) f_s(l \rightarrow v) \cos(\theta l)}{P_k}$$

### BRDF:

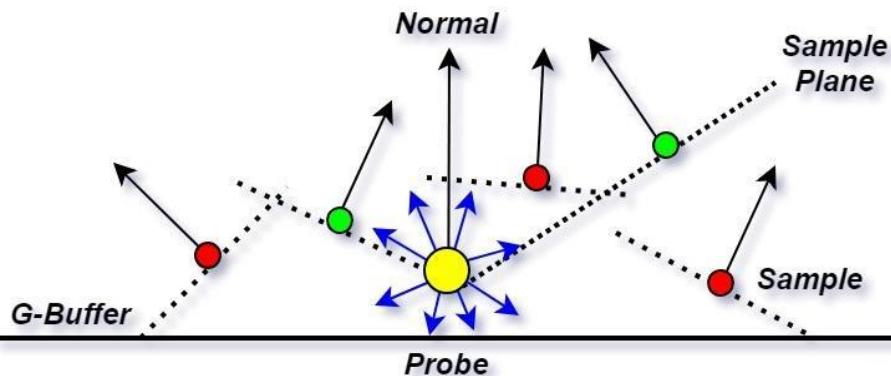
- For a probe that's placed on a flat wall, about half of its sphere having a zero BRDF
- Accumulate from pixels that will use this Screen Probe





# Nearby Normal Accumulation

- Gather 64 neighbor pixels around current probe's pixel in a 32x32 pixel range
- Accept pixel if its depth weight > 0.1
- Accumulate these pixels' world normal into SH



```
float3 PixelPosition = GetWorldPositionFromScreenUV(PixelScreenUV, Material.SceneDepth);
float4 PixelPlane = float4(Material.WorldNormal, dot(Material.WorldNormal, PixelPosition));
float3 ProbeWorldPosition = GetWorldPositionFromScreenUV(ScreenUV, ProbeSceneDepth);

float PlaneDistance = abs(dot(float4(ProbeWorldPosition, -1), PixelPlane));
float RelativeDepthDifference = PlaneDistance / ProbeSceneDepth;
float DepthWeight = exp2(-10000.0f * (RelativeDepthDifference * RelativeDepthDifference));
```

```
if (DepthWeight > .1f || bCenterSample)
{
    uint Index;
    InterlockedAdd(NumSphericalHarmonics, 1, Index);

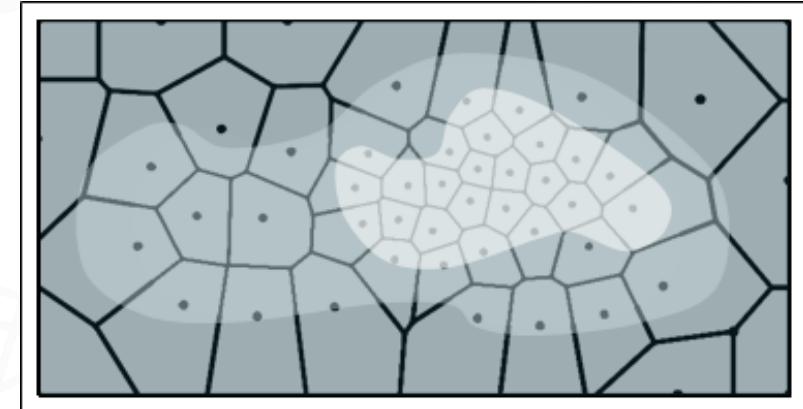
    FThreeBandSHVector BRDF;
    if (HasSphericalVisibility(Material))
    {
        // Avoid culling directions that the shading models will sample
        BRDF = (FThreeBandSHVector)0;
        BRDF.V0.x = 1.0f;
    }
    else
    {
        BRDF = CalcDiffuseTransferSH3(Material.WorldNormal, 1.0f);
    }
    WriteGroupSharedSH(BRDF, Index);
}
```



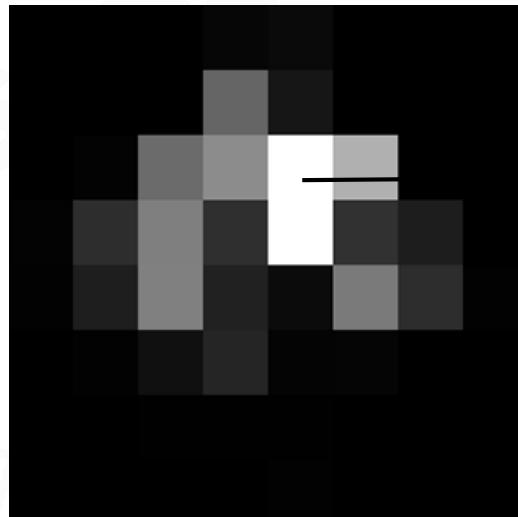
## Structured Importance Sampling

- Assigns a small number of samples to hierarchically structured areas of the Probability Density Function (PDF)
- Achieves good global stratification
- Sample placement requires offline algorithm

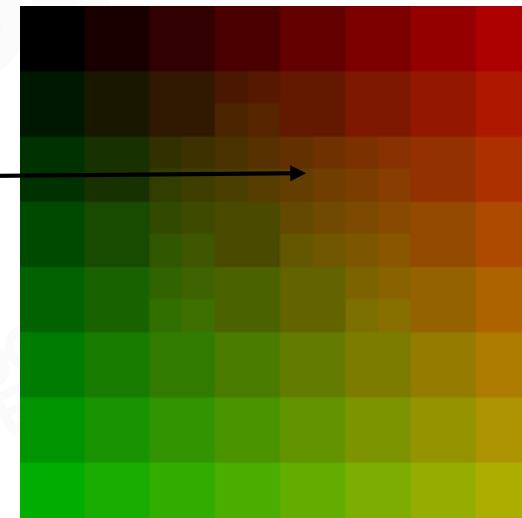
Maps perfectly to Octahedral mip quadtree!



PDF



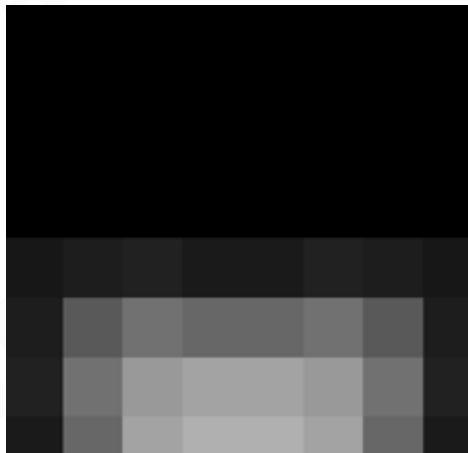
Subdivided Octahedral texels



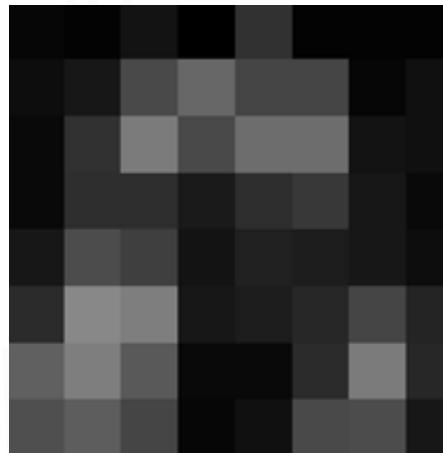


# Fix Budget Importance Sampling based on Lighting and BRDF

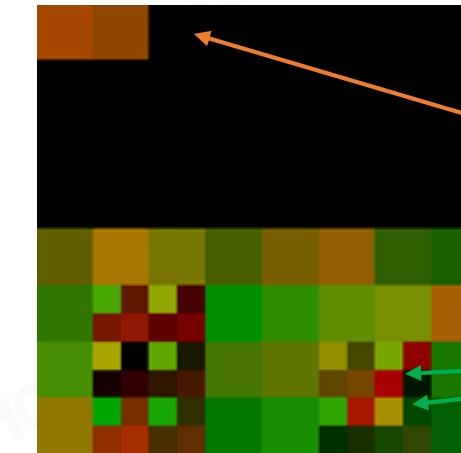
- Start with uniformly distributed probe ray directions
- Fixed probe tracing ray count=64
- Calculate BRDF PDF \* Lighting PDF for each Octahedral texel
- Sort rays by PDF from low to high
- For every 3 rays with PDF below cull threshold, supersample the matching highest PDF ray



BRDF PDF



Lighting PDF



Culled directions

Supersampled  
directions



uniform ray directions



ray directions after importance sampling







## Denoising and Spatial Probe Filtering



## Denoise: Spatial filtering for Probe

Large spatial filter for cheap

- Each probe cover 16x16 pixels, 3x3 filtering kernel in probe space equals 48x48 in screen space

Can ignore normal differences between spatial neighbors

- Only depth weighting

```
float GetFilterPositionWeight(float ProbeDepth, float SceneDepth)
{
    float DepthDifference = abs(ProbeDepth - SceneDepth);
    float RelativeDepthDifference = DepthDifference / SceneDepth;
    return ProbeDepth >= 0 ? exp2(-SpatialFilterPositionWeightScale * (RelativeDepthDifference * RelativeDepthDifference)) : 0;
}
```



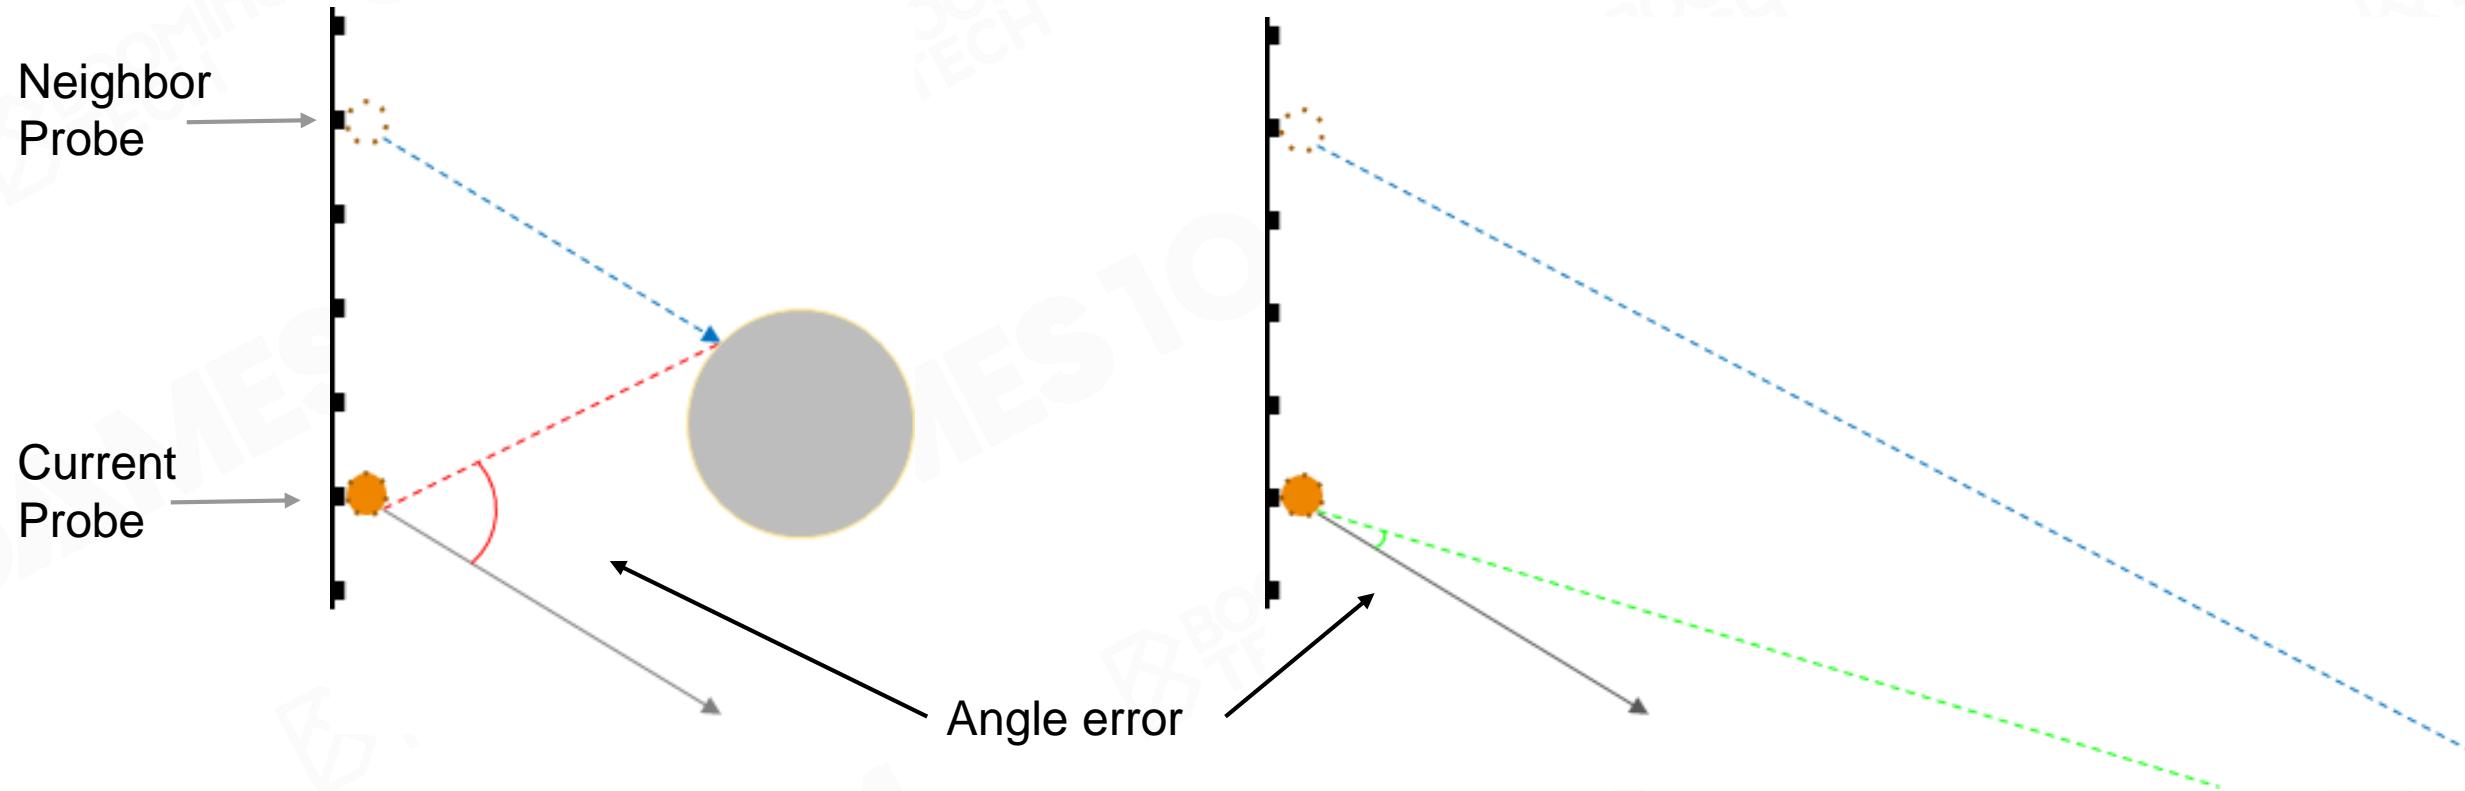


## Denoise: Gather Radiance from neighbors

Gather radiance from matching Octahedral cell in neighbor probes

Error weighting:

- Angle error from reprojected neighbor ray hits (less than 10 degree)
- Filters distant lighting, preserves local shadowing





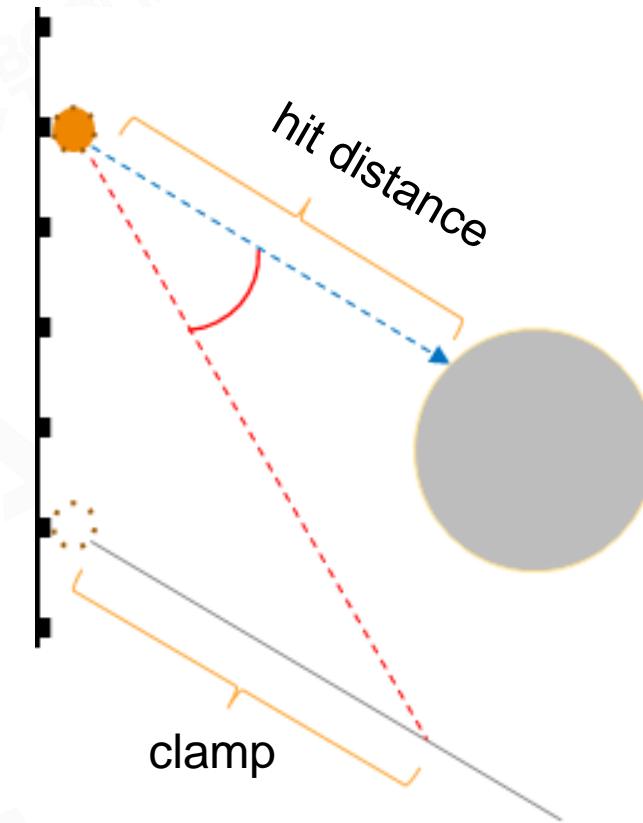
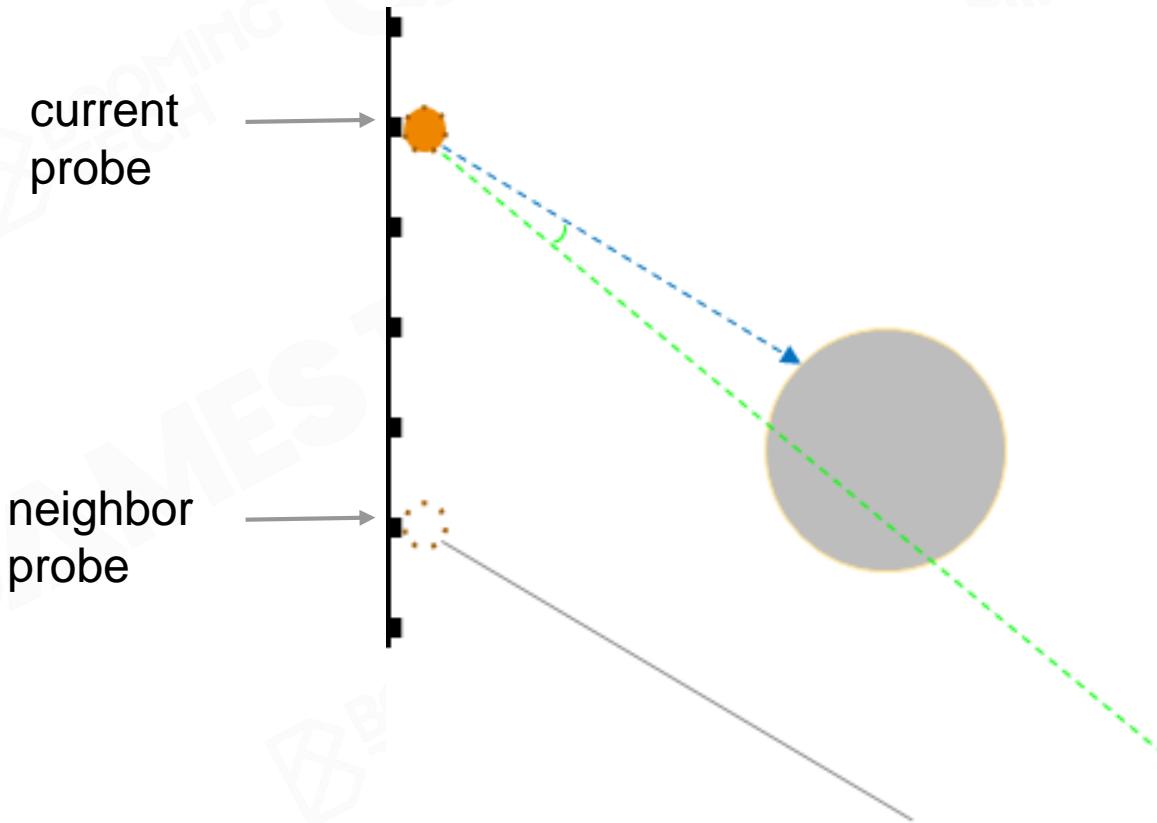


# Clamp Distance Mismatching

Angle error biases toward distant light = leaking

- Distant light has no parallax and never gets rejected

Solution: clamp neighbor hit distance to our own before reprojection







## Final filtering compare





## World Space Probes and Ray Connecting



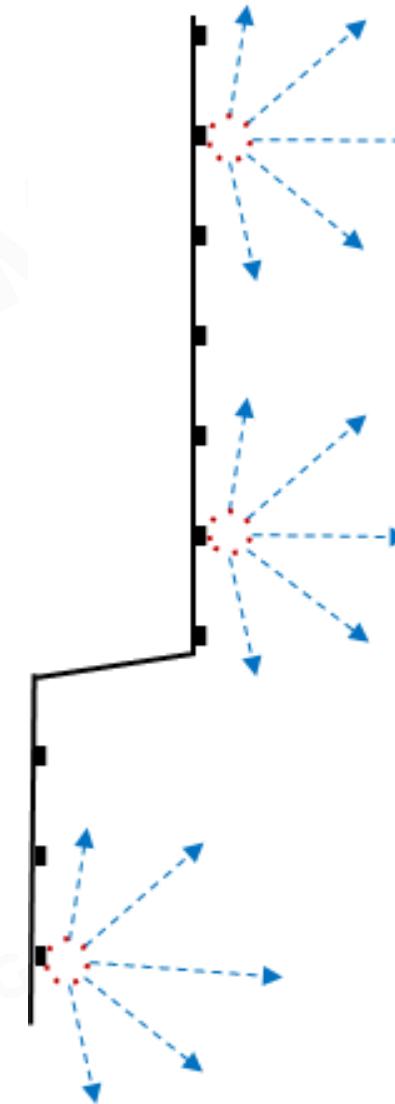
## World Space Radiance Cache

### Problem: distant lighting

- Noise from small bright feature increases with distance
- Long incoherent traces are slow
- Distant lighting is changing slowly - opportunity to cache
- Redundant operations for nearby Screen Probes

### Solution: separate sampling for distant Radiance

- World space Radiance Caching for distant lighting
- Stable error since world space - easy to hide



Screen Radiance Cache



World Radiance Cache



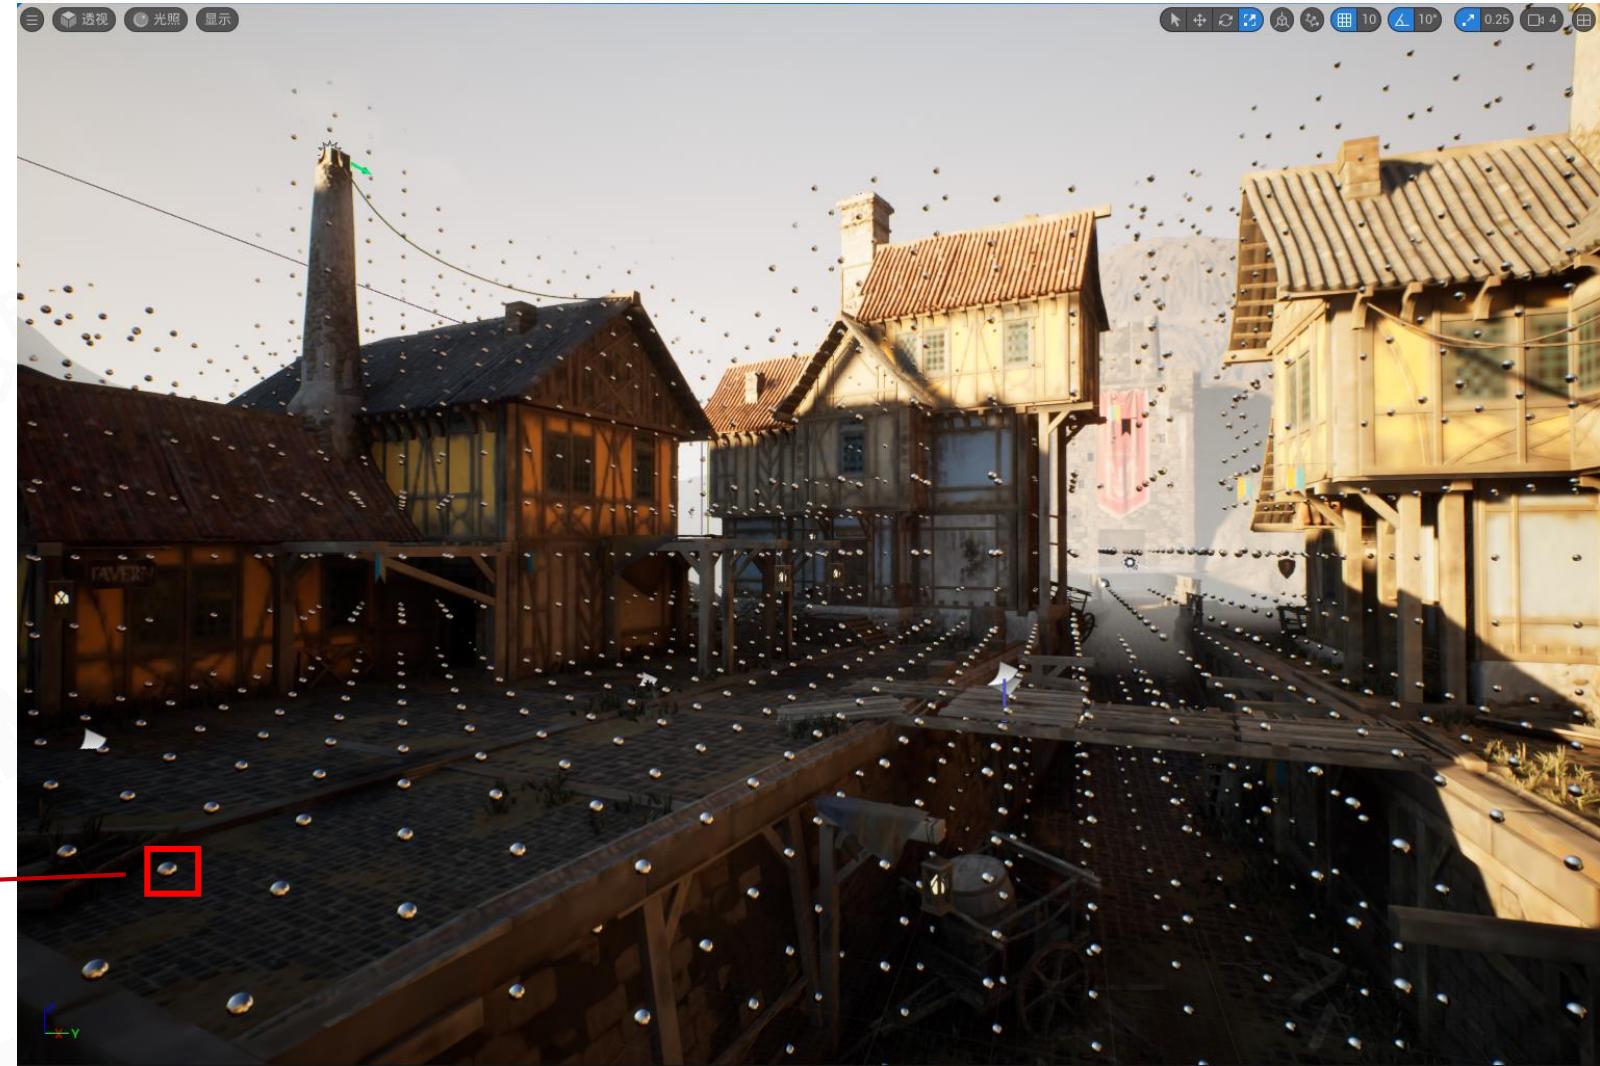
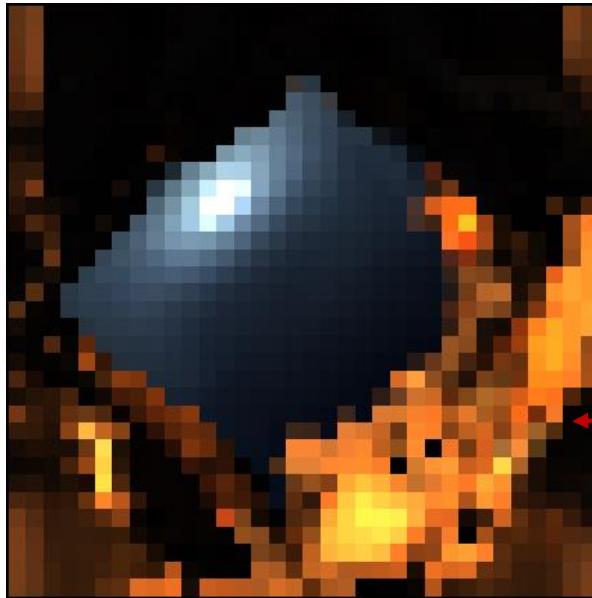
## World Space Radiance Cache

### Placement

- 4 level clipmaps around camera
- default resolution is  $48^3$
- clipmap 0 size is  $50m^3$

### Radiance

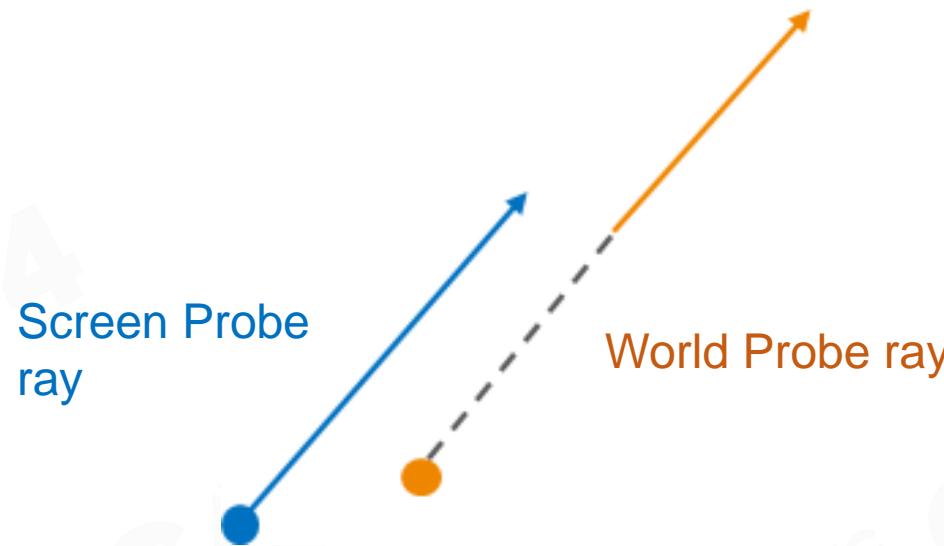
- $32 \times 32$  atlas a per probe





## Connecting rays

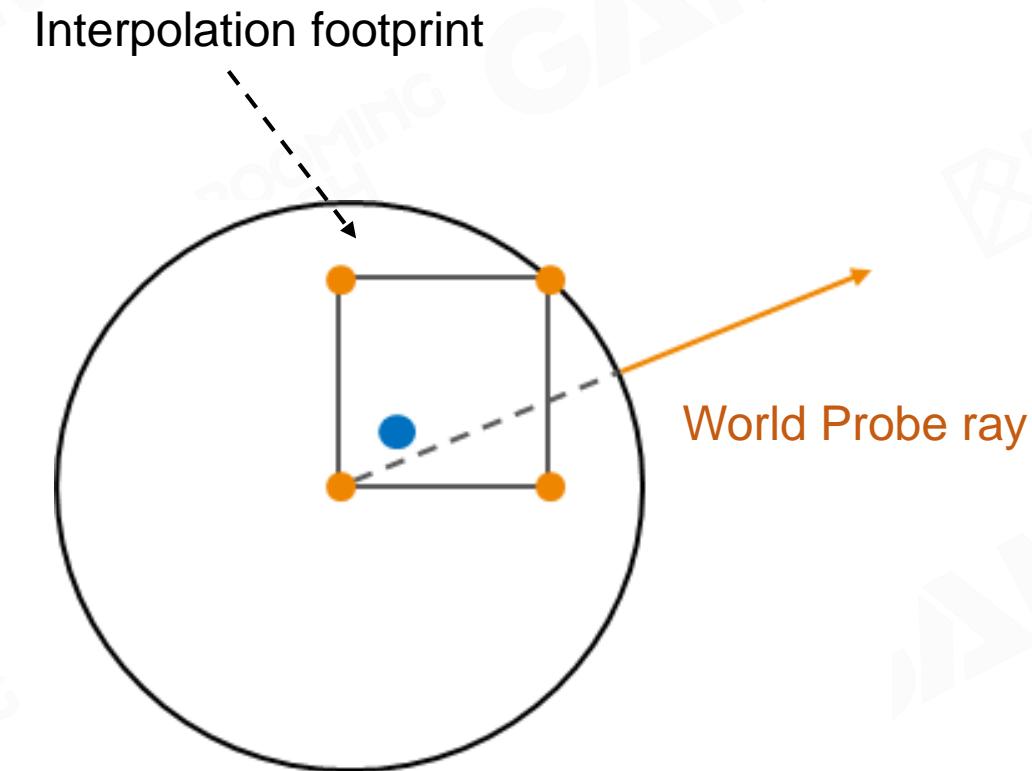
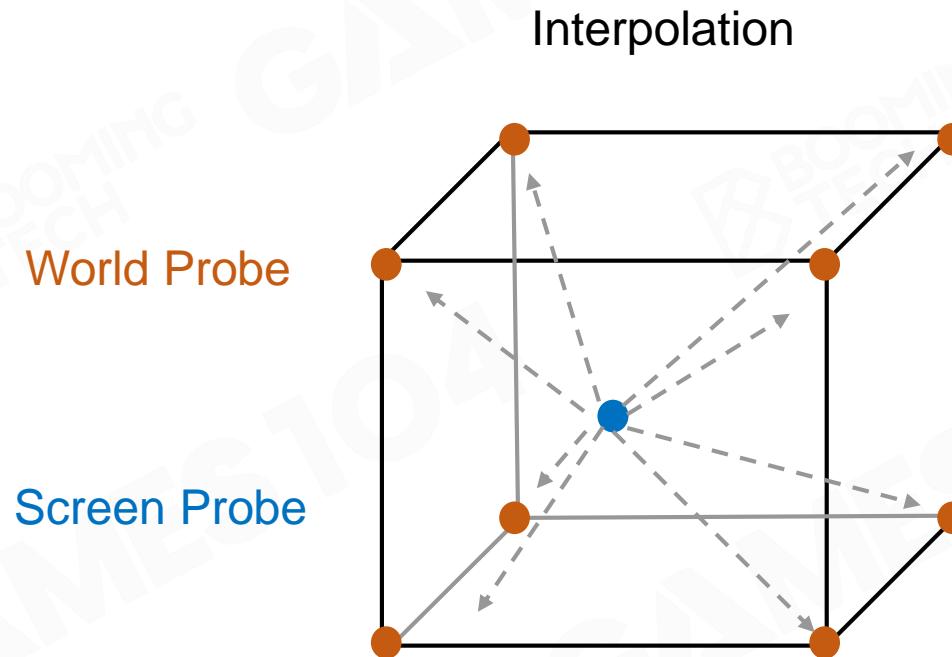
- How to connect **Screen Probe ray** and **World Probe ray**





## Connecting rays

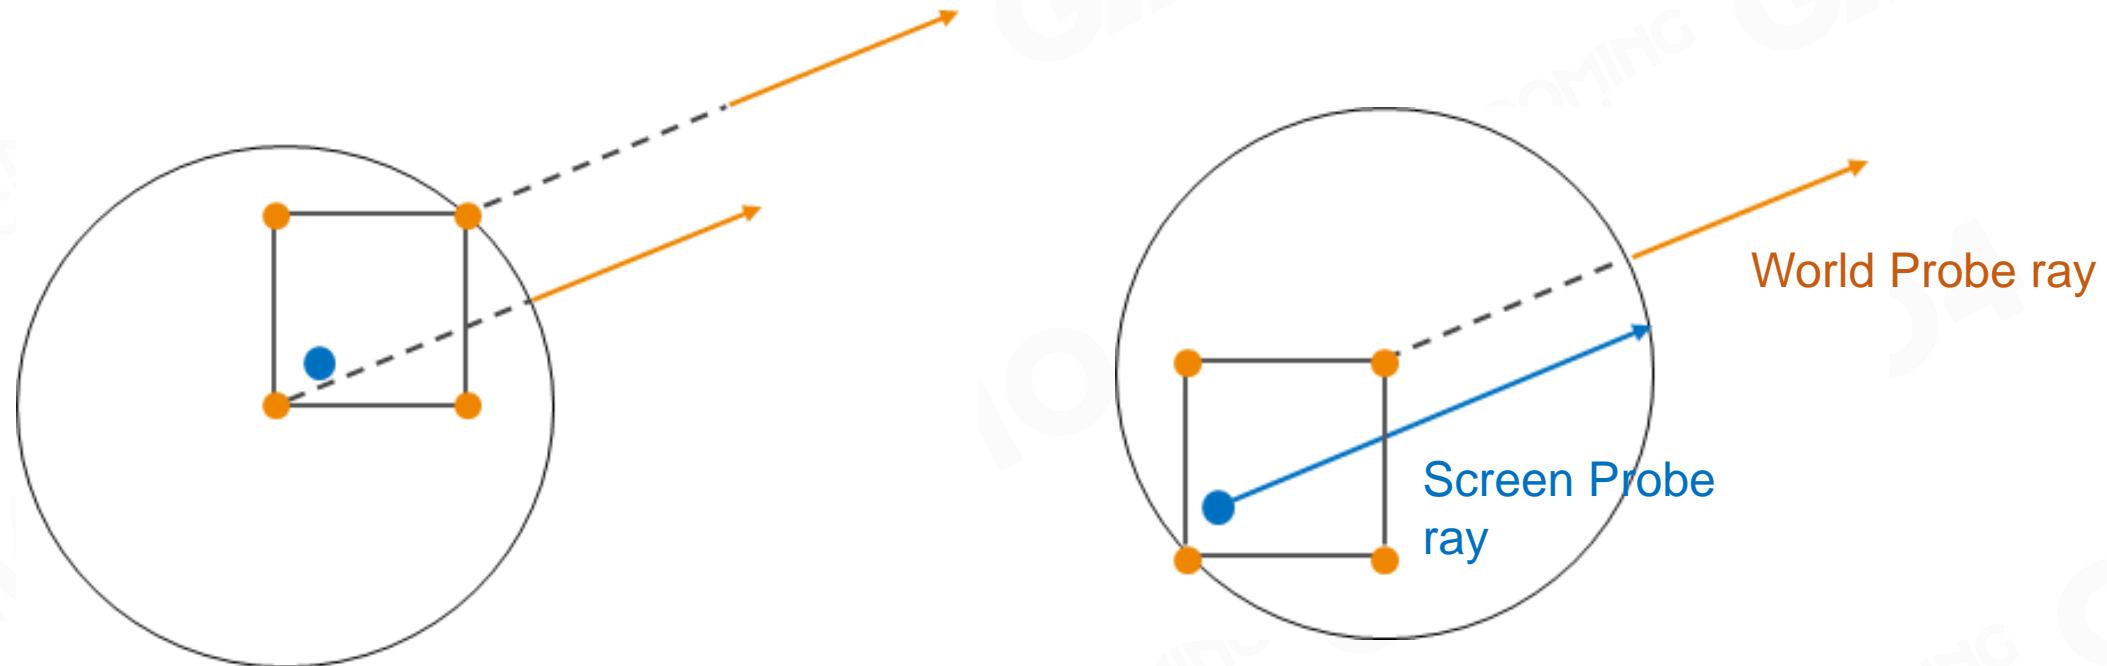
- World Probe ray must skip the interpolation footprint





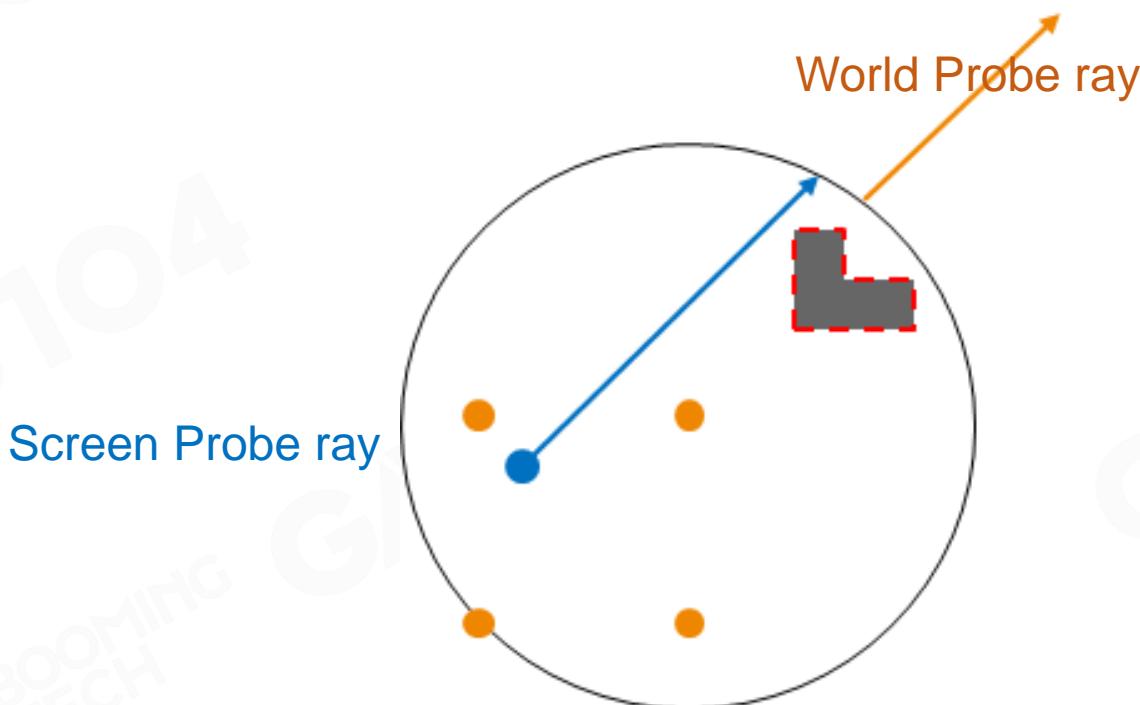
## Connecting rays

- Screen Probe ray must cover interpolation footprint + skipped distance



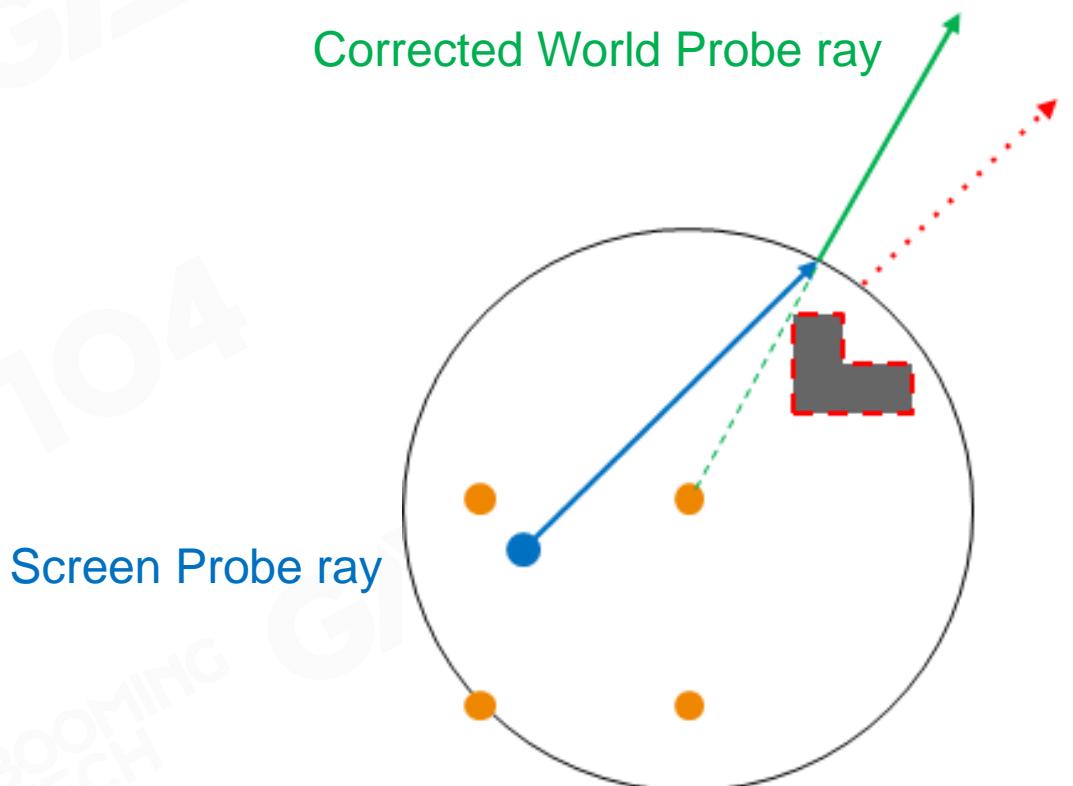


- Problem: leaking!
- World probe radiance should have been occluded
  - But wasn't due to incorrect parallax





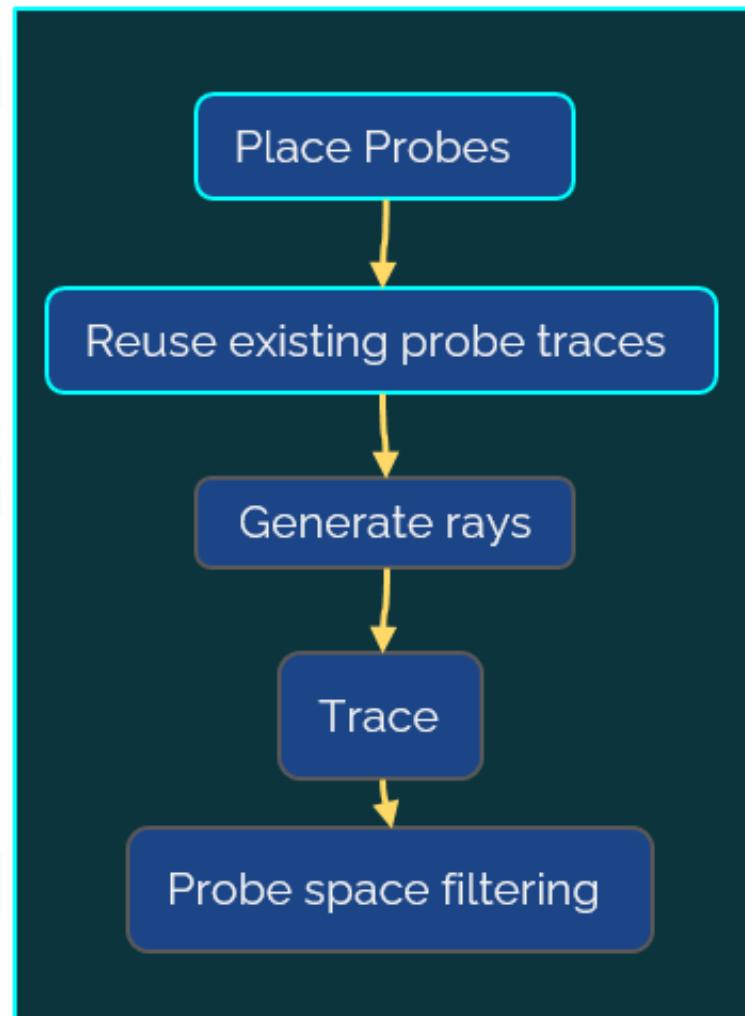
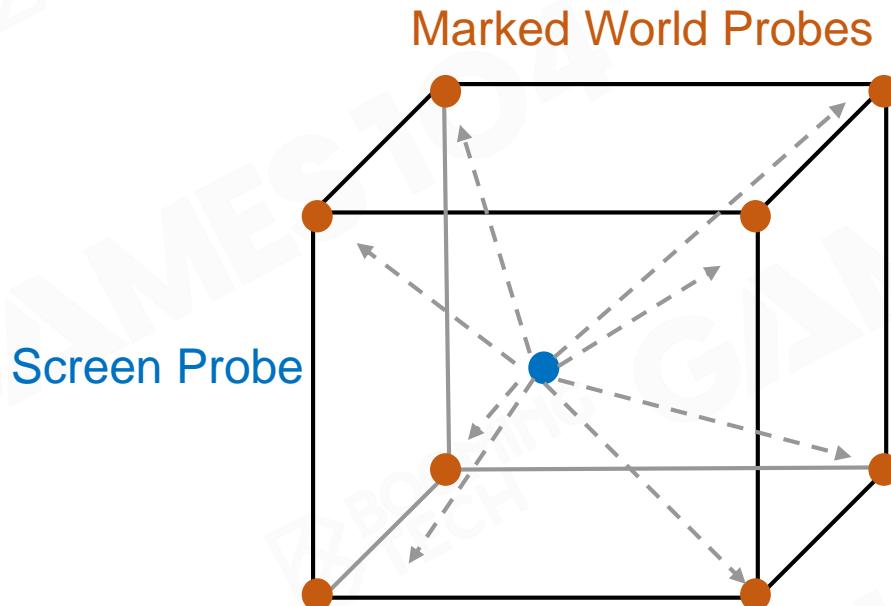
- Solution: simple sphere parallax
- Reproject Screen Probe ray intersection with World Probe sphere





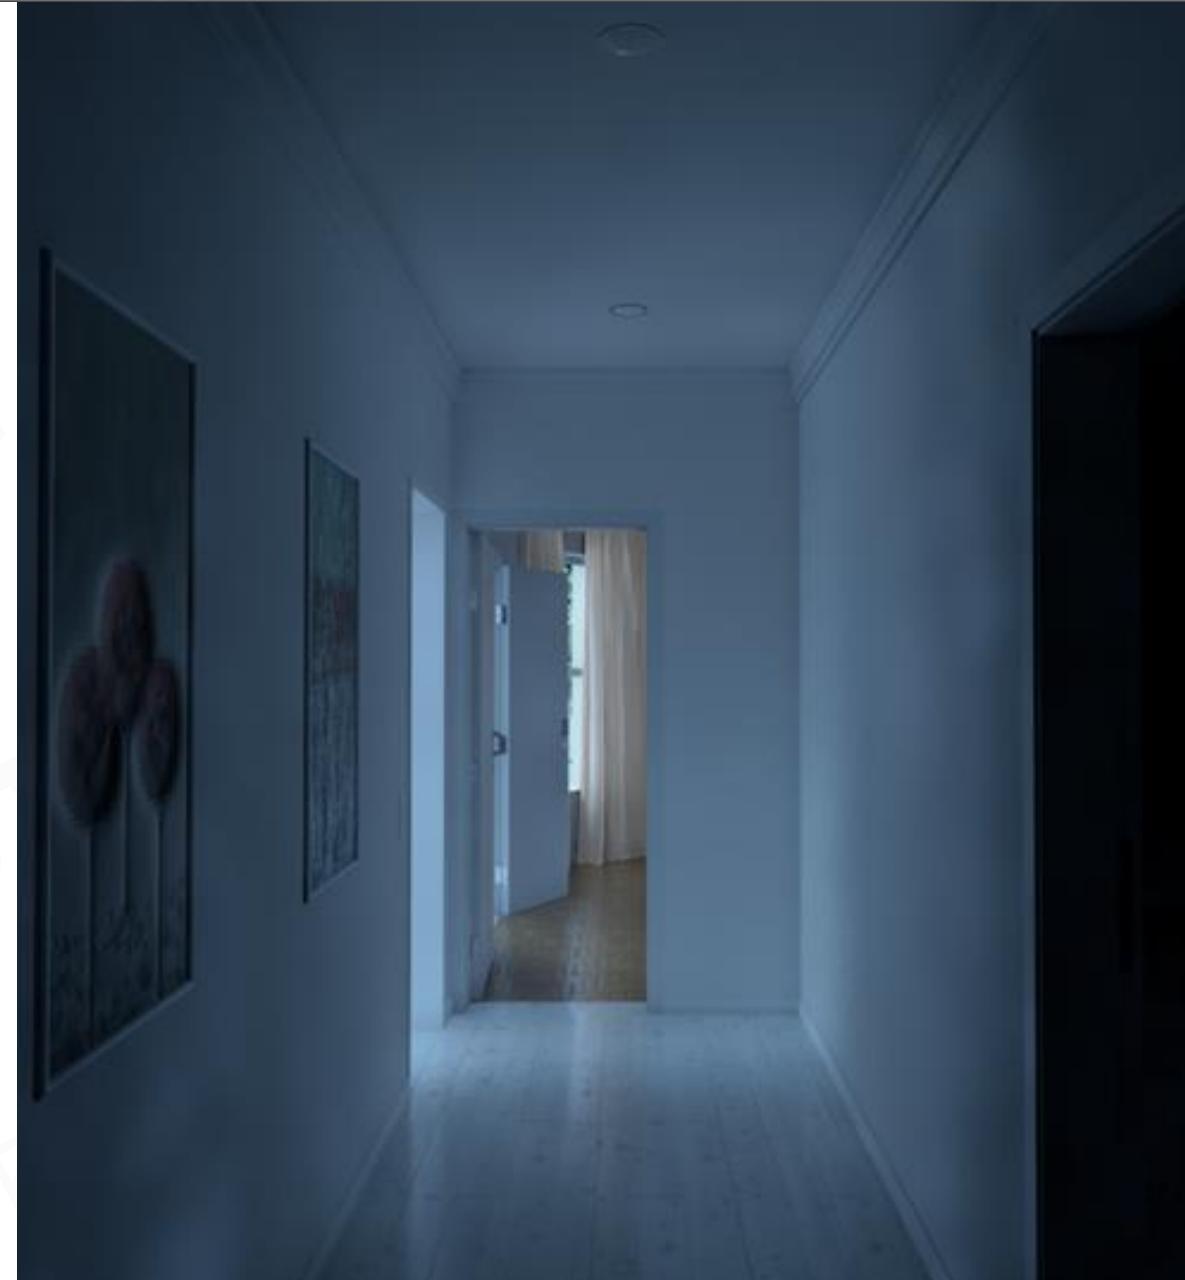
## Placement and caching

- Mark any position that we will interpolate from later in clipmap indirections
- For each marked world probe:
  - Reuse traces from last frame, or allocate new probe index
  - Re-trace a subset of cache hits to propagate lighting changes



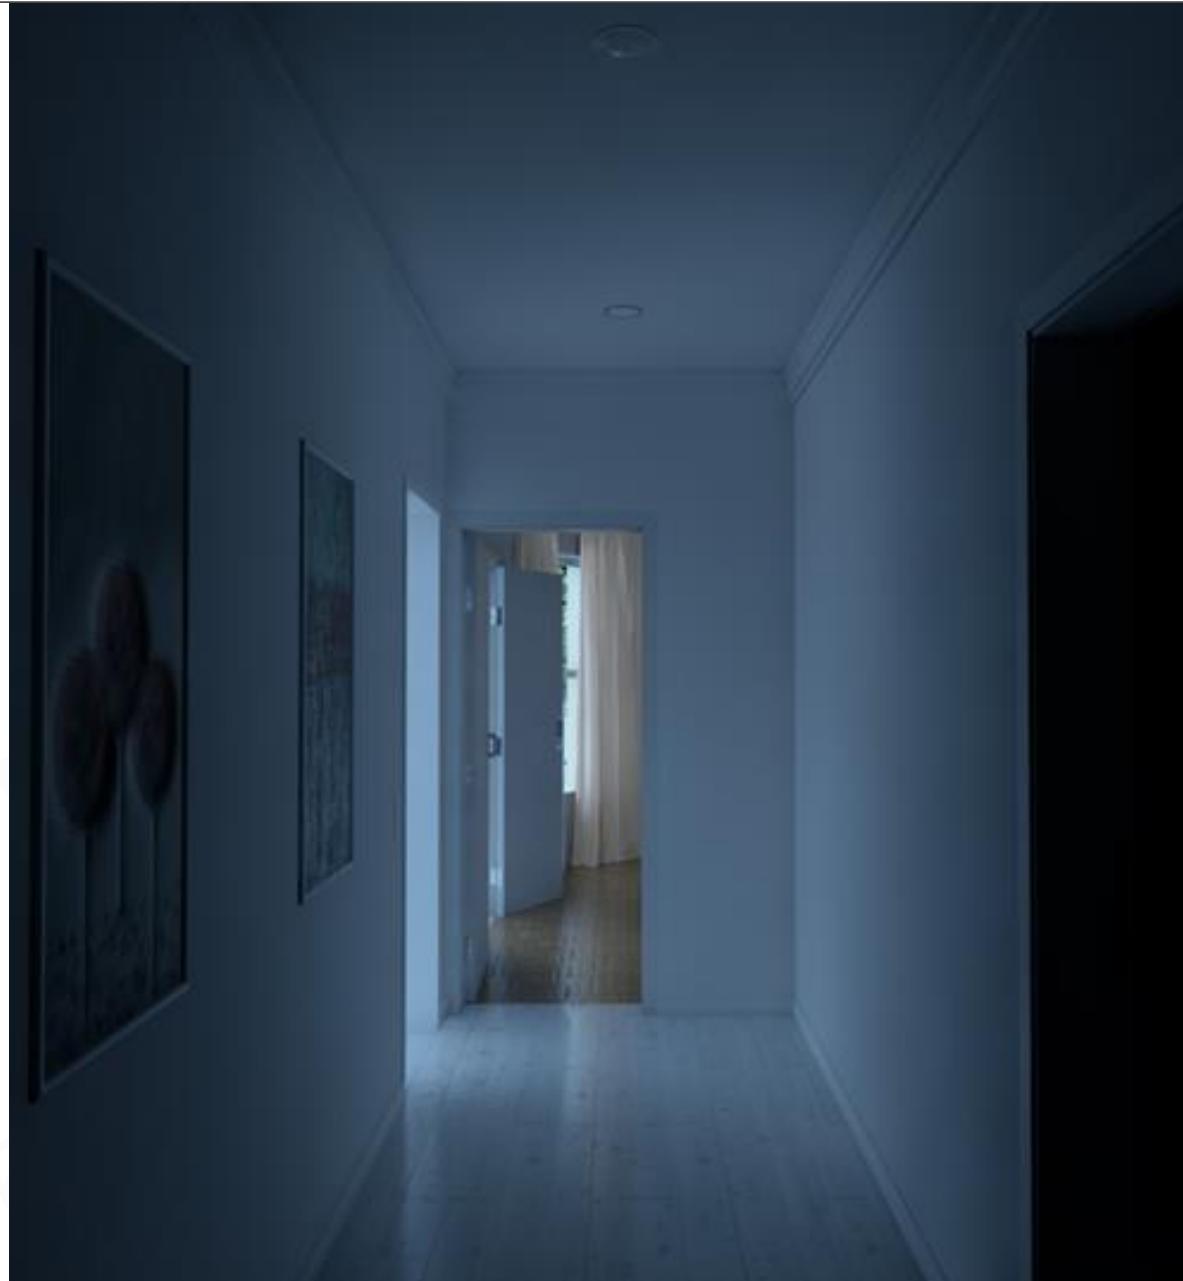


- Without World Space Probes





- Screen Radiance Cache for the first **2 meters**
- World Radiance Cache for any lighting further than that.





## Phase 4 : Shading Full Pixels with Screen Space Probes



## Convert Probe Radiance to 3rd order Spherical Harmonic:

- SH is calculated per Screen Probe
- Full res pixels load SH coherently
- SH Diffuse integration cheap and high quality

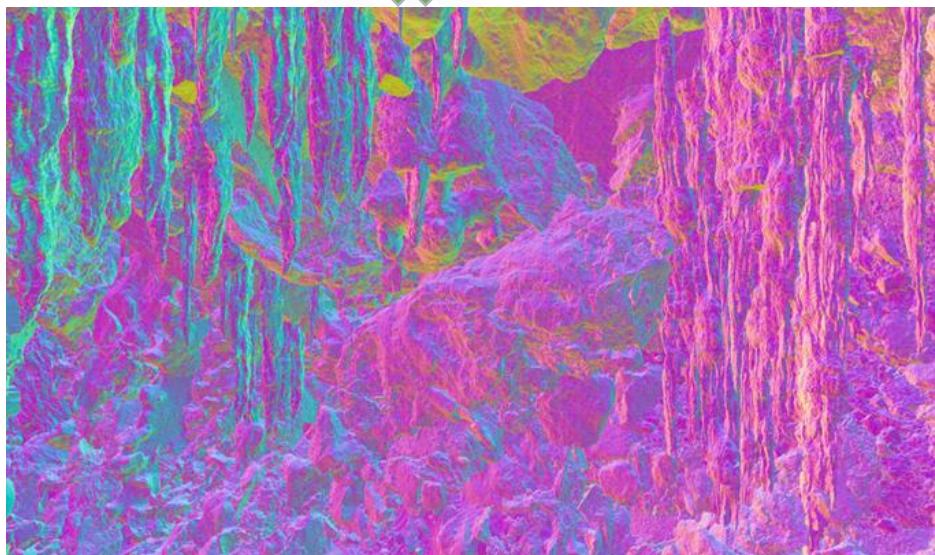
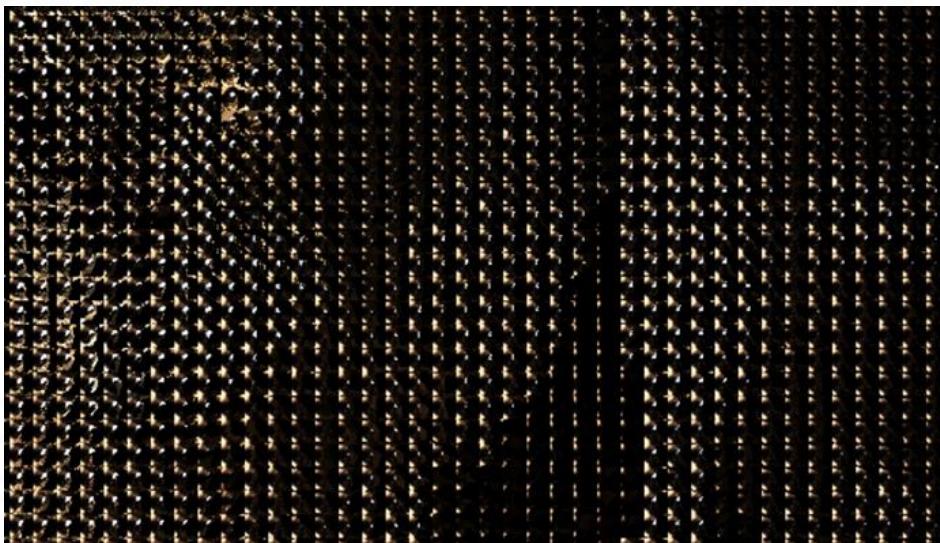
Spherical Harmonic

importance sample the BRDF  
to get ray directions, and then  
sample the Radiance Cache.





## Final integration with SH

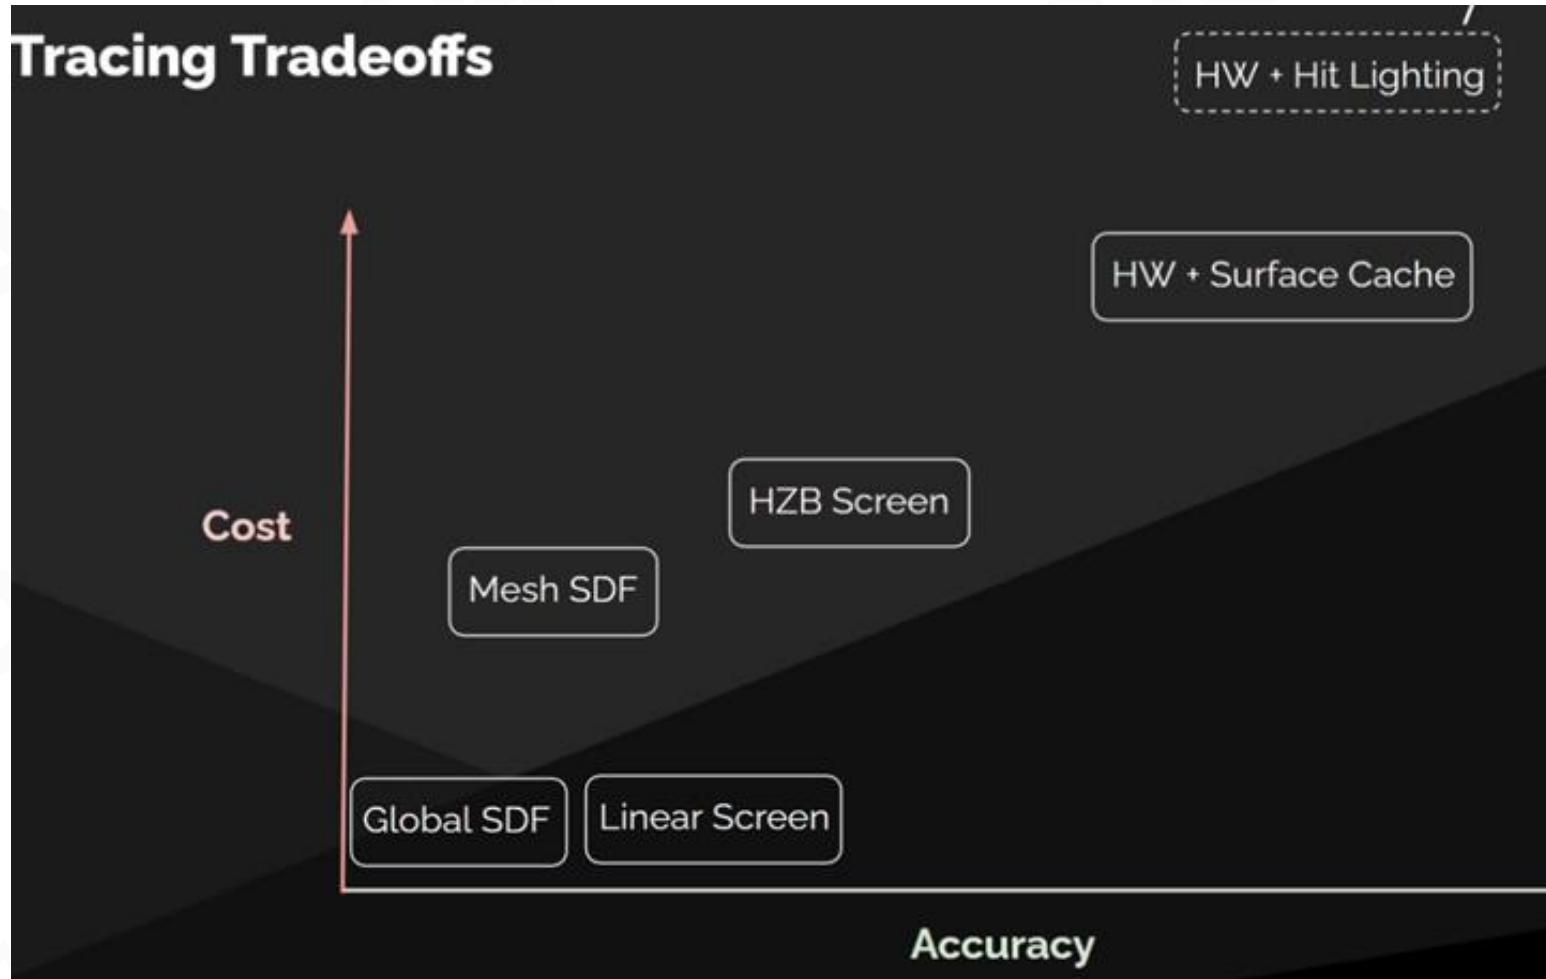




## Overall, Performance and Result



## Speed of Different Tracing Methods





Red – Screen Space Trace

fail to hit



Green – Mesh SDF Trace

fail to hit



Blue – Global SDF Trace





Trace Method	Trace Parameter	Hit Information	Sampling
Screen Space Trace	HZB Max Step=50	Hit Screen UV	Previous Frame SceneColorTexture
Mesh SDF Trace	Max Trace Distance= 1.8m Position in 40m Radius of Camera	Mesh ID Hit World Position Normal	Final Lighting
Global SDF Trace	Max Trace Distance = 200m	Hit World Position Normal	Voxel Lighting
Cubemap	Infinite	N/A	Sky Cube Color



SSGI OFF





SSGI ON





## Performance

Playstation 5

1080p internal resolution

Temporal Super Resolution to 4k

**½ ray per pixel**

**Total: 3.74ms**

Screen Bent Normal - .39ms

Place Probes - .13ms

Generate rays - .35ms

Trace - 1.07ms

Probe space filtering - .24ms

Interpolate and Integrate - .62ms

Temporal filter - .32ms

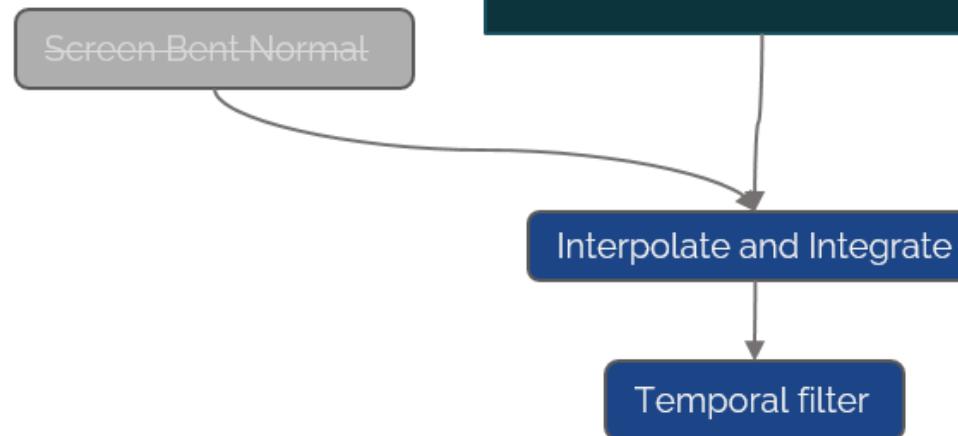
World Radiance  
Cache

.53ms



**1/8 ray per pixel**

**Total: 2.15ms**





$\frac{1}{8}$  ray per pixel  
2.15ms Final Gather





$\frac{1}{2}$  ray per pixel  
3.74ms Final Gather



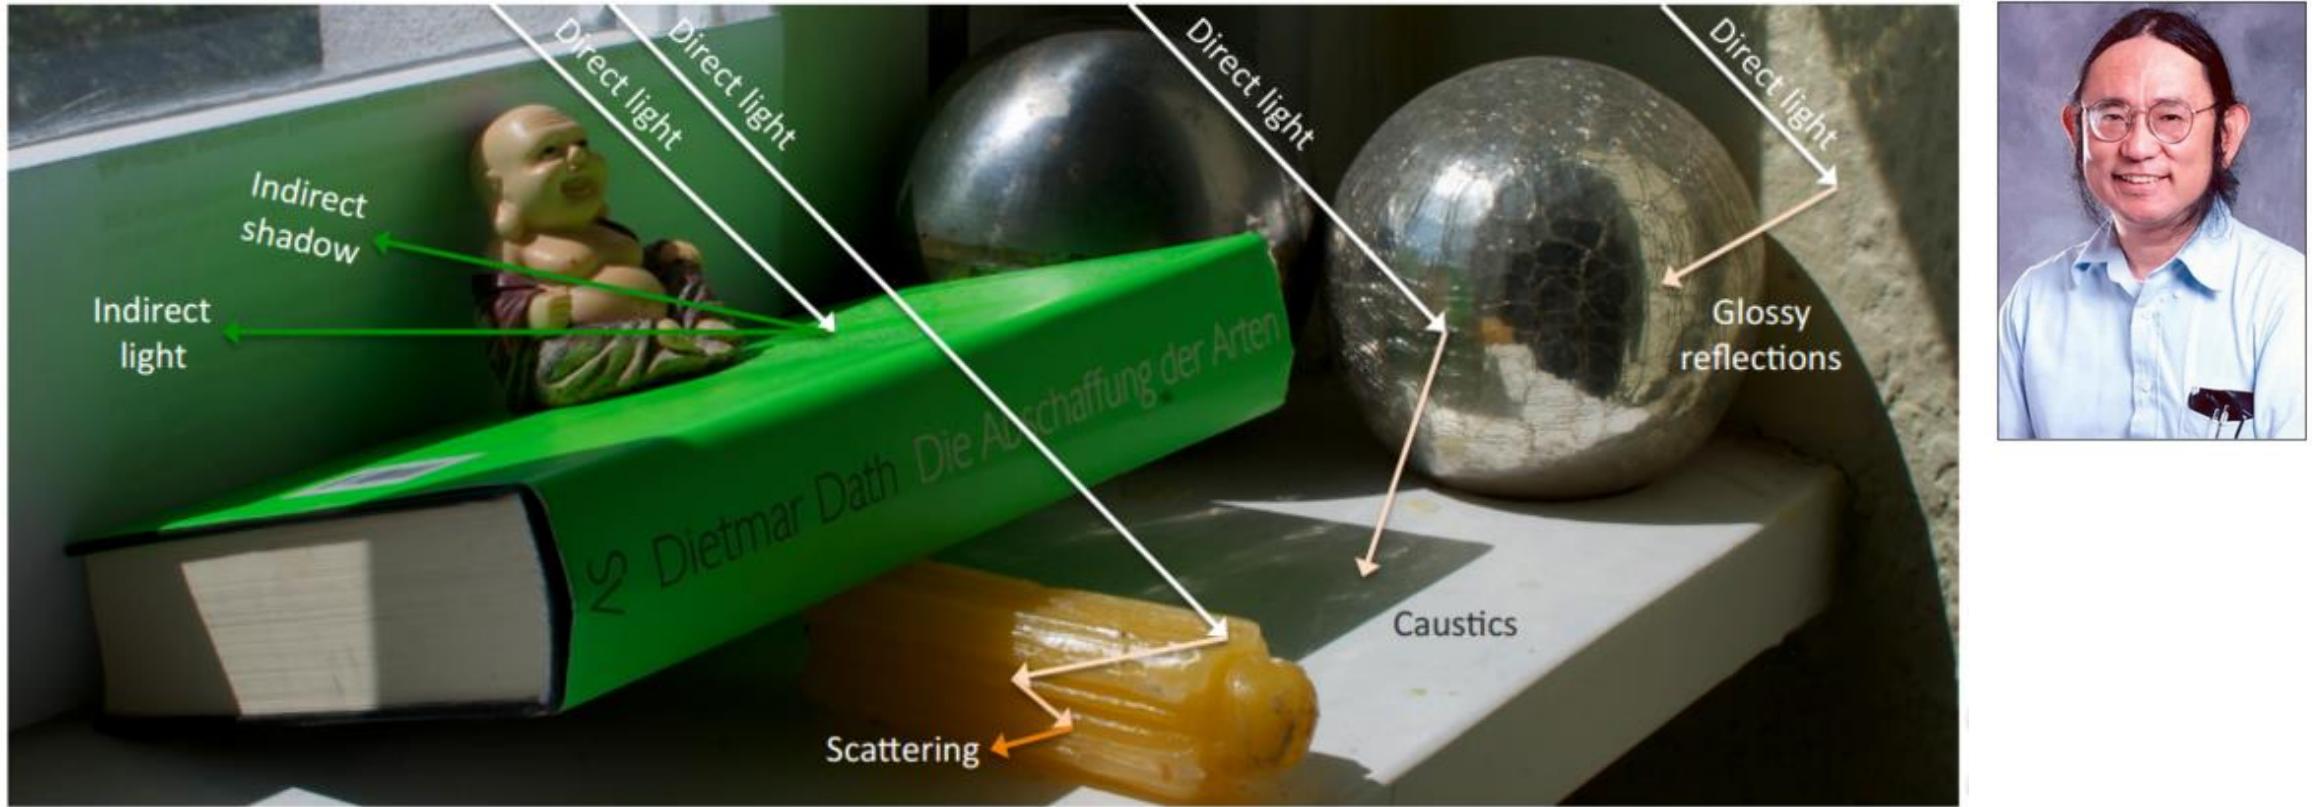






# Conclusion

# Complexity of Real Rendering



$$L_o(x, \omega_o) = L_e(x, \omega_o) + \int_{H^2} f_r(x, \omega_o, \omega_i) L_i(x, \omega_i) \cos \theta_i d\omega_i$$



# References



## Monte Carlo Integration

- Importance Sampling: <https://patapom.com/blog/Math/ImportanceSampling/>
- Notes on importance sampling: [https://www.tobias-franke.eu/log/2014/03/30/notes\\_on\\_importance\\_sampling.html](https://www.tobias-franke.eu/log/2014/03/30/notes_on_importance_sampling.html)
- Chinagraph 2020会前课程  
3:[https://www.bilibili.com/video/BV1my4y1z76s?p=3&vd\\_source=5e38e5c84aa6ff6cff3b802b3eecb8bc](https://www.bilibili.com/video/BV1my4y1z76s?p=3&vd_source=5e38e5c84aa6ff6cff3b802b3eecb8bc)
- Probability density function: [https://en.wikipedia.org/wiki/Probability\\_density\\_function](https://en.wikipedia.org/wiki/Probability_density_function)
- Microfacet Models for Refraction through Rough Surfaces:  
<https://www.cs.cornell.edu/~srm/publications/EGSR07-btdf.pdf>



## GI

- SSGI: <https://www.ea.com/frostbite/news/stochastic-screen-space-reflections>
- DDGI: <https://zhuanlan.zhihu.com/p/404520592>
- DDGI:  
<https://developer.download.nvidia.com/video/gputechconf/gtc/2019/presentation/s9900-irradiance-fields-rtx-diffuse-global-illumination-for-local-and-cloud-graphics.pdf>
- DDGI: <https://www.jcgt.org/published/0008/02/01/paper-lowres.pdf>



## GI

- Reflective Shadow Maps:

<https://users.soe.ucsc.edu/~pang/160/s13/proposal/mijallen/proposal/media/p203-dachsba...pdf>

- Light Propagation Volumes: <https://ericpolman.com/>

- VXGI: <https://on-demand.gputechconf.com/gtc/2014/presentations/S4552-rt-voxel-based-global-illumination-gpus.pdf>

- Lumen:

[https://advances.realtimerendering.com/s2021/Radiance%20Caching%20for%20real-time%20Global%20Illumination%20\(SIGGRAPH%202021\).pptx](https://advances.realtimerendering.com/s2021/Radiance%20Caching%20for%20real-time%20Global%20Illumination%20(SIGGRAPH%202021).pptx)



## Hardware Ray Tracing

- Ray Tracing Gems II: <https://link.springer.com/content/pdf/10.1007/978-1-4842-7185-8.pdf>
- The six levels of ray tracing acceleration:  
<https://f.hubspotusercontent10.net/hubfs/2426966/Gated%20Files/imagination-raytracing-primer-sept2020.pdf>
- Hybrid Rendering for Real-Time Ray Tracing:  
[https://link.springer.com/content/pdf/10.1007/978-1-4842-4427-2\\_25.pdf](https://link.springer.com/content/pdf/10.1007/978-1-4842-4427-2_25.pdf)
- Ray Traced Reflections in 'Wolfenstein: Youngblood':  
<https://www.gdcvault.com/play/1026723/Ray-Traced-Reflections-in-Wolfenstein>
- DirectX Raytracing (DXR) Functional Spec: <https://microsoft.github.io/DirectX-Specs/d3d/Raytracing.html>
- [VulkanRayTracingFinalSpecificationRelease](#)



## Signed Distance Field

- Dynamic Occlusion With Signed Distance Fields:

<http://advances.realtimerendering.com/s2015/DynamicOcclusionWithSignedDistanceFields.pdf>

- Lecture 5 Real-time Environment Mapping:

[https://www.bilibili.com/video/BV1YK4y1T7yY?p=5&vd\\_source=5e38e5c84aa6ff6cff3b802b3eebc8bc](https://www.bilibili.com/video/BV1YK4y1T7yY?p=5&vd_source=5e38e5c84aa6ff6cff3b802b3eebc8bc)

- DX12 渲染管线(5) - 距离场：建场: [https://zhuanlan.zhihu.com/p/89701518?utm\\_id=0](https://zhuanlan.zhihu.com/p/89701518?utm_id=0)

- GPU Gems 2 - Chapter 8. Per-Pixel Displacement Mapping with Distance Functions:

<https://developer.nvidia.com/gpugems/gpugems2/part-i-geometric-complexity/chapter-8-pixel-displacement-mapping-distance-functions>

- Unreal Engine 5.0 Documentation: <https://docs.unrealengine.com/5.0/en-US/mesh-distance-fields-in-unreal-engine/>

- Advances in Real-Time Rendering in Games: Part I - SIGGRAPH 2022:

<https://advances.realtimerendering.com/s2022/index.html>



## Lecture 20 Contributors

- 砚书  
- 灵鑫

- 坤  
- 光哥

- 小老弟

- 峻铖



# Q&A



# Enjoy ;) Coding



Course Wechat

*Please follow us for  
further information*